

OSTEP

Event-based Concurrency (Advanced)

Event-based Concurrency

- A different style of **concurrent programming**
 - Used in *GUI-based applications*, some types of *internet servers*.
- **The problem** that event-based concurrency addresses is two-fold.
 - Managing concurrency correctly in multi-threaded applications.
 - Missing locks, deadlock, and other nasty problems can arise.
 - The developer has little or no control over what is scheduled at a given moment in time.

The Basic Idea: An Event Loop

- The approach:
 - **Wait** for something (i.e., an "*event*") to occur.
 - When it does, **check** what type of event it is.
 - **Do** the small amount of work it requires.
- Example:

```
1  while(1){  
2      events = getEvents();  
3      for( e in events )  
4          processEvent(e); // event handler  
5  }
```

How exactly does an event-based server determine which events are **taking place**.

An Important API: `select()` (or `poll()`)

- Check whether there is any **incoming I/O** that should be attended to.

- `select()`

```
int select(int nfd,  
           fd_set * restrict readfds,  
           fd_set * restrict writefds,  
           fd_set * restrict errorfds,  
           struct timeval * restrict timeout);
```

- Lets a server determine that a **new packet has arrived** and is in need of processing.
- Let the service know when **it is OK to reply**.
- `timeout`
 - `NULL`: Cause `select()` to *block indefinitely* until some descriptor is ready.
 - `0`: Use the call to `select()` to *return immediately*.

Using `select()`

- How to use `select()` to see which network descriptors have incoming messages upon them.

```
1  #include <stdio.h>
2  #include <stdlib.h>
3  #include <sys/time.h>
4  #include <sys/types.h>
5  #include <unistd.h>
6
7  int main(void) {
8      // open and set up a bunch of sockets (not shown)
9      // main loop
10     while (1) {
11         // initialize the fd_set to all zero
12         fd_set readFDs;
13         FD_ZERO(&readFDs);
14
15         // now set the bits for the descriptors
16         // this server is interested in
17         // (for simplicity, all of them from min to max)
18         ...
```

Using `select()` (Cont.)

```
18         int fd;
19         for (fd = minFD; fd < maxFD; fd++)
20             FD_SET(fd, &readFDs);
21
22         // do the select
23         int rc = select(maxFD+1, &readFDs, NULL, NULL, NULL);
24
25         // check which actually have data using FD_ISSET()
26         int fd;
27         for (fd = minFD; fd < maxFD; fd++)
28             if (FD_ISSET(fd, &readFDs))
29                 processFD(fd);
30     }
31 }
```

Why Simpler? No Locks Needed

- The event-based server **cannot be interrupted** by another thread.
 - With a single CPU and an event-based application.
 - It is decidedly **single threaded**.
 - Thus, *concurrency bugs* common in threaded programs **do not manifest** in the basic event-based approach.

A Problem: Blocking System Calls

- What if an event requires that you issue a **system call** that might block?
 - There are no other threads to run: *just the main event loop*
 - The entire server will do just that: *block until the call completes.*
 - Huge potential waste of resources

In event-based systems: no blocking calls are allowed.

A Solution: Asynchronous I/O

- Enable an application to issue an I/O request and **return control immediately** to the caller, before the I/O has completed.
 - Example:

```
struct aiocb {  
    int aio_fildes;           /* File descriptor */  
    off_t aio_offset;         /* File offset */  
    volatile void *aio_buf;   /* Location of buffer */  
    size_t aio_nbytes;        /* Length of transfer */  
};
```

trol

A Solution: Asynchronous I/O (Cont.)

- Asynchronous API:
 - To issue an asynchronous read to a file

```
int aio_read(struct aiocb *aiocbp);
```

- If successful, it returns right away and the application can continue with its work.
- Checks whether the request referred to by `aiocbp` has completed.

```
int aio_error(const struct aiocb *aiocbp);
```

- An application can **periodically pool** the system via `aio_error()`.
 - If it has completed, returns success.
 - If not, `EINPROGRESS` is returned.

A Solution: Asynchronous I/O (Cont.)

- Interrupt
 - Remedy **the overhead to check** whether an I/O has completed
 - Using **UNIX signals** to inform applications when an asynchronous I/O completes.
 - Removing the need to *repeatedly ask the system*.

Another Problem: State Management

- The code of event-based approach is generally **more complicated** to write than *traditional thread-based* code.
 - It must package up some program state for the next event handler to use when the I/O completes.
 - The state the program needs is on the stack of the thread. → **manual stack management**

Another Problem: State Management (Cont.)

- **Example** (an event-based system):

```
int rc = read(fd, buffer, size);  
rc = write(sd, buffer, size);
```

- First **issue** the read asynchronously.
- Then, **periodically check** for completion of the read.
- That call informs us that the **read is complete**.
- How does the event-based server know **what to do**?

Another Problem: State Management (Cont.)

- **Solution:** continuation
 - **Record** the needed information to finish processing this event *in some data structure*.
 - When the event happens (i.e., when the disk I/O completes), **look up** the needed information and process the event.

What is still difficult with Events.

- Systems moved from a single CPU to **multiple CPUs**.
 - Some of the simplicity of the event-based approach disappeared.
- It **does not integrate well** with certain kinds of systems activity.
 - **Ex. Paging:** A server will not make progress until page fault completes (implicit blocking).
- Hard to manage overtime: The exact semantics of various routines changes.
- Asynchronous disk I/O **never quite integrates with asynchronous network I/O** in as simple and uniform a manner as you might think.

ASIDE: Unix Signals

- Provide a way to **communicate with a process**.
 - *HUP* (hang up), *INT*(interrupt), *SEGV*(segmentation violation), and etc.
 - **Example:** When your program encounters a *segmentation violation*, the OS sends it a *SIGSEGV*.

```
#include <stdio.h>
#include <signal.h>
void handle(int arg) {
    printf("stop wakin' me up...\n");
}

int main(int argc, char *argv[]) {
    signal(SIGHUP, handle);
    while (1)
        ; // doin' nothin' except catchin' some sigs
    return 0;
}
```


ASIDE: Unix Signals (Cont.)

- You can send signals to it with the **kill command** line tool.
 - Doing so will *interrupt the main while loop* in the program and run the handler code `handle()`.

```
prompt> ./main &  
[3] 36705  
prompt> kill -HUP 36705  
stop wakin' me up...  
prompt> kill -HUP 36705  
stop wakin' me up...  
prompt> kill -HUP 36705  
stop wakin' me up...
```