

# 운영 체 제

4190.307


2017.2학기

월/수 14:00-15:15

# Prerequisites..

- C Programming
- 컴퓨터구조
- 시스템 프로그래밍
- 자료구조

# Textbook



## Operating Systems

Three Easy Pieces

by

Remzi Arpaci-Dusseau



Andrea Arpaci-Dusseau

PDF file available from



<http://pages.cs.wisc.edu/~remzi/OSTEP/>

Or Hardcopy (come to 302-319-1)

# References



# References



# Teaching Staff

- Instructor : Yeom, Heonyoung (염헌영)
  - Email : yeom AT snu.ac.kr
  - Office : 302-321
  - Office Hours : M W 13-14pm
- TAs
  - Park, Jiwoong (박지웅)    pjbear AT gmail.com
  - Min, Cheol Gi (민철기)    m47341 AT gmail.com
  - Office : 302-319-1

# Course Components

- Lectures
- Programming Projects using Artik-10 board + Tizen 3.0
  - Provide in-depth understanding of the internals of OS
  - Learning by doing !
  - Could be done in 2-3 person teams
- Midterm Exam (10/23)
- Final Exam (12/13)
- No class on 10/4(추석) classes on 10/2, 10/9
- Lab(302-310-2) sessions : to be determined...

# When you need help

- Class eTL
  - Lecture slides
  - Q & A



# Class Rules

- No open laptops
  - Unless you are asked to do something...
- No cell phones in my sight
  - Better have them mute if they are in your possession.
- Use **common senses** !
- No more **begging** for better grades !
  - I have to report the incidents to the president !!!

# Questions answered in this lecture:

What will you do in this course?

What is an OS and why do you want one?

Why study operating systems?

# What happens when a program runs?


A running program executes instructions.

1. The processor **fetches** an instruction from memory.
2. **Decode**: Figure out which instruction this is
3. **Execute**: i.e., add two numbers, access memory, check a condition, jump to function, and so forth.
4. The processor moves on to the **next instruction** and so on.

# Operating System (OS)

Responsible for

- Making it easy to **run** programs
- Allowing programs to **share** memory
- Enabling programs to **interact** with devices



OS

- Software that converts hardware into a useful form for applications
- in charge of making sure the system operates correctly and efficiently

# What DOES OS Provide?

Role #1: Abstraction - Provide standard library for resources

What is a resource?

Anything valuable (e.g., CPU, memory, disk)

What abstraction does modern OS typically provide for each resource?

- CPU: process and/or thread
- Memory: address space
- Disk: files

Advantages of OS providing abstraction?

- Allow applications to reuse common facilities
- Make different devices look the same
- Provide higher-level or more useful functionality

Challenges

- What are the correct abstractions?
- How much of hardware should be exposed?

# What DOES OS Provide?

Role #2: Resource management – Share resources well

Advantages of OS providing resource management?

- Protect applications from one another
- Provide efficient access to resources (cost, time, energy)
- Provide fair access to resources

Challenges

- What are the correct mechanisms?
- What are the correct policies?

# OS Organization

How to cover all the topics relevant to operating systems?

# OS Organization

How to cover all the topics relevant to operating systems?

Virtualization

Concurrency

persistence



# Three PIECES: FIRST

## Virtualization

Make each application believe it has each resource to itself

# Virtualization

The OS takes a **physical resource**(Processor, Memory, Disk...) and transforms it into a **virtual form** of itself.

- The virtual form is more general, powerful and easy-to-use. Sometimes, we refer to the OS as a **virtual machine**.

# System call

System call allows user **to tell the OS what to do.**

The OS provides some interface (APIs, standard library).

A typical OS exports a few hundred system calls.

- Run programs
- Access memory
- Access devices

# The OS is a resource manager

The OS **manage resources** such as

*CPU, memory and disk.*

The OS allows

- Many programs to run → Sharing the CPU
- Many programs to *concurrently* access their own instructions and data → Sharing memory
- Many programs to access devices → Sharing disks

# Virtualizing the CPU

The system has a very large number of virtual CPUs.

- Turning a single CPU into a seemingly infinite number of CPUs.
- Allowing many programs to seemingly run at once  
→ **Virtualizing the CPU**

# Virtualizing the CPU (Cont.)

```
1      #include <stdio.h>
2      #include <stdlib.h>
3      #include <sys/time.h>
4      #include <assert.h>
5      #include "common.h"
6
7      int
8      main(int argc, char *argv[])
9      {
10         if (argc != 2) {
11             fprintf(stderr, "usage: cpu <string>\n");
12             exit(1);
13         }
14         char *str = argv[1];
15         while (1) {
16             Spin(1); // Repeatedly checks the time and
17                      // returns once it has run for a second
18             printf("%s\n", str);
19         }
20         return 0;
21     }
```

**Simple Example(cpu.c): Code That Loops and Prints**

# Virtualizing the CPU (Cont.)

## Execution result 1.

```
prompt> gcc -o cpu cpu.c -Wall
prompt> ./cpu "A"
A
A
A
^C
prompt>
```

**Run forever; Only by pressing "Control-c" can we halt the program**

# Virtualizing the CPU (Cont.)

## Execution result 2.

```
prompt> ./cpu A & ; ./cpu B & ; ./cpu C & ; ./cpu D &  
[1] 7353  
[2] 7354  
[3] 7355  
[4] 7356  
A  
B  
D  
C  
A  
B  
D  
C  
A  
C  
B  
D  
...
```

Even though we have only **one processor**, all four of programs seem to be running **at the same time!**



# Virtualizing Memory

The physical memory is an array of bytes.

A program keeps all of its data structures in memory.

**Read memory** (load):

Specify an address to be able to access the data

**Write memory** (store):

Specify the data to be written to the given address

# Virtualizing Memory (Cont.)

## A program that Accesses Memory (mem.c)

```
1      #include <unistd.h>
2      #include <stdio.h>
3      #include <stdlib.h>
4      #include "common.h"
5
6      int
7      main(int argc, char *argv[])
8      {
9          int *p = malloc(sizeof(int)); // a1: allocate some
                                         memory
10         assert(p != NULL);
11         printf("(%d) address of p: %08x\n",
12              getpid(), (unsigned) p); // a2: print out the
                                         address of the memory
13         *p = 0; // a3: put zero into the first slot of the memory
14         while (1) {
15             Spin(1);
16             *p = *p + 1;
17             printf("(%d) p: %d\n", getpid(), *p); // a4
18         }
19         return 0;
20     }
```

# Virtualizing Memory (Cont.)

The output of the program `mem.c`

```
prompt> ./mem
(2134) memory address of p: 00200000
(2134) p: 1
(2134) p: 2
(2134) p: 3
(2134) p: 4
(2134) p: 5
^C
```

The newly allocated memory is at address `00200000`.  
It updates the value and prints out the result.

# Virtualizing Memory (Cont.)

## Running mem.c multiple times

```
prompt> ./mem & ./mem &  
[1] 24113  
[2] 24114  
(24113) memory address of p: 00200000  
(24114) memory address of p: 00200000  
(24113) p: 1  
(24114) p: 1  
(24114) p: 2  
(24113) p: 2  
(24113) p: 3  
(24114) p: 3  
...
```

It is as if each running program has its **own private memory**.

- Each running program has allocated memory at the same address.
- Each seems to be updating the value at 00200000 independently.

# Virtualizing Memory (Cont.)

Each process accesses its own private **virtual address space**.

- The OS maps **address space** onto the **physical memory**.
- A memory reference within one running program does not affect the address space of other processes.
- Physical memory is a shared resource, managed by the OS.

Three PIECES: SECOND

# **Concurrency:**

How can you handle all the different things at once ?

# Three PIECES: SECOND

## Concurrency:

- Events are occurring simultaneously and may interact with one another
- OS must be able to handle concurrent events

### Easier case

Hide concurrency from **independent** processes

### Trickier case

Manage concurrency with **interacting** processes

- Provide abstractions (locks, semaphores, condition variables, shared memory, critical sections) to processes
- Ensure processes do not deadlock

### Demo

Interacting threads must coordinate access to shared data

# The problem of Concurrency

The OS is juggling **many things at once**, first running one process, then another, and so forth.

Modern **multi-threaded programs** also exhibit the concurrency problem.



# Concurrency Example

## A Multi-threaded Program (thread.c)

```
1      #include <stdio.h>
2      #include <stdlib.h>
3      #include "common.h"
4
5      volatile int counter = 0;
6      int loops;
7
8      void *worker(void *arg) {
9          int i;
10         for (i = 0; i < loops; i++) {
11             counter++;
12         }
13         return NULL;
14     }
15     ...
```

# Concurrency Example (Cont.)

```
16     int
17     main(int argc, char *argv[])
18     {
19         if (argc != 2) {
20             fprintf(stderr, "usage: threads <value>\n");
21             exit(1);
22         }
23         loops = atoi(argv[1]);
24         pthread_t p1, p2;
25         printf("Initial value : %d\n", counter);
26
27         Pthread_create(&p1, NULL, worker, NULL);
28         Pthread_create(&p2, NULL, worker, NULL);
29         Pthread_join(p1, NULL);
30         Pthread_join(p2, NULL);
31         printf("Final value : %d\n", counter);
32         return 0;
33     }
```

The main program creates **two threads**.

- Thread: a function running within the same memory space. Each thread start running in a routine called `worker()`.
- `worker()`: increments a counter

# Concurrency Example (Cont.)

`loops` determines how many times each of the two workers will **increment the shared counter** in a loop.

`loops: 1000.`

```
prompt> gcc -o thread thread.c -Wall -pthread
prompt> ./thread 1000
Initial value : 0
Final value : 2000
```

- `loops: 100000.`

```
prompt> ./thread 100000
Initial value : 0
Final value : 143012 // huh??
prompt> ./thread 100000
Initial value : 0
Final value : 137298 // what the??
```

# Why is this happening?

Increment a shared counter → take three instructions.

1. Load the value of the counter from memory into register.
2. Increment it
3. Store it back into memory

These three instructions do not execute **atomically**.  
→ Problem of **concurrency** happen.

# Three PIECES: Third

## **Persistence:**

**The data should be there even after you turn off the computer !**

# Three PIECES: Third

**Persistence:** Access information permanently

- Lifetime of information is longer than lifetime of any one process
- Machine may be rebooted, machine may lose power or crash unexpectedly

Issues:

- Provide abstraction so applications do not know how data is stored : Files, directories (folders), links
- Correctness with unexpected failures
- Performance: disks are very slow; many optimizations needed!

Demo

File system does work to ensure data updated correctly

# Persistence

Devices such as DRAM store values in a volatile.

*Hardware* and *software* are needed to store data **persistently**.

- **Hardware:** I/O device such as a hard drive, solid-state drives(SSDs)
- **Software:**
  - File system manages the disk.
  - File system is responsible for storing any files the user creates.

# Persistence (Cont.)

Create a file (`/tmp/file`) that contains the string “hello world”

```
1      #include <stdio.h>
2      #include <unistd.h>
3      #include <assert.h>
4      #include <fcntl.h>
5      #include <sys/types.h>
6
7      int
8      main(int argc, char *argv[])
9      {
10         int fd = open("/tmp/file", O_WRONLY | O_CREAT
                        | O_TRUNC, S_IRWXU);
11         assert(fd > -1);
12         int rc = write(fd, "hello world\n", 13);
13         assert(rc == 13);
14         close(fd);
15         return 0;
16     }
```

`open()`, `write()`, and `close()` system calls are routed to the part of OS called the file system, which handles the requests



# Persistence (Cont.)

What does OS do in order to write to disk?

- Figure out **where** on disk this new data will reside
- **Issue I/O** requests to the underlying storage device

File system handles system crashes during write.

- **Journaling** or **copy-on-write**
- Carefully ordering writes to disk

# Design Goals

## Build up **abstraction**

- Make the system convenient and easy to use.

## Provide high **performance**

- Minimize the overhead of the OS.
- OS must strive to provide virtualization without excessive overhead.

## **Protection** between applications

- Isolation: Bad behavior of one does not harm other and the OS itself.

# Design Goals (Cont.)

## High degree of **reliability**

- The OS must also run non-stop.

## Other issues

- Energy-efficiency
- Security
- Mobility

# Advanced Topics

## Current systems

- Multiprocessors
- Networked and distributed systems
- Virtual machines

# Why study Operating Systems ?

Build, modify, or administer an operating system

Understand system performance

- Behavior of OS impacts entire machine
- Tune workload performance
- Apply knowledge across many layers

Computer architecture, programming languages, data structures and algorithms, and performance modeling

Fun and challenging to understand large, complex systems

Disclaimer: This lecture slide set was made by merging/modifying/revising below three.

- 1) Andrea's at University of Wisconsin.
- 2) Tyler's at University of Wisconsin
- 3) Youzip's Operating System course in Computer Science Dept. at Hanyang University