

# OSTEP

## Memory Virtualization

### Virtual Memory

#### **Questions answered in this lecture:**

What is in the address space of a process (review)?

What are the different ways that that OS can virtualize memory?

Time sharing, static relocation, dynamic relocation

(base, base + bounds, segmentation)

What hardware support is needed for dynamic relocation?

# More Virtualization

1<sup>st</sup> part of course: Virtualization

Virtual CPU: *illusion* of **private CPU registers**

- 2 lectures (mechanism + policy)

Virtual RAM: *illusion* of **private memory**

- 5 lectures

# Memory Virtualization

What is **memory virtualization**?

- OS virtualizes its physical memory
- OS provides an **illusion memory space** per each process
- It seems to be seen like **each process uses the whole memory**

# Benefit of Memory Virtualization

Ease of use in programming


Memory efficiency in terms of **times** and **space**

The guarantee of isolation for processes as well as OS

- Protection from **errant accesses** of other processes

# Motivation for Virtualization

Uniprogramming: One process runs at a time



## Disadvantages:

- Only one process runs at a time
- Process can destroy OS
- Poor utilization and efficiency

# Multiprogramming Goals

## Transparency

- Processes are not aware that memory is shared
- Works regardless of number and/or location of processes

## Protection

- Cannot corrupt OS or other processes
- Privacy: Cannot read data of other processes

## Efficiency

- Do not waste memory resources (minimize fragmentation)

## Sharing

- Cooperating processes can share portions of address space


# Multiprogramming and Time Sharing

**Load multiple processes** in memory.

- Execute one for a short while.
- Switch processes between them in memory.
- Increase utilization and efficiency.

Cause an important **protection issue**.

- Errant memory accesses from other processes



**Physical Memory**

# Abstraction: Address Space

Address space: Each process has set of addresses that **map** to bytes

Problem:

How can OS provide illusion of **private** address space to each process?


Review: What is in an **address space**?

Extend LDE (limited direct execution)

Address space has static and dynamic components

Static: Code and some global variables

Dynamic: Stack and Heap





# Address Space(CONT.)

## Code


- Where instructions live

## Heap

- Dynamically allocate memory.  
    `malloc` in C language  
    `new` in object-oriented language

## Stack

- Store return addresses or values
- Contain local variables arguments to routines



# Virtual Address

**Every address** in a running program is virtual.

- OS translates the virtual address to physical address

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]){

    printf("location of code   : %p\n", (void *) main);
    printf("location of heap   : %p\n", (void *) malloc(1));
    int x = 3;
    printf("location of stack  : %p\n", (void *) &x);

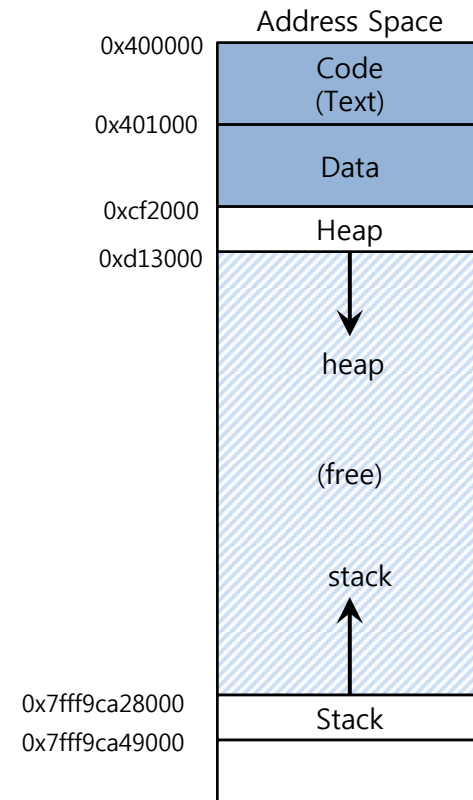
    return x;
}
```

**A simple program that prints out addresses**

# Virtual Address(Cont.)

The output in 64-bit Linux machine

```
location of code   : 0x40057d
location of heap   : 0xcf2010
location of stack  : 0x7fff9ca45fcc
```



# Motivation for Dynamic Memory

Why do processes need dynamic allocation of memory?

- Do not know amount of memory needed at compile time
- Must be pessimistic when allocate memory statically
  - Allocate enough for worst possible case; Storage is used inefficiently

Recursive procedures

- Do not know how many times procedure will be nested

Complex data structures: lists and trees

- `struct my_t *p = (struct my_t *)malloc(sizeof(struct my_t));`

Two types of dynamic allocation

- Stack
- Heap

# Stack Organization

Definition: Memory is freed in opposite order from allocation

```
alloc(A);  
alloc(B);  
alloc(C);  
free(C);  
alloc(D);  
free(D);  
free(B);  
free(A);
```

Simple and efficient implementation:

Pointer separates allocated and freed space

Allocate: Increment pointer

Free: Decrement pointer

No fragmentation

# Where Are stacks Used?

OS uses stack for procedure call frames (local variables and parameters)

```
main () {  
    int A = 0;  
    foo (A);  
    printf("A: %d\n", A);  
}  
  
void foo (int Z) {  
    int A = 2;  
    Z = 5;  
    printf("A: %d Z: %d\n", A, Z);  
}
```

# Heap Organization

Definition: Allocate from any random location: malloc(), new()

- Heap memory consists of allocated areas and free areas (holes)
- Order of allocation and free is unpredictable

## Advantage

- Works for all data structures

## Disadvantages

- Allocation can be slow
- End up with small chunks of free space - fragmentation
- Where to allocate 12 bytes? 16 bytes? 24 bytes??
- What is OS's role in managing heap?
  - OS gives big chunk of free memory to process; library manages individual allocations



# Quiz: Match that Address Location

```
int x;  
int main(int argc, char *argv[]) {  
    int y;  
    int *z = malloc(sizeof(int));  
}
```

Possible segments: static data, code, stack, heap

What if no static data segment?

Address	Location
x	Static data → Code
main	Code
y	Stack
z	Stack
*z	Heap




# Memory Accesses

```
#include <stdio.h>
#include <stdlib.h>

int main(int argc, char *argv[]) {
    int x;
    x = x + 3;
}
```

otool -tv demo1.o  
(or objdump on Linux)



0x10: movl 0x8(%rbp), %edi  
0x13: addl \$0x3, %edi  
0x19: movl %edi, 0x8(%rbp)

**%rbp** is the base pointer:  
points to base of current stack frame

# Quiz: Memory Accesses?

Initial %rip = 0x10

%rbp = 0x200

➡ 0x10: movl 0x8(%rbp), %edi  
➡ 0x13: addl \$0x3, %edi  
➡ 0x19: movl %edi, 0x8(%rbp)

**%rbp** is the base pointer:  
points to base of current stack frame

**%rip** is instruction pointer (or program counter)

**Memory Accesses to what addresses?**

Fetch instruction at addr **0x10**

Exec:

load from addr **0x208**

Fetch instruction at addr **0x13**

Exec:

no memory access

Fetch instruction at addr **0x19**

Exec:

store to addr **0x208**

# How to Virtualize Memory?

Problem: How to run multiple processes simultaneously?

Addresses are “hardcoded” into process binaries

How to avoid collisions?

Possible Solutions for Mechanisms (covered today):

1. Time Sharing
2. Static Relocation
3. Base
4. Base+Bounds
5. Segmentation



# 1) Time Sharing of Memory

Try similar approach to how OS virtualizes CPU


Observation:



OS gives illusion of many virtual CPUs by saving **CPU registers** to **memory** when a process isn't running


Could give illusion of many virtual memories by saving **memory** to **disk** when process isn't running




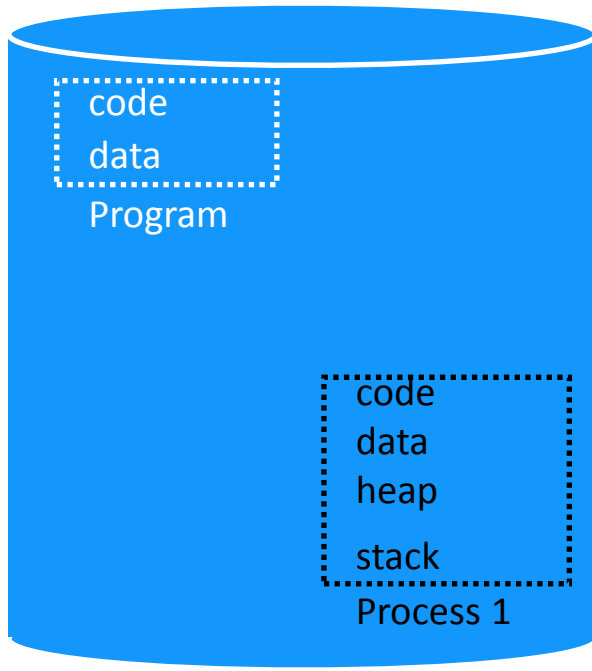
# Time Share Memory: Example

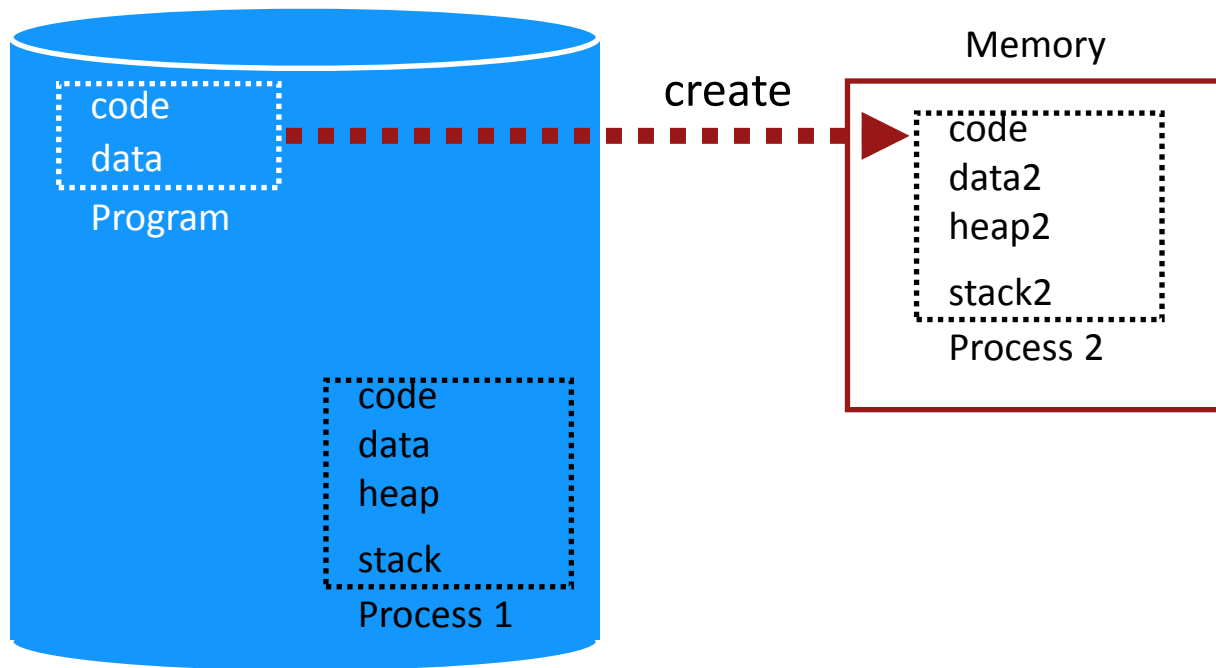









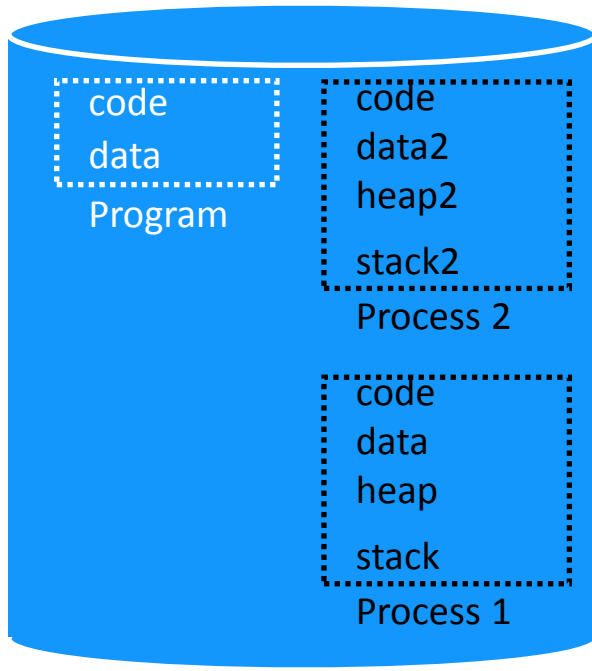










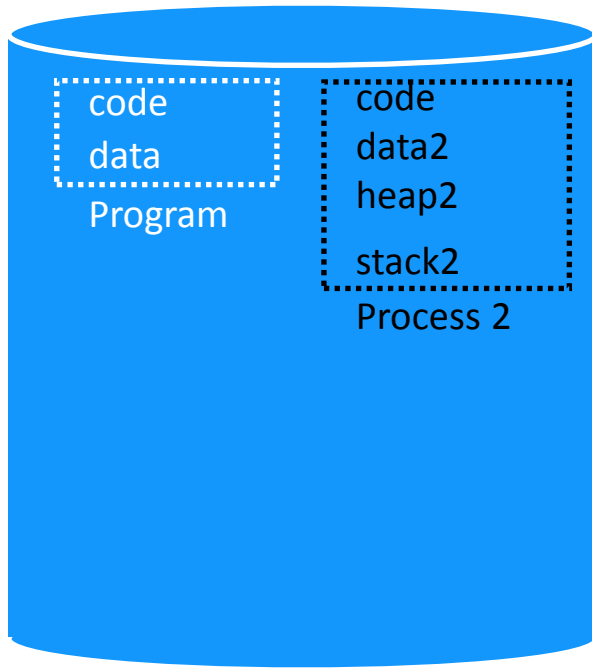




Memory







# Problems with Time Sharing Memory

Problem: Ridiculously poor performance

Better Alternative: space sharing

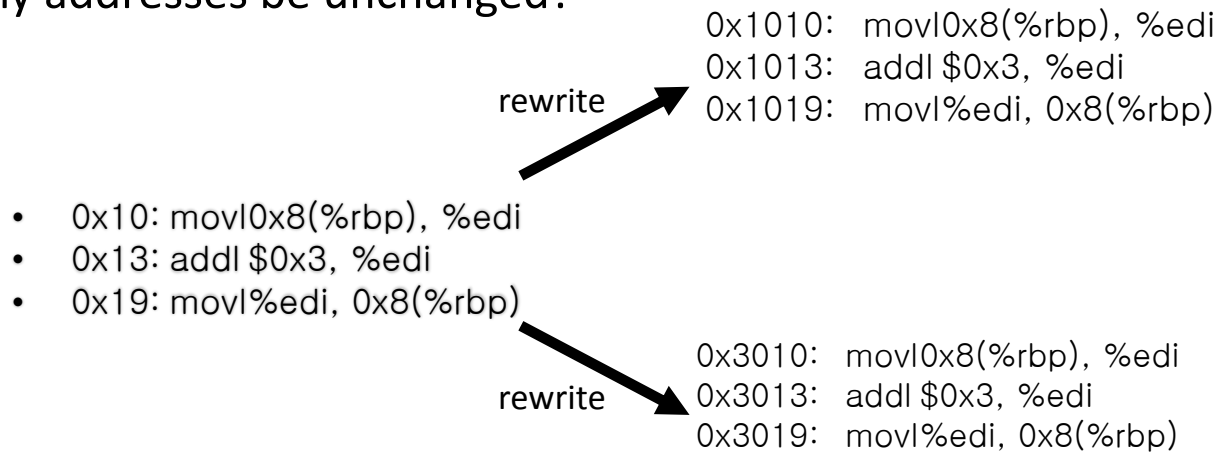
- At same time, space of memory is divided across processes

Remainder of solutions all use space sharing

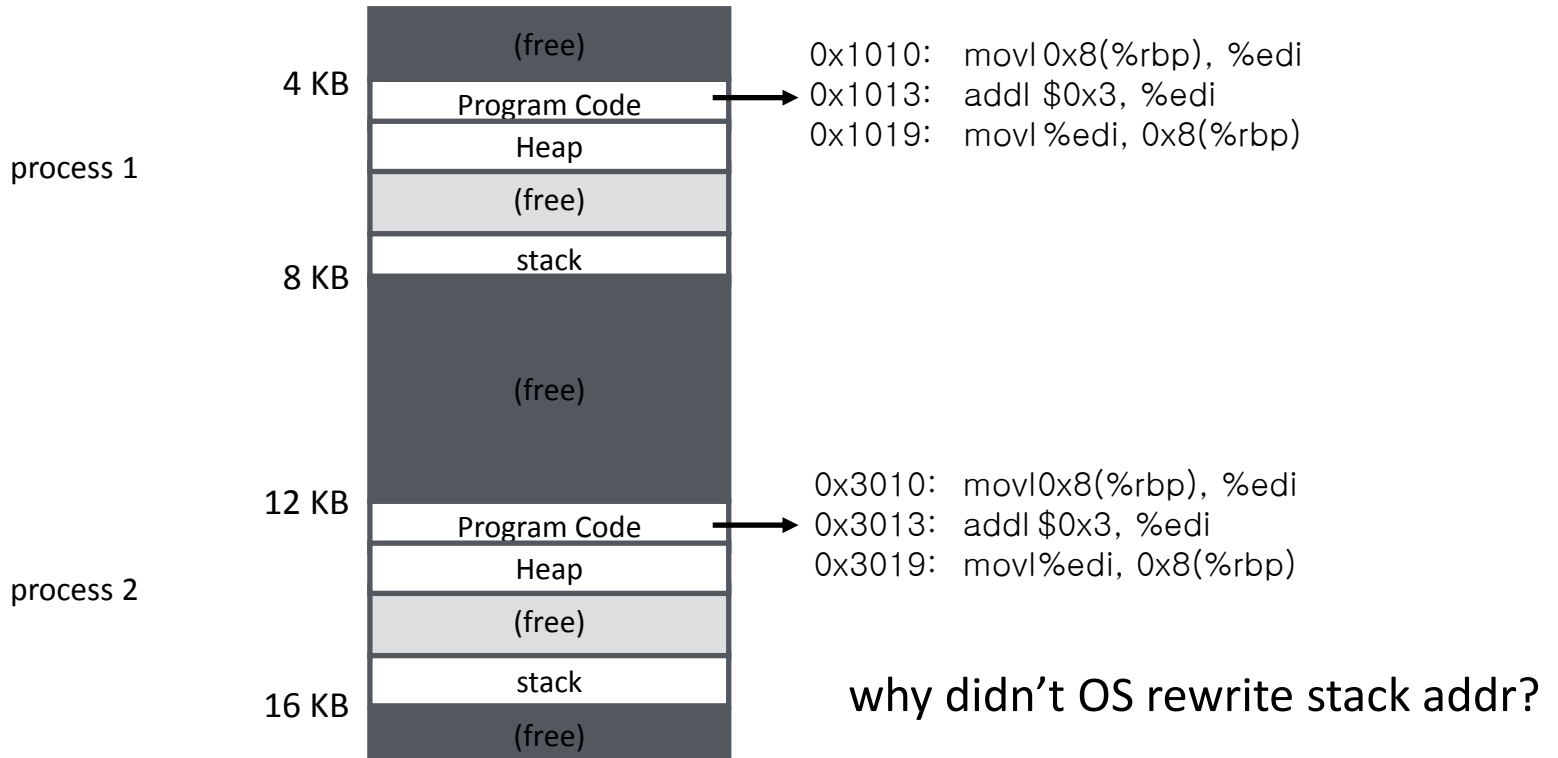


## 2) Static Relocation

- Idea: OS rewrites each program before loading it as a process in memory
- Each rewrite for different process uses different addresses and pointers
- Change jumps, loads of static data  
Can any addresses be unchanged?



# Static: Layout in Memory



# Static Relocation:

## Disadvantages

### No protection

- Process can destroy OS or other processes
- No privacy

### Cannot move address space after it has been placed

- May not be able to allocate new process

# 3) Dynamic Relocation

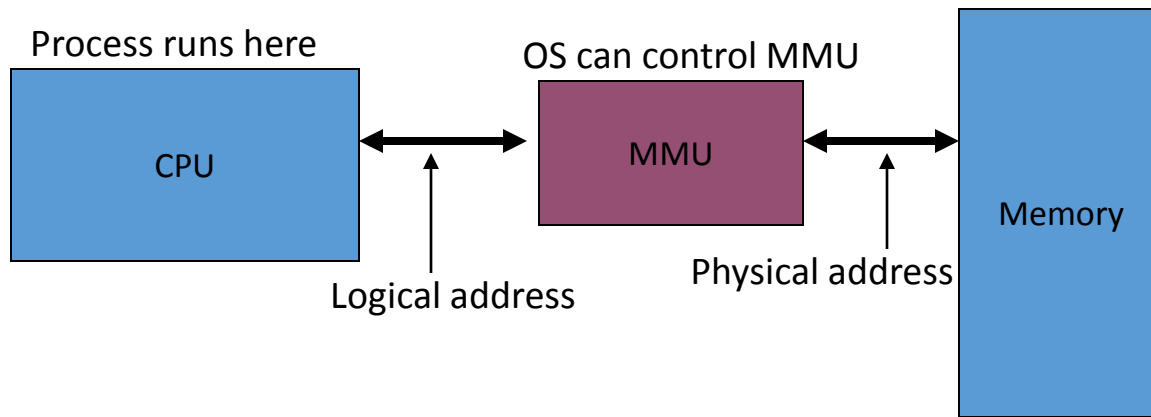
Goal: Protect processes from one another

Requires hardware support

- Memory Management Unit (MMU)

MMU dynamically changes process address at every memory reference

- Process generates **logical** or **virtual** addresses (in their address space)
- Memory hardware uses **physical** or **real** addresses



# Hardware Support for Dynamic Relocation

## Two operating modes

Privileged (protected, kernel) mode: OS runs

- When enter OS (trap, system calls, interrupts, exceptions)
- Allows certain instructions to be executed
  - Can manipulate contents of MMU
- **Allows OS to access all of physical memory**

User mode: User processes run

- **Perform translation of logical address to physical address**


Minimal MMU contains **base register** for translation

base: start location for address space

# Implementation of Dynamic Relocation: BASE REG

Translation on every memory access of user process

- MMU adds base register to logical address to form physical address




# Dynamic Relocation with Base Register

Idea: translate virtual addresses to physical by adding a fixed offset each time.

Store offset in base register


Each process has different value in base register




VISUAL Example of DYNAMIC RELOCATION:

BASE REGISTER






P1 is running




P2 is running




(Decimal notation)


Virtual	Physical
P1: load 100, R1	




Virtual	Physical	
P1: load 100, R1	load 1124, R1	(1024 + 100)




Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	



Virtual	Physical	
P1: load 100, R1	load 1124, R1	
P2: load 100, R1	load 4196, R1	(4096 + 100)




Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	




Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1





Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 100, R1	




Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 1000, R1	load 2024, R1

# Quiz: Who Controls the Base Register?

What entity should **do translation** of addresses with base register?  
(1) process, (2) OS, or (3) HW

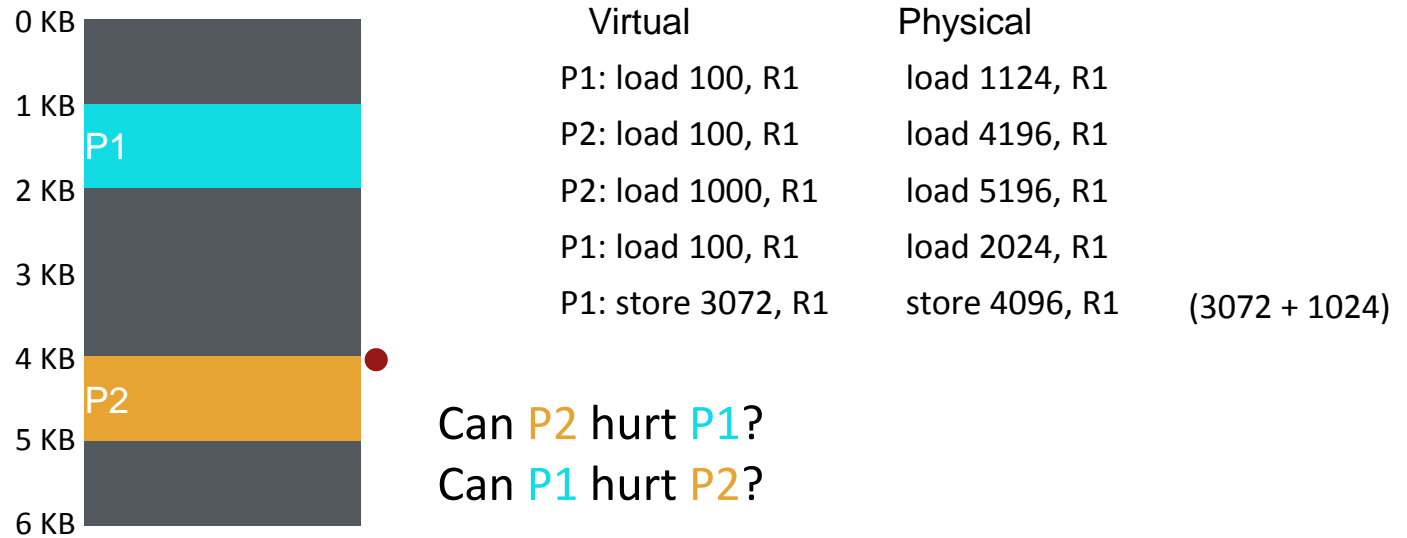
What entity should **modify** the base register?  
(1) process, (2) OS, or (3) HW



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 100, R1	load 2024, R1

Can P2 hurt P1?  
Can P1 hurt P2?

How well does dynamic relocation do with base register for protection?



How well does dynamic relocation do with base register for protection?

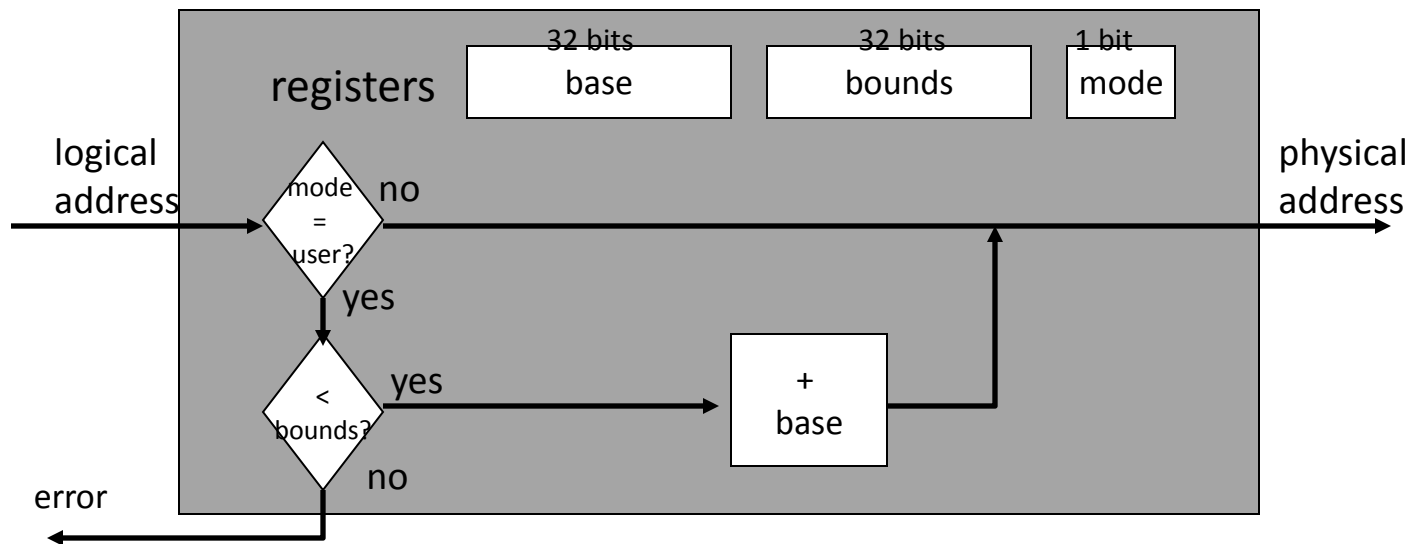
## 4) Dynamic with Base+Bounds

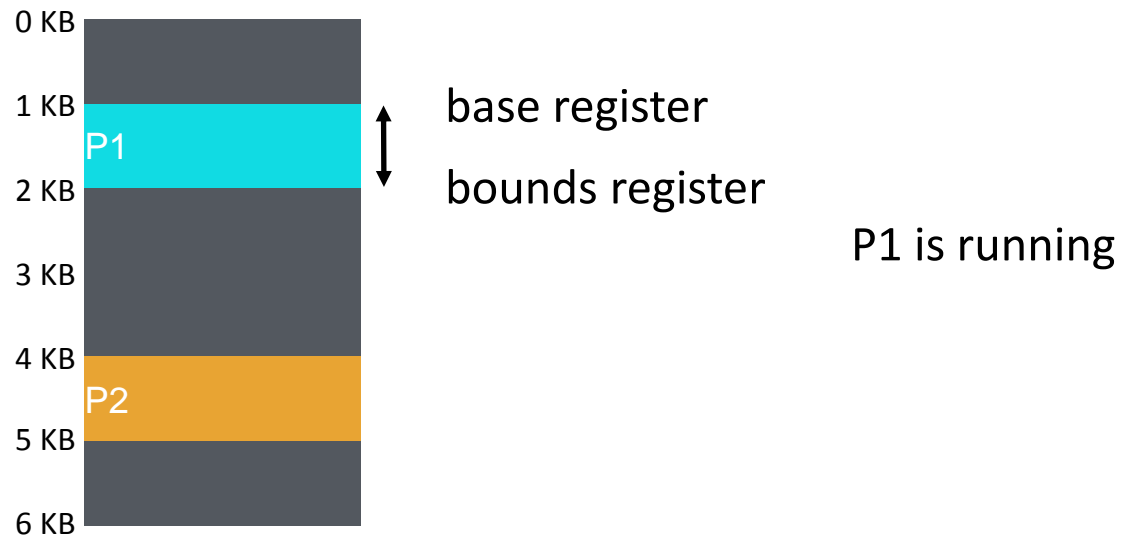
- Idea: limit the address space with a bounds register
- **Base register**: smallest physical addr (or starting location)
- **Bounds register**: size of this process's virtual address space
  - Sometimes defined as largest physical address (base + size)
- What happens if you load/store after bounds?  
OS kills process if process loads/stores beyond bounds

# Implementation of BASE+BOUNDS


Translation on every memory access of user process

- MMU compares logical address to bounds register
  - if logical address is greater, then generate error
- MMU adds base register to logical address to form physical address











base register  
bounds register

P2 is running



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 100, R1	load 2024, R1
P1: store 3072, R1	


Can P1 hurt P2?



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 100, R1	load 2024, R1
P1: store 3072, R1	interrupt OS!

3072 > 1024

Can P1 hurt P2?



Virtual	Physical
P1: load 100, R1	load 1124, R1
P2: load 100, R1	load 4196, R1
P2: load 1000, R1	load 5196, R1
P1: load 100, R1	load 2024, R1
P1: store 3072, R1	interrupt OS!

Can P1 hurt P2?

# Managing Processes with Base and Bounds

## Context-switch

- Add base and bounds registers to PCB
- Steps
  - Change to privileged mode
  - Save base and bounds registers of old process
  - Load base and bounds registers of new process
  - Change to user mode and jump to new process

What if don't change base and bounds registers when switch?

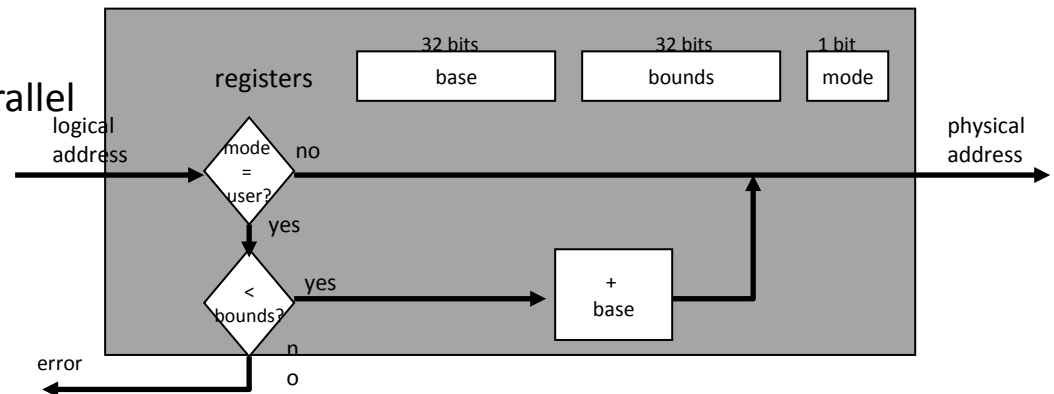
## Protection requirement

- User process cannot change base and bounds registers
- User process cannot change to privileged mode

# Base and Bounds

## Advantages


- Provides protection (both read and write) across address spaces
- Supports dynamic relocation
  - Can place process at different locations initially and also move address spaces
- Simple, inexpensive implementation
  - Few registers, little logic in MMU
- Fast
  - Add and compare in parallel



# Base and Bounds

## Disadvantages

- Each process must be allocated contiguously in physical memory
  - Must allocate memory that may not be used by process
- No partial sharing: Cannot share limited parts of address space




# 5) Segmentation

Divide address space into logical segments

- Each segment corresponds to logical entity in address space
  - code, stack, heap

Each segment can independently:

- be placed separately in physical memory
- grow and shrink
- be protected  
(separate read/write/execute protection bits)



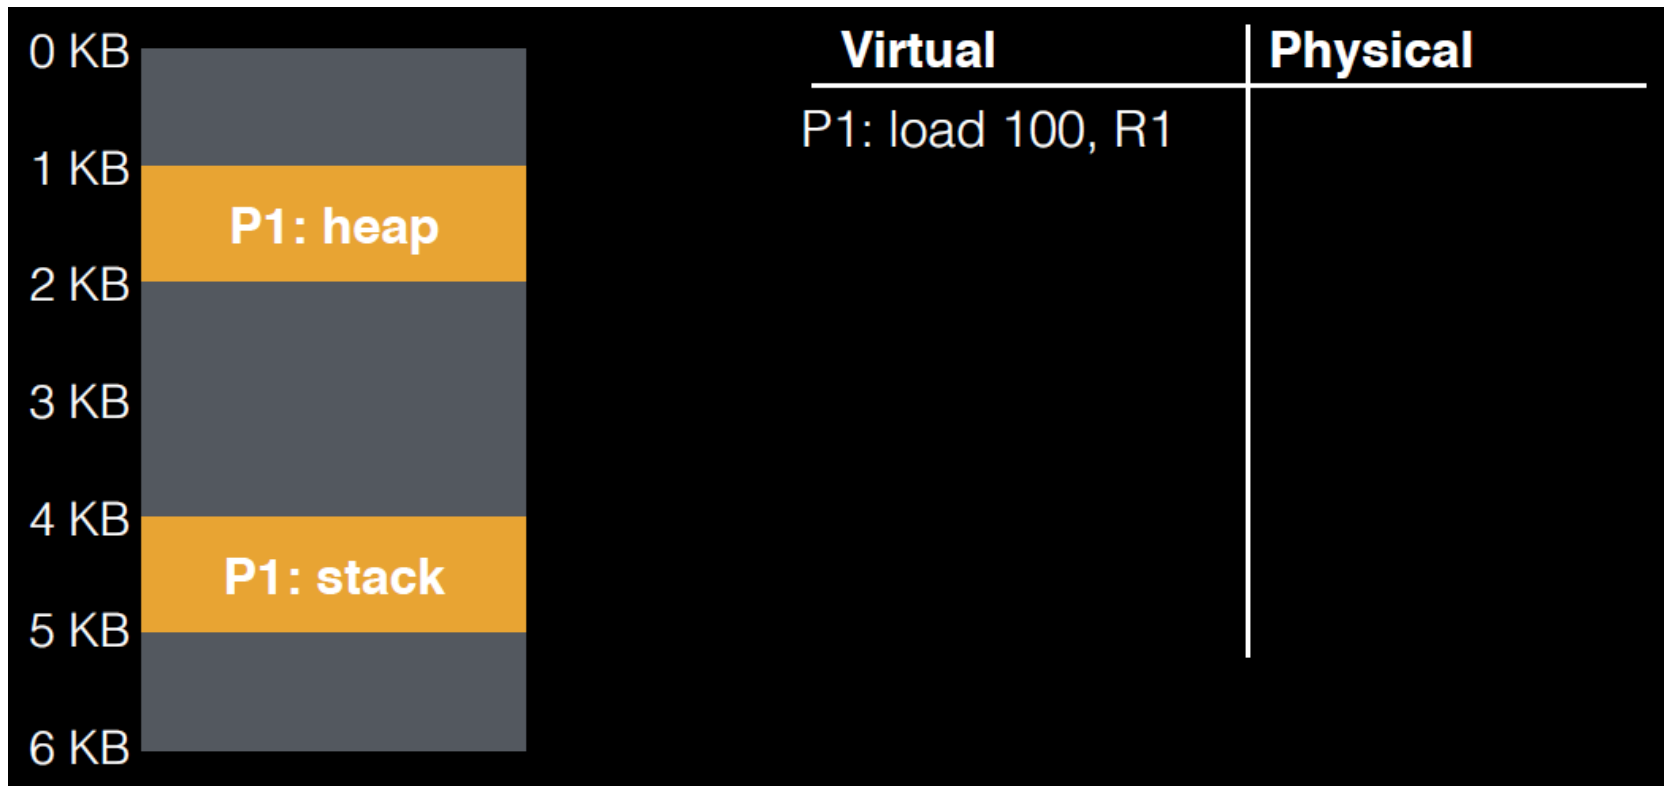


# Multi-segment translation

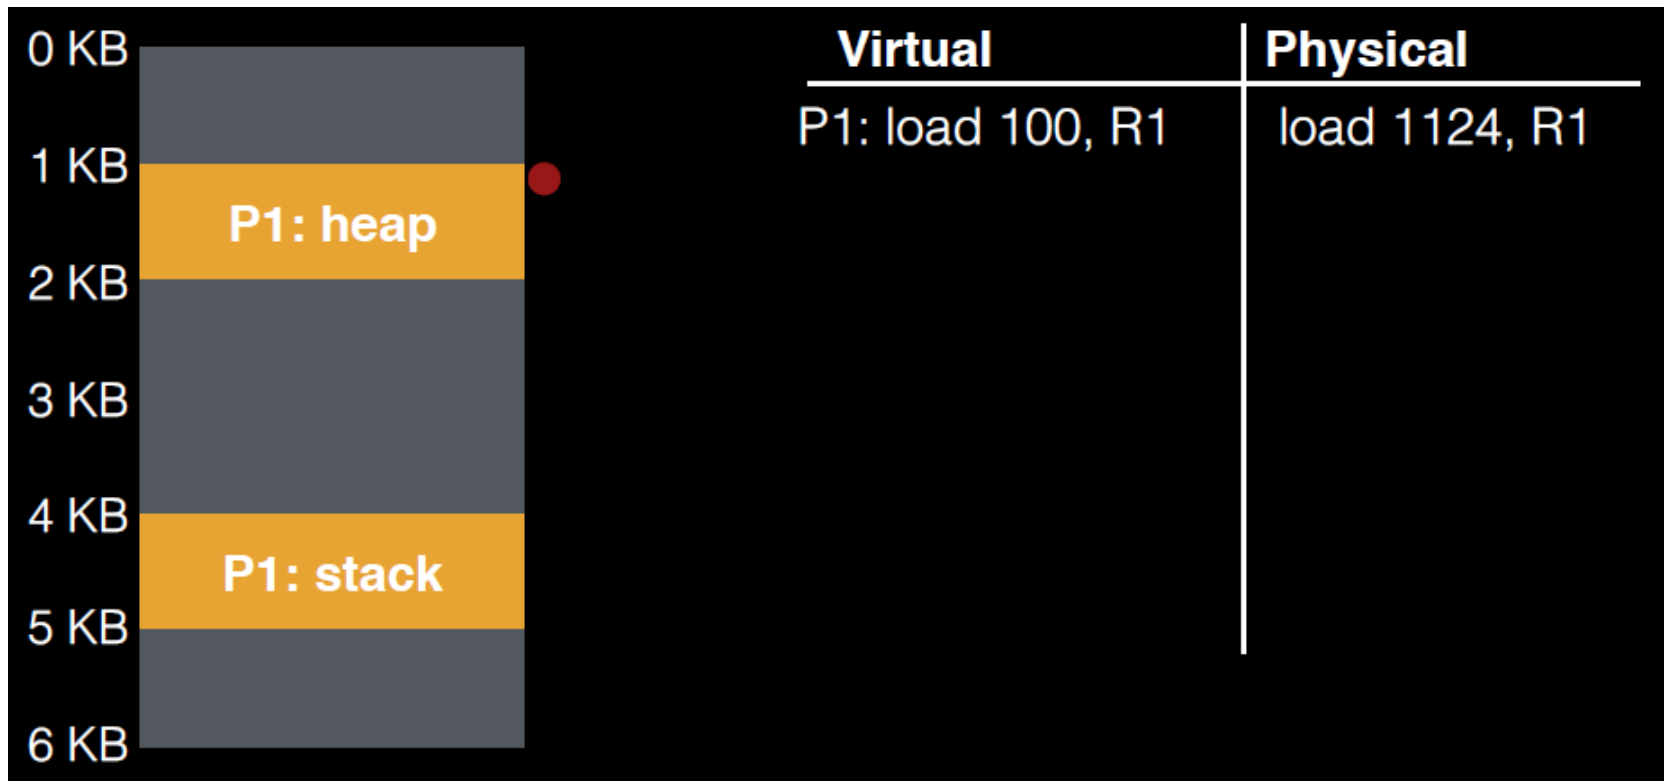
One (broken) approach:

- have **no gaps** in virtual addresses
- map as many low addresses to the first segment as possible, then as **many** as possible to the second (on so on)

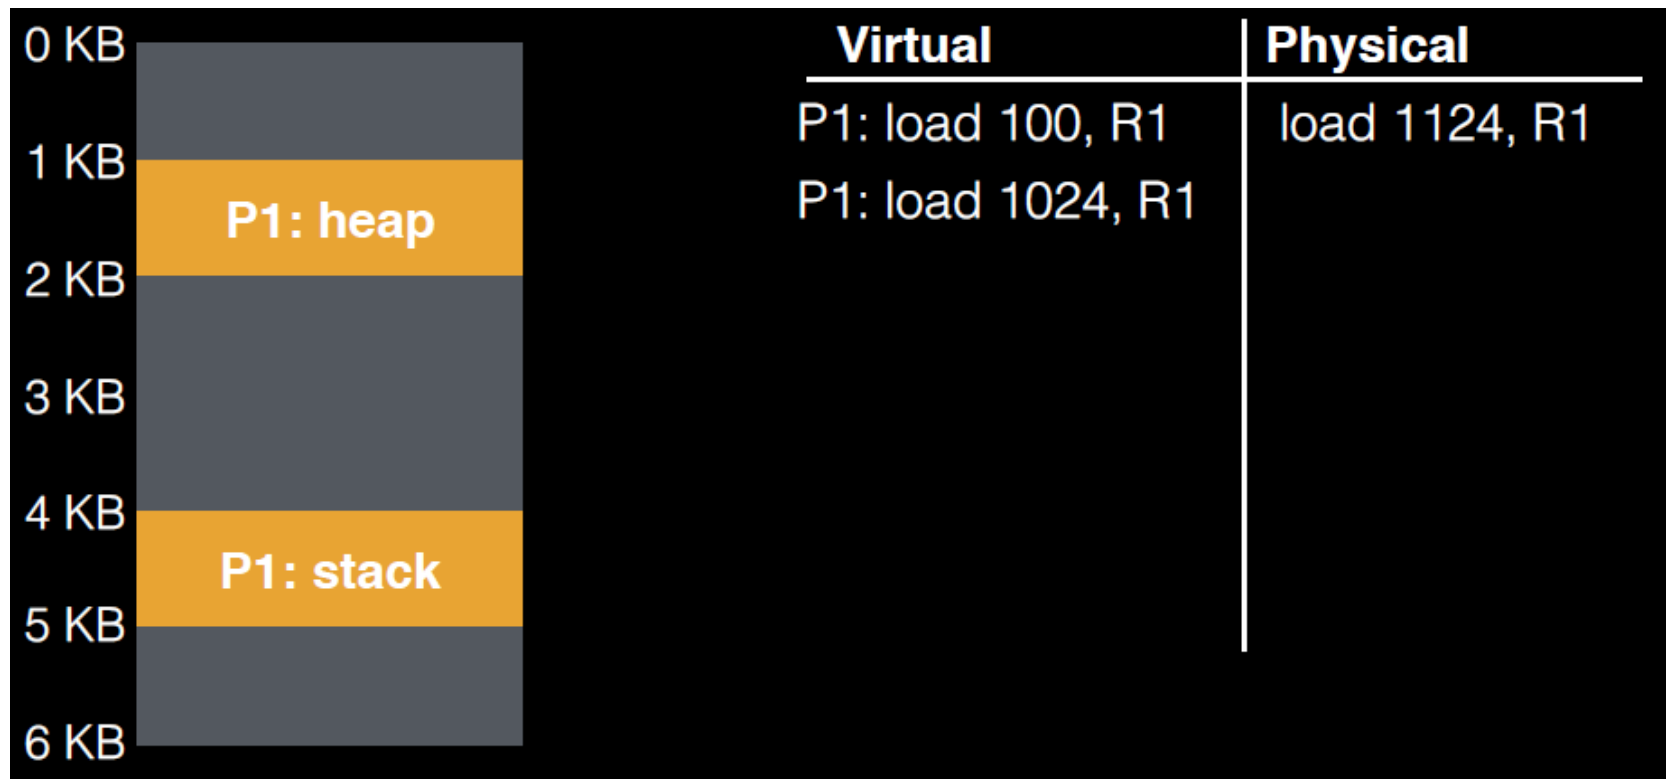
# A tricky example



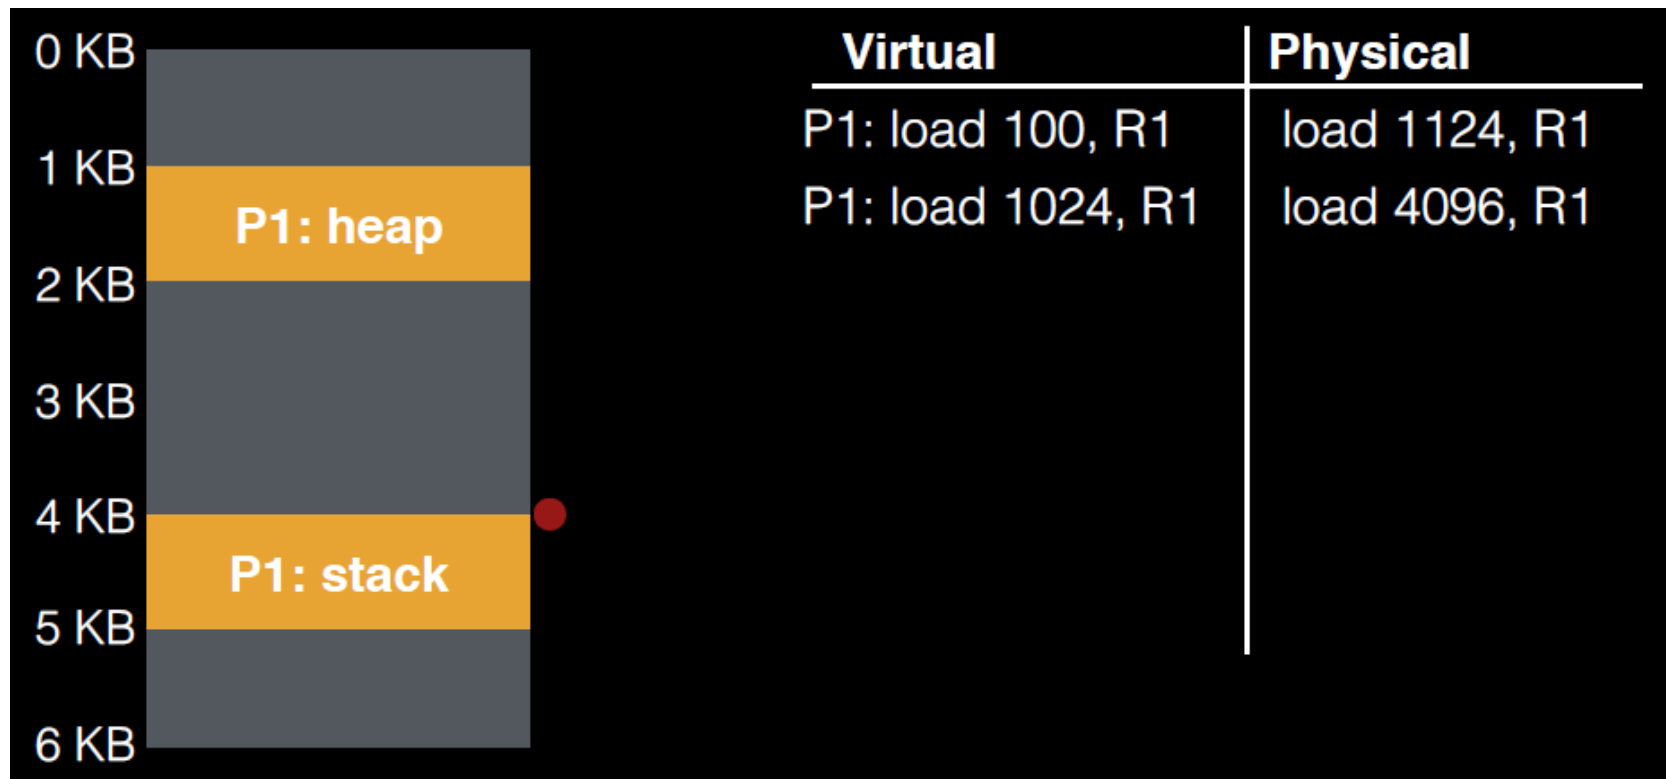
# A tricky example



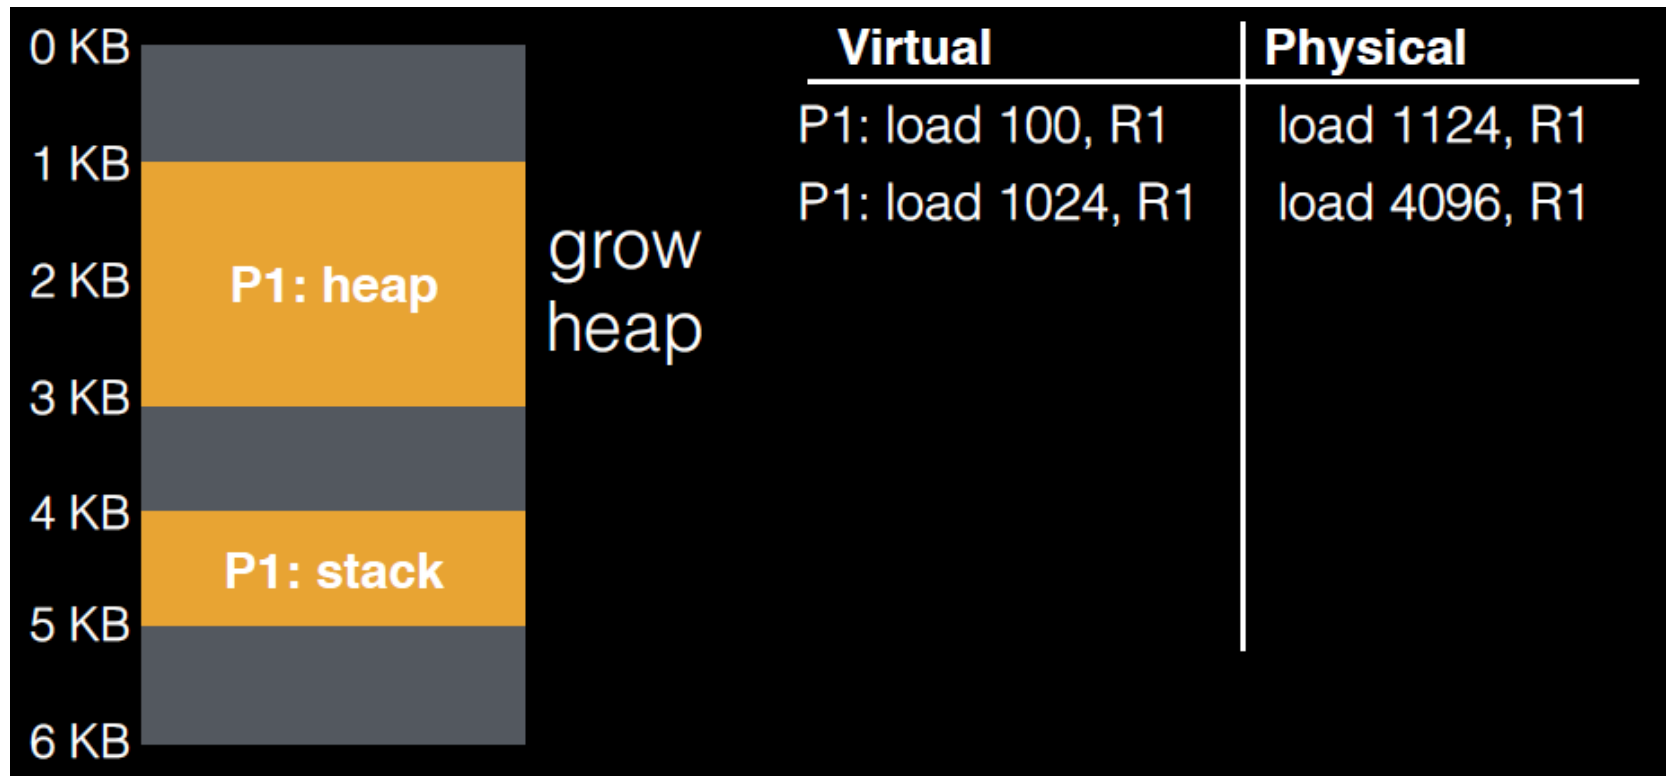
# A tricky example




# A tricky example



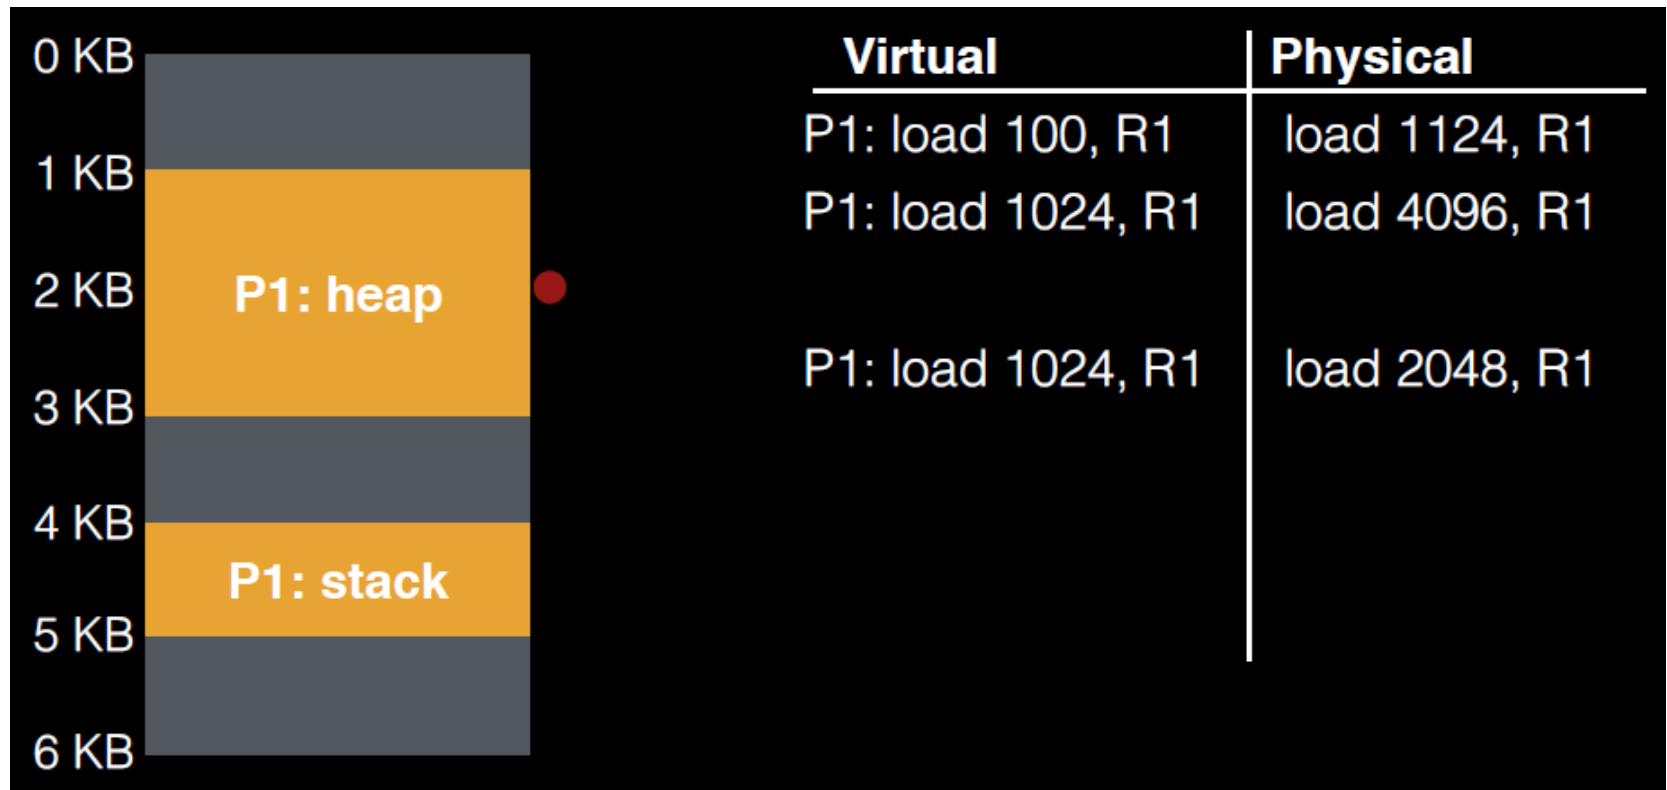
# A tricky example



# A tricky example



# A tricky example





# Multi-segment translation

One (correct) approach:

- break virtual addresses into **two parts**
- one part indicates **segment**
- one part indicates **offset within segment**

# Virtual Address

For example, say addresses are 14 bits.

Use 2 bits for **segment**, 12 bits for **offset**


An address might look like 201E

# Virtual Address

For example, say addresses are 14 bits.

Use 2 bits for **segment**, 12 bits for **offset**

An address might look like **2 01E**




# Virtual Address

For example, say addresses are 14 bits.

Use 2 bits for **segment**, 12 bits for **offset**

An address might look like **2 01E**



Choose some segment numbering, such as

0: code+data

1: heap

2: stack

# What is the segment/offset?

Segment numbers:

0: code+data


1: heap

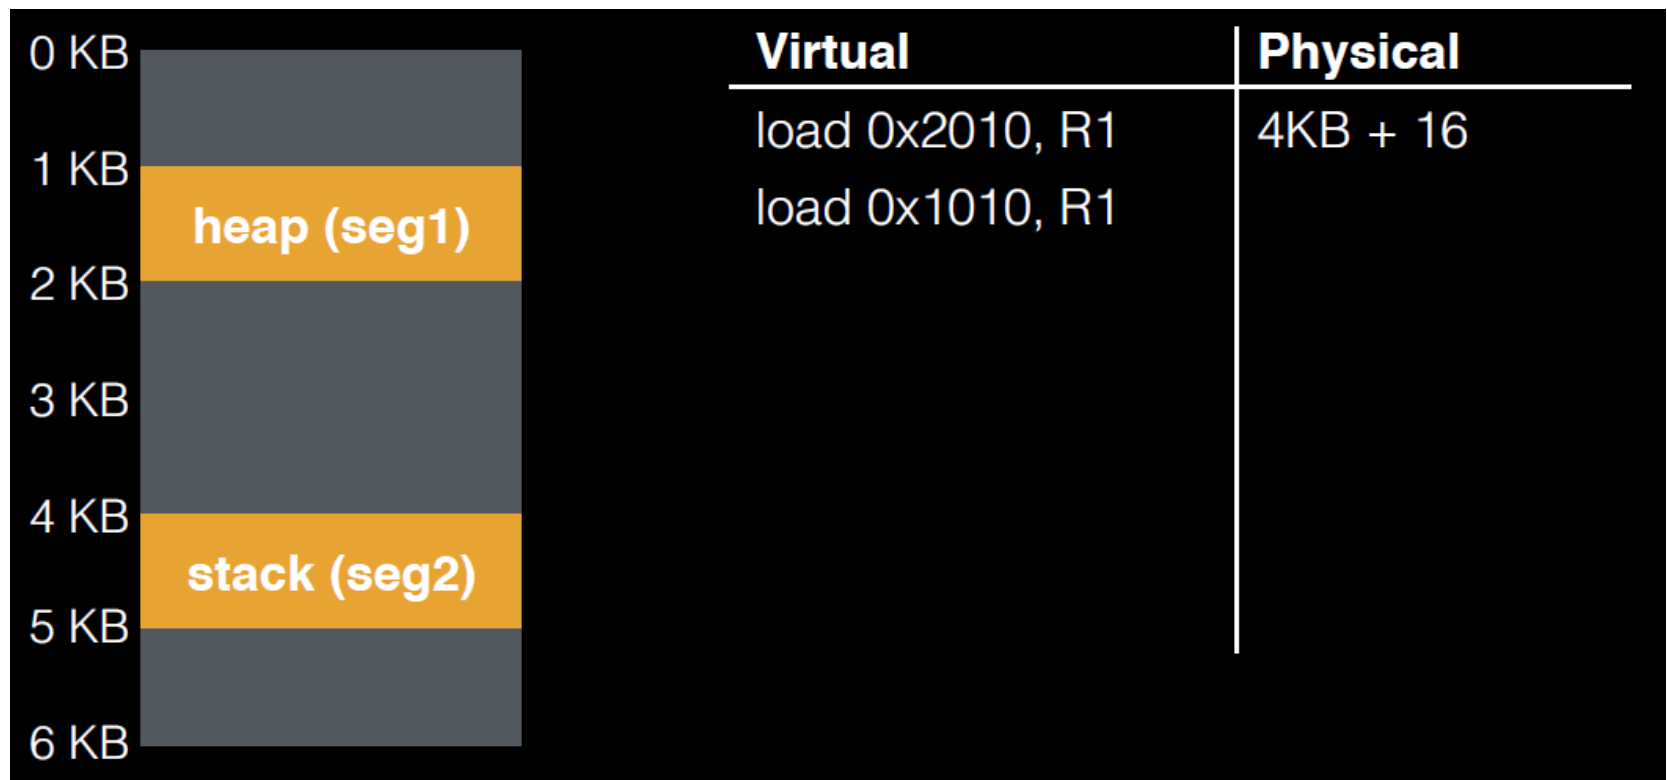
2: stack

10 0000 0001 0001 (binary)

110A (hex)

4096 (decimal)



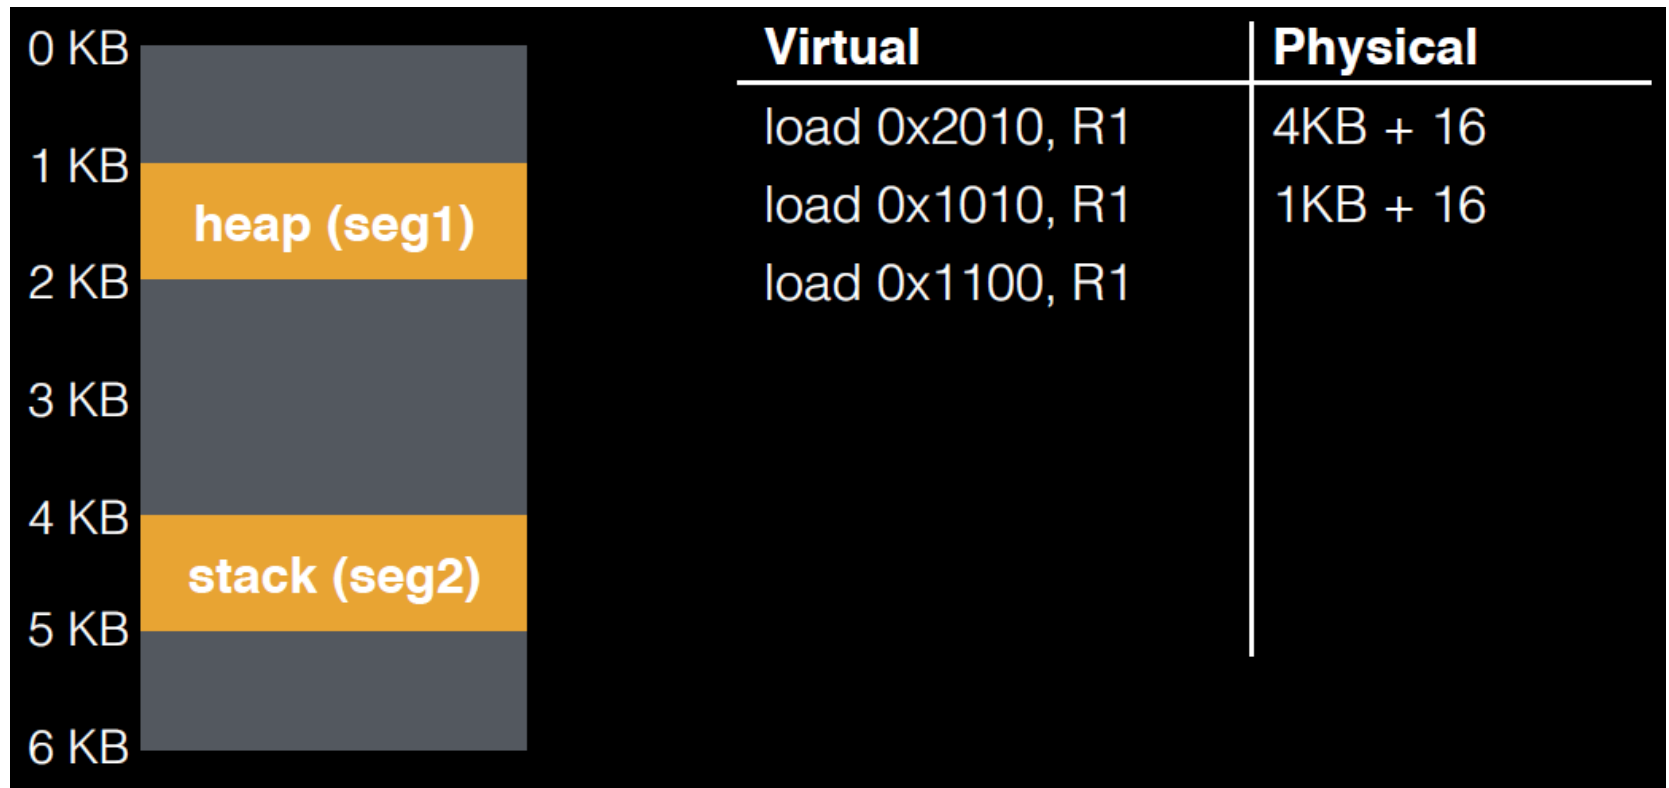


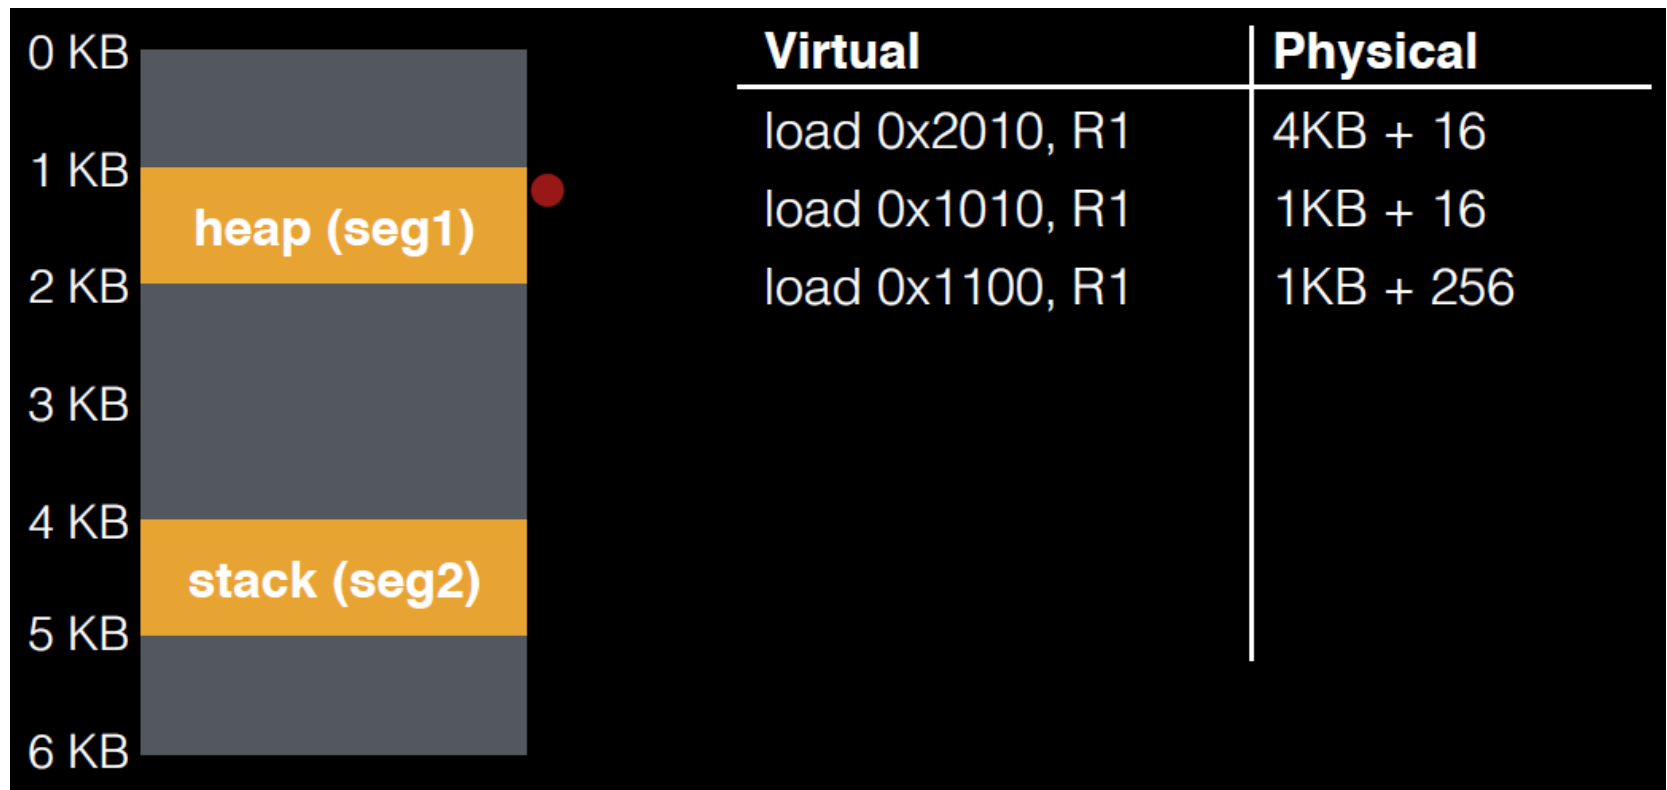


Virtual
load 0x2010, R1
load 0x1010, R1

Physical
4KB + 16
1KB + 16








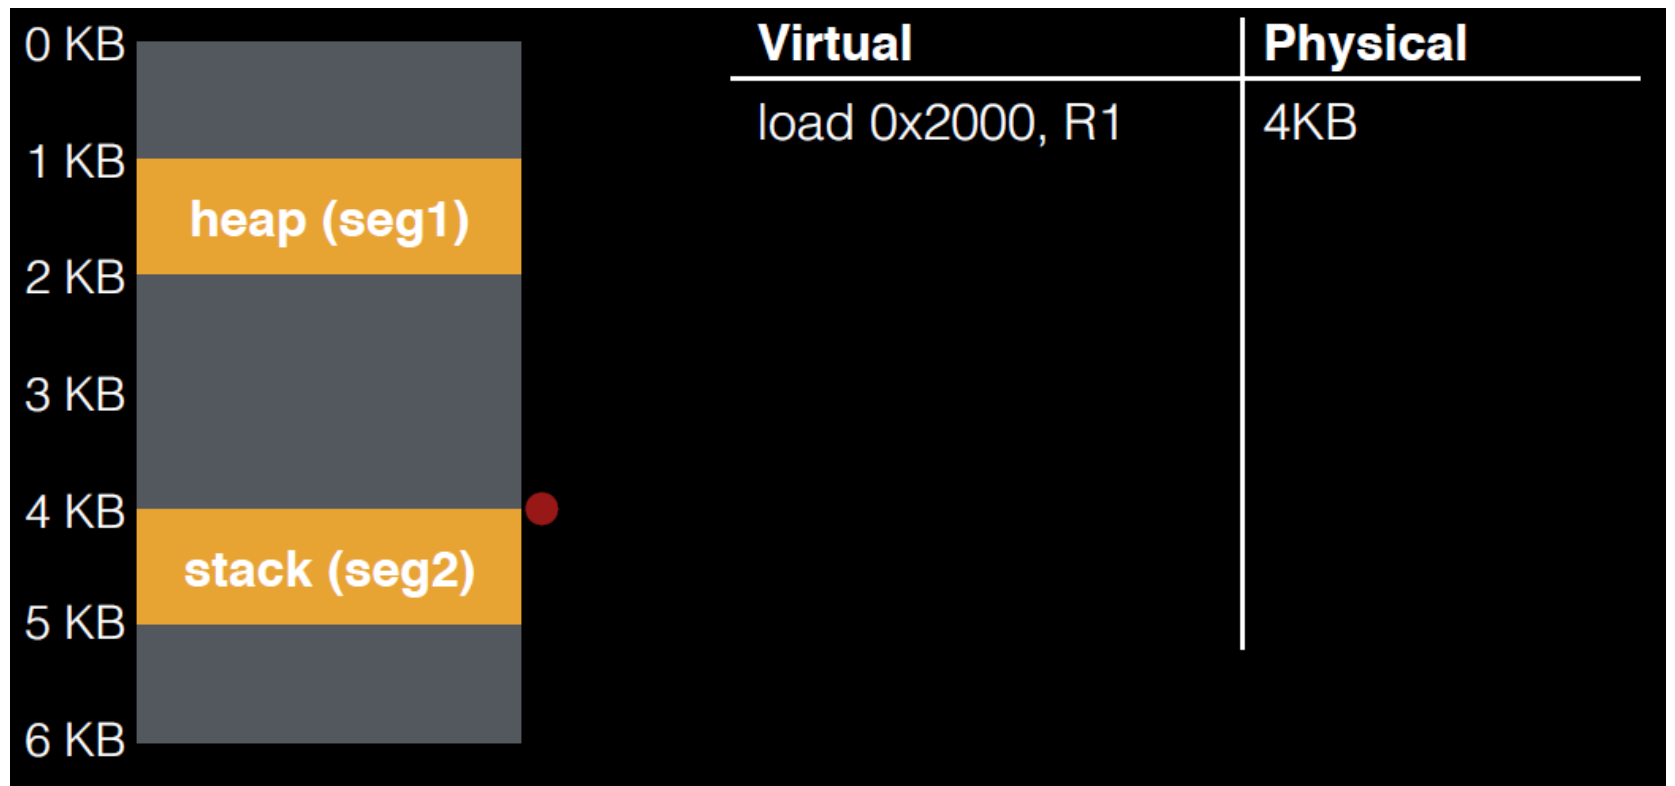
0 KB		<b>Virtual</b>	<b>Physical</b>
1 KB	heap (seg1)	load 0x2010, R1	4KB + 16
2 KB		load 0x1010, R1	1KB + 16
3 KB		load 0x1100, R1	1KB + 256
4 KB	stack (seg2)	load 0x1400, R1	
5 KB			
6 KB			

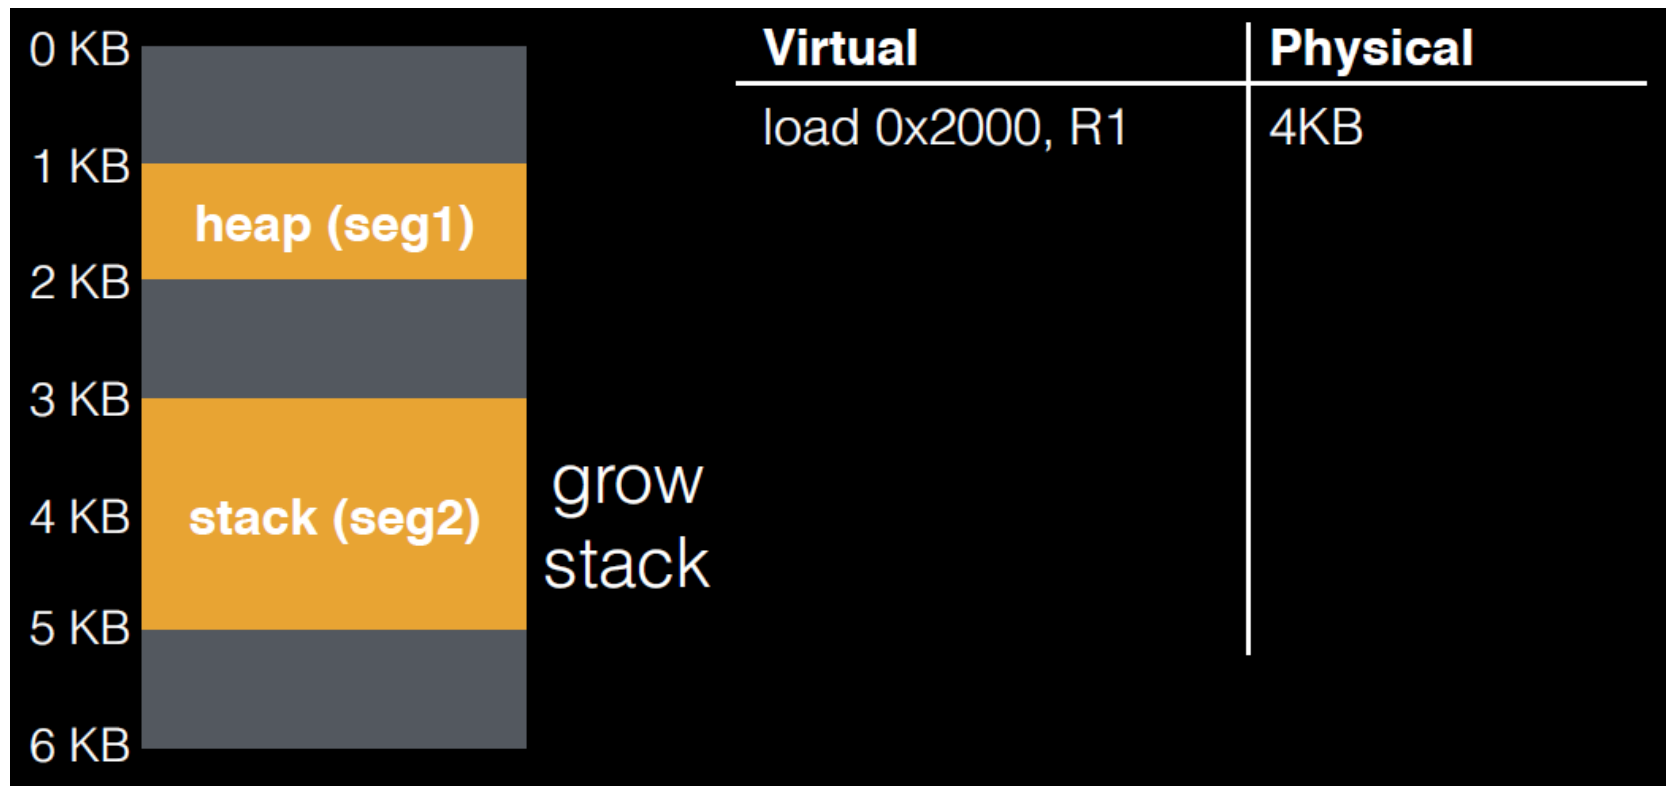
0 KB		Virtual	Physical
1 KB	heap (seg1)	load 0x2010, R1	4KB + 16
2 KB		load 0x1010, R1	1KB + 16
3 KB		load 0x1100, R1	1KB + 256
4 KB	stack (seg2)	load 0x1400, R1	interrupt OS!
5 KB			
6 KB			

# Stack Growth Problem

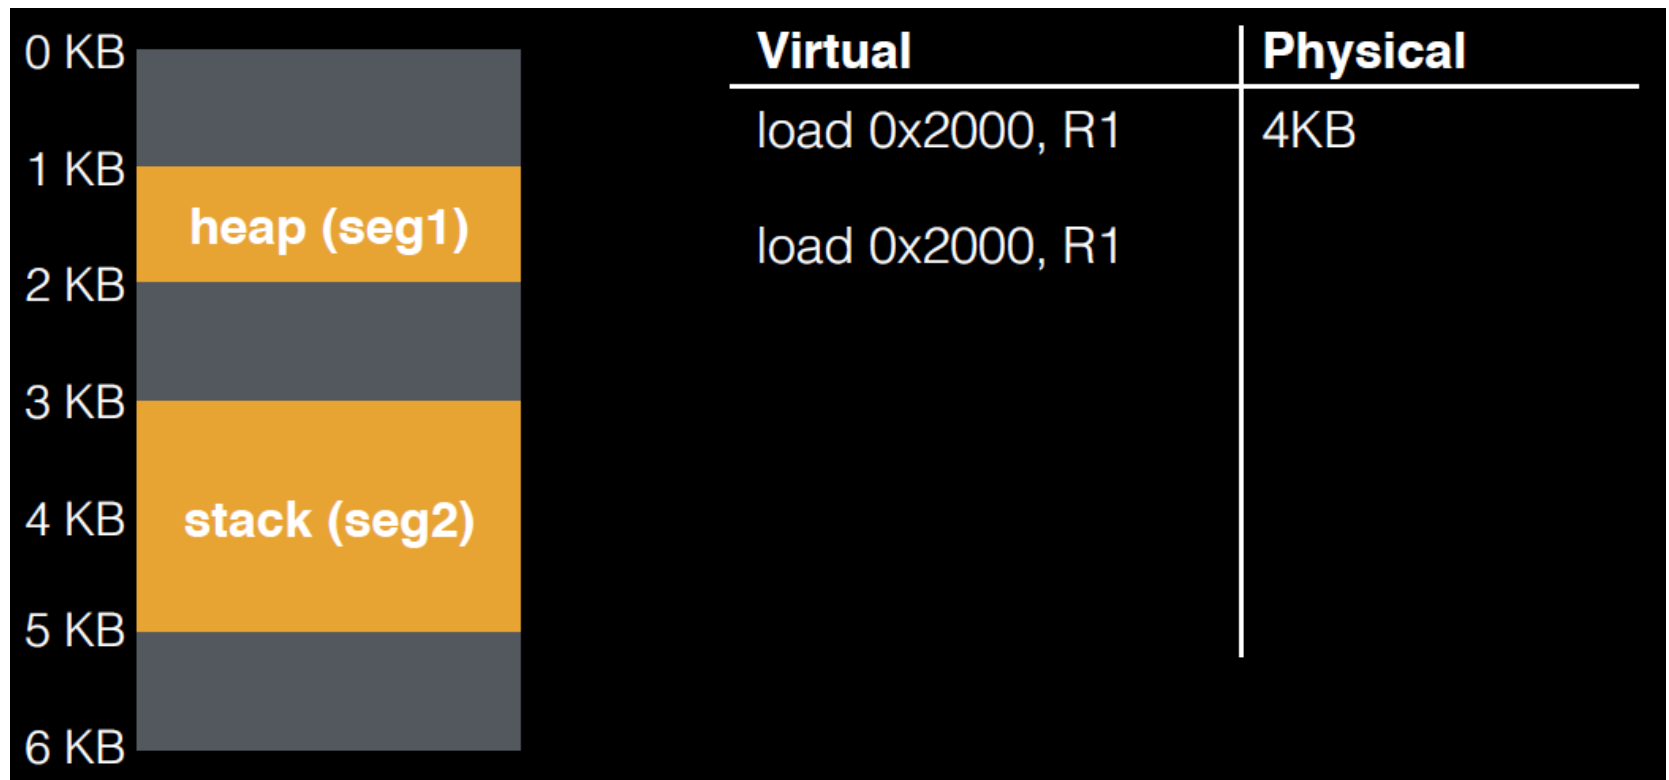
Example...

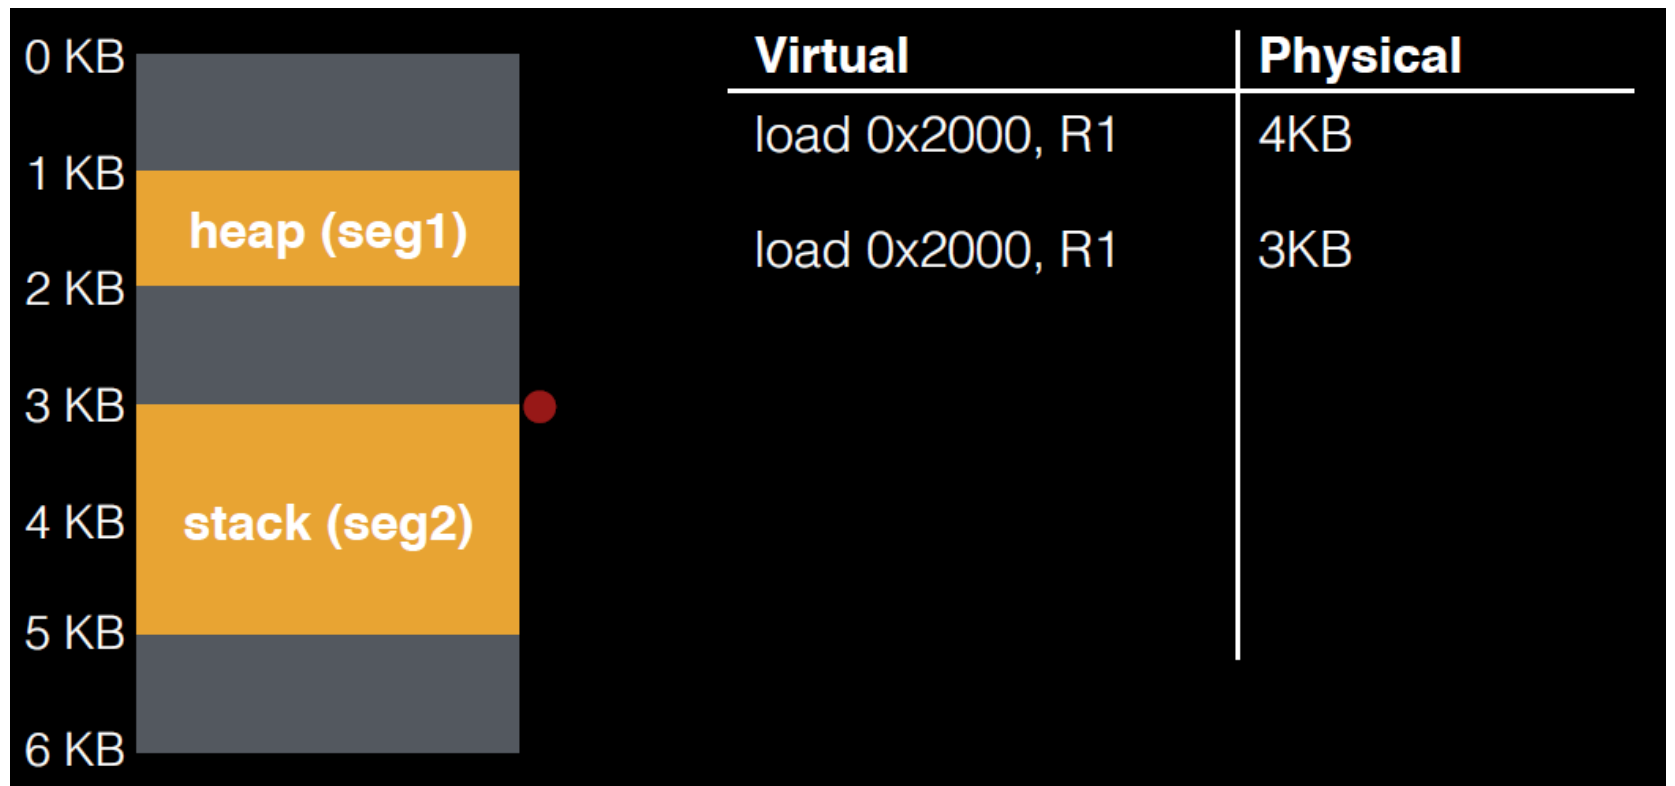












# Stack Growth Problem

Example...

Problem:  $\text{phys} = \text{virt\_offset} + \text{base\_reg}$

phys is anchored to `base_reg`, which moves

# Stack Growth Problem

Example...

Problem:  $\text{phys} = \text{virt\_offset} + \text{base\_reg}$

phys is anchored to `base_reg`, which moves

Solution: anchor heap segment to `bounds_reg`:

$\text{phys} = \text{bounds\_reg} - (\text{max\_offset} - \text{virt\_offset})$

# Stack Growth Problem

Example...

Problem:  $\text{phys} = \text{virt\_offset} + \text{base\_reg}$

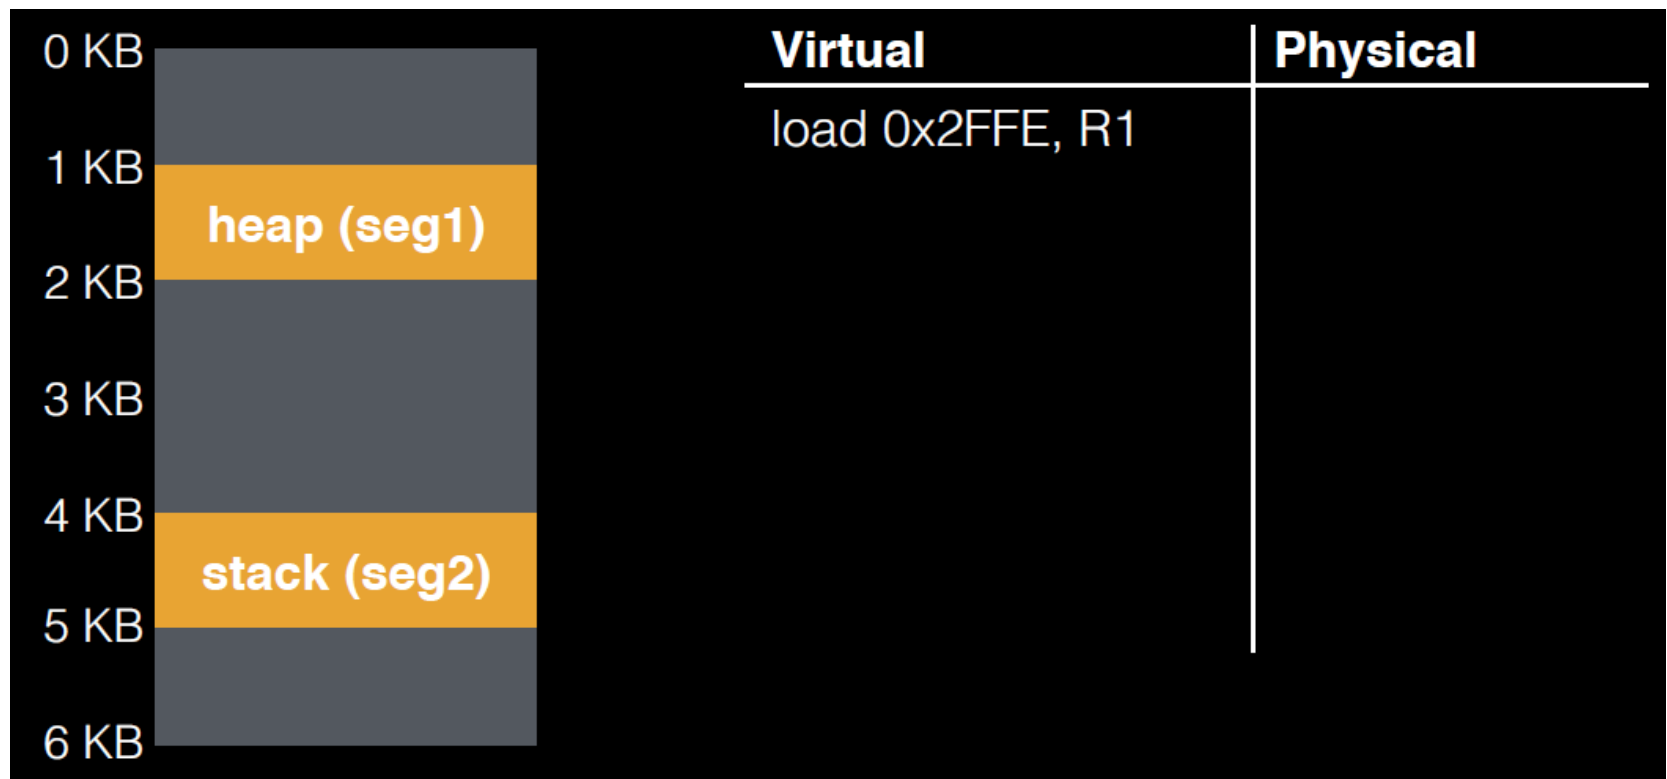
phys is anchored to `base_reg`, which moves

Solution: anchor heap segment to `bounds_reg`:

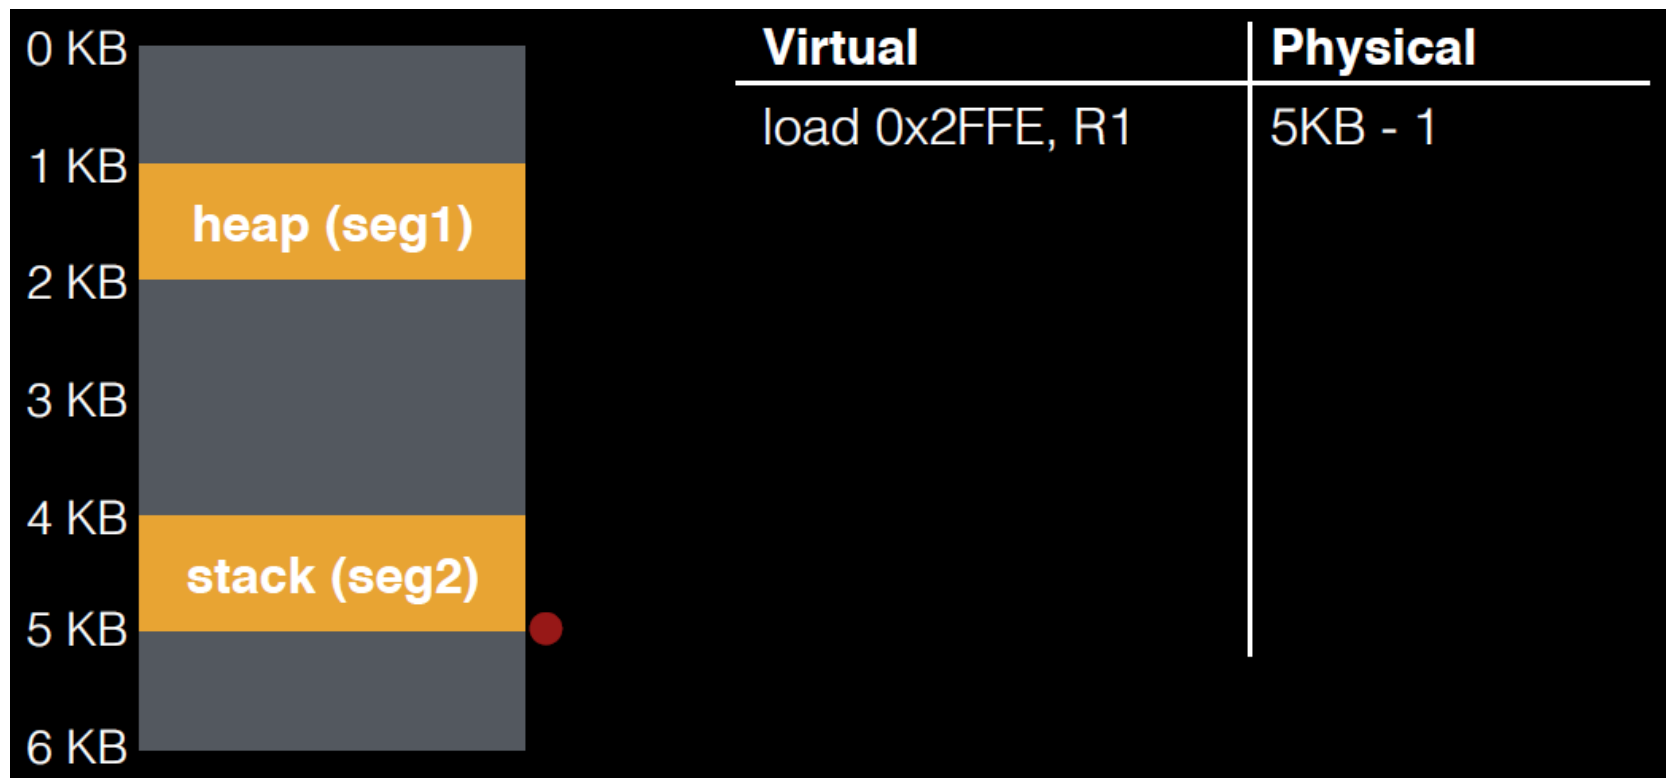
$\text{phys} = \text{bounds\_reg} - (\text{max\_offset} - \text{virt\_offset})$

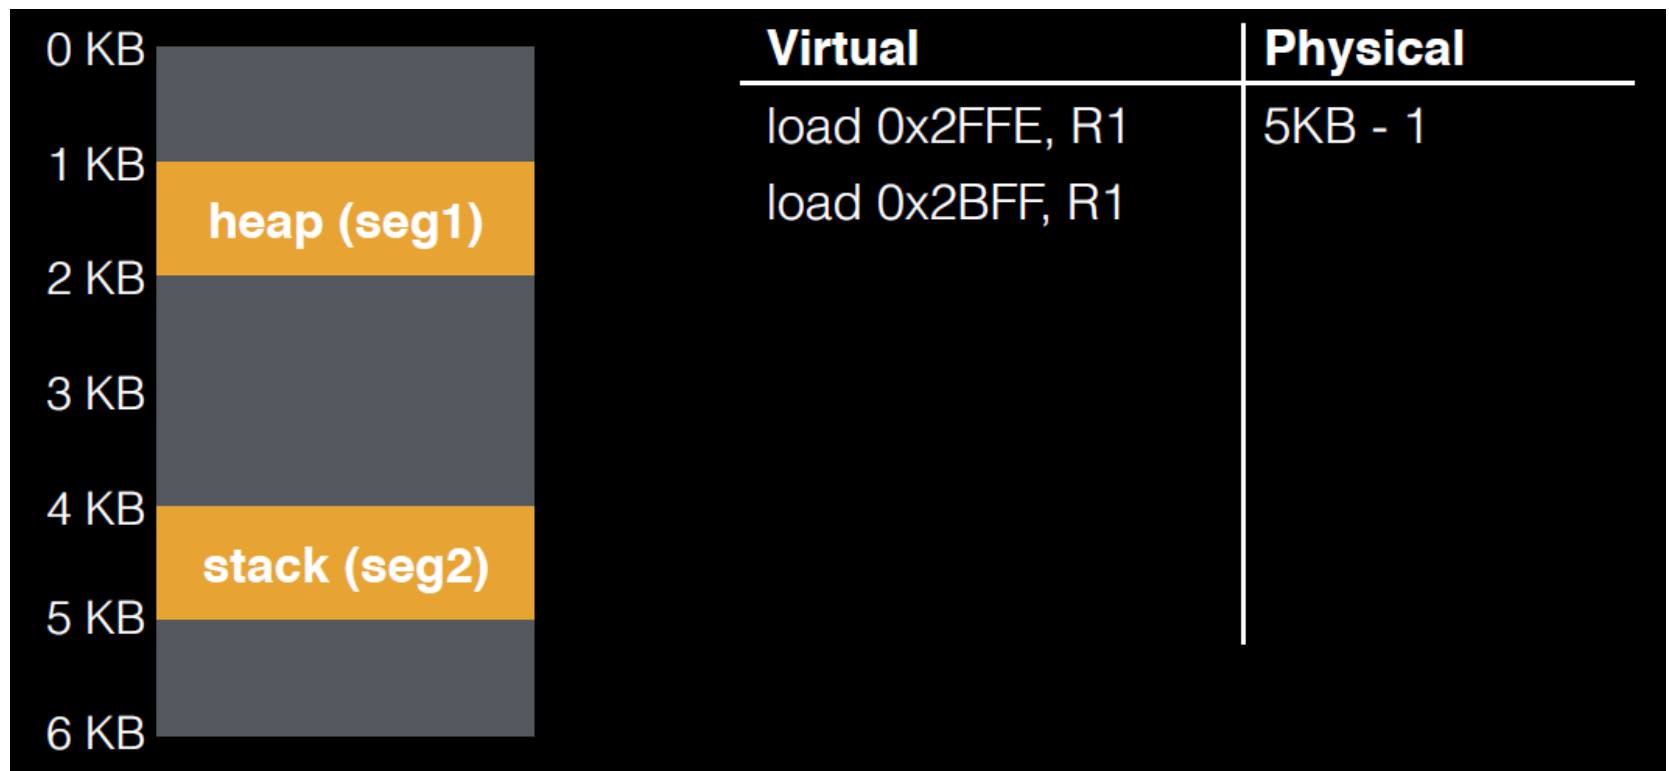
Example (with `max_offset = FFF`)...

stack's max\_offset = FFF

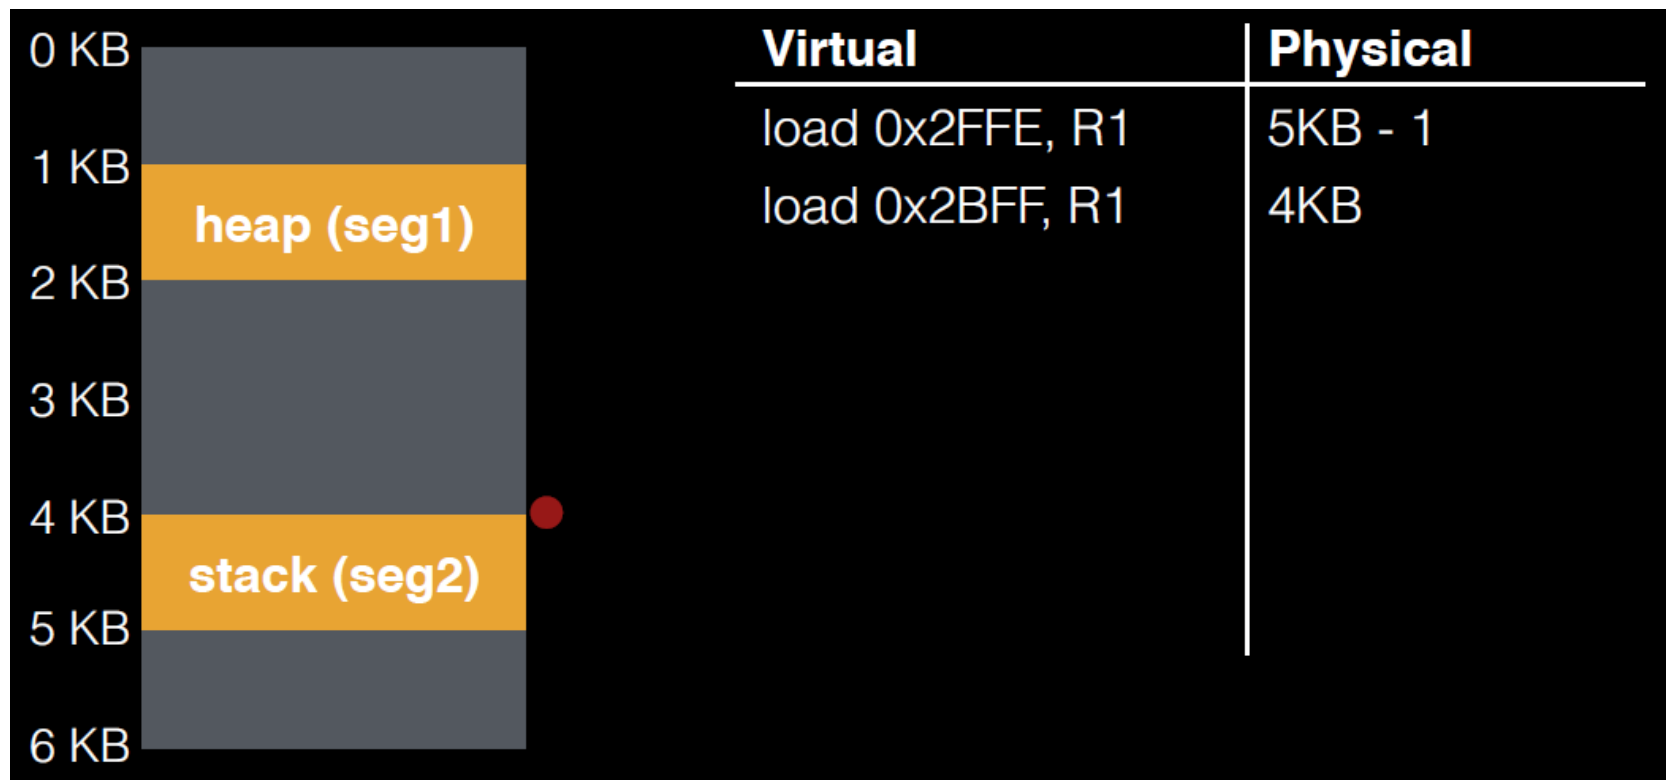


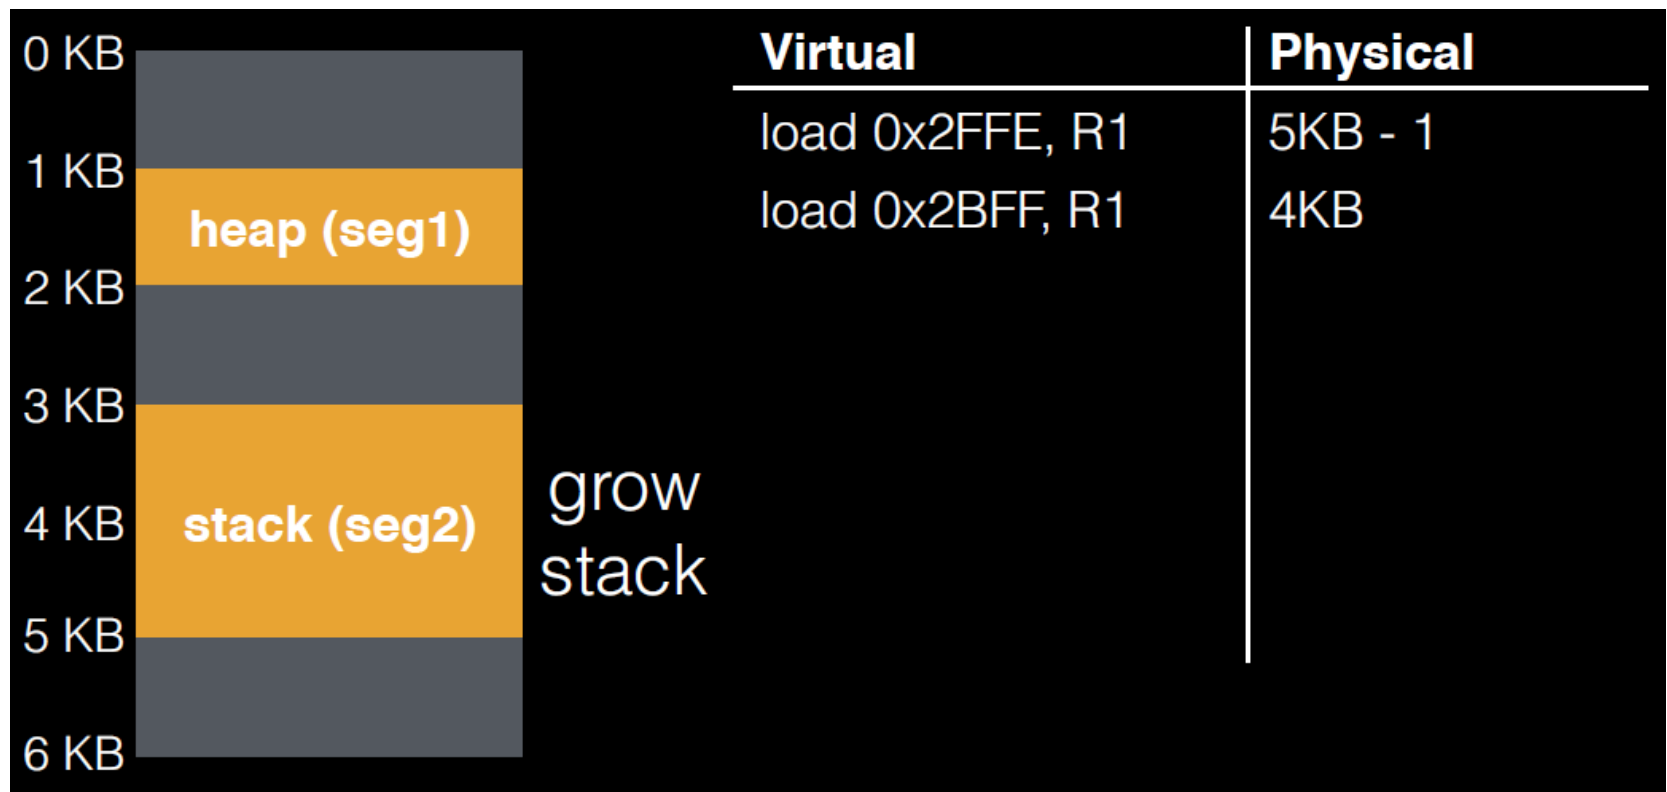
stack's max\_offset = FFF



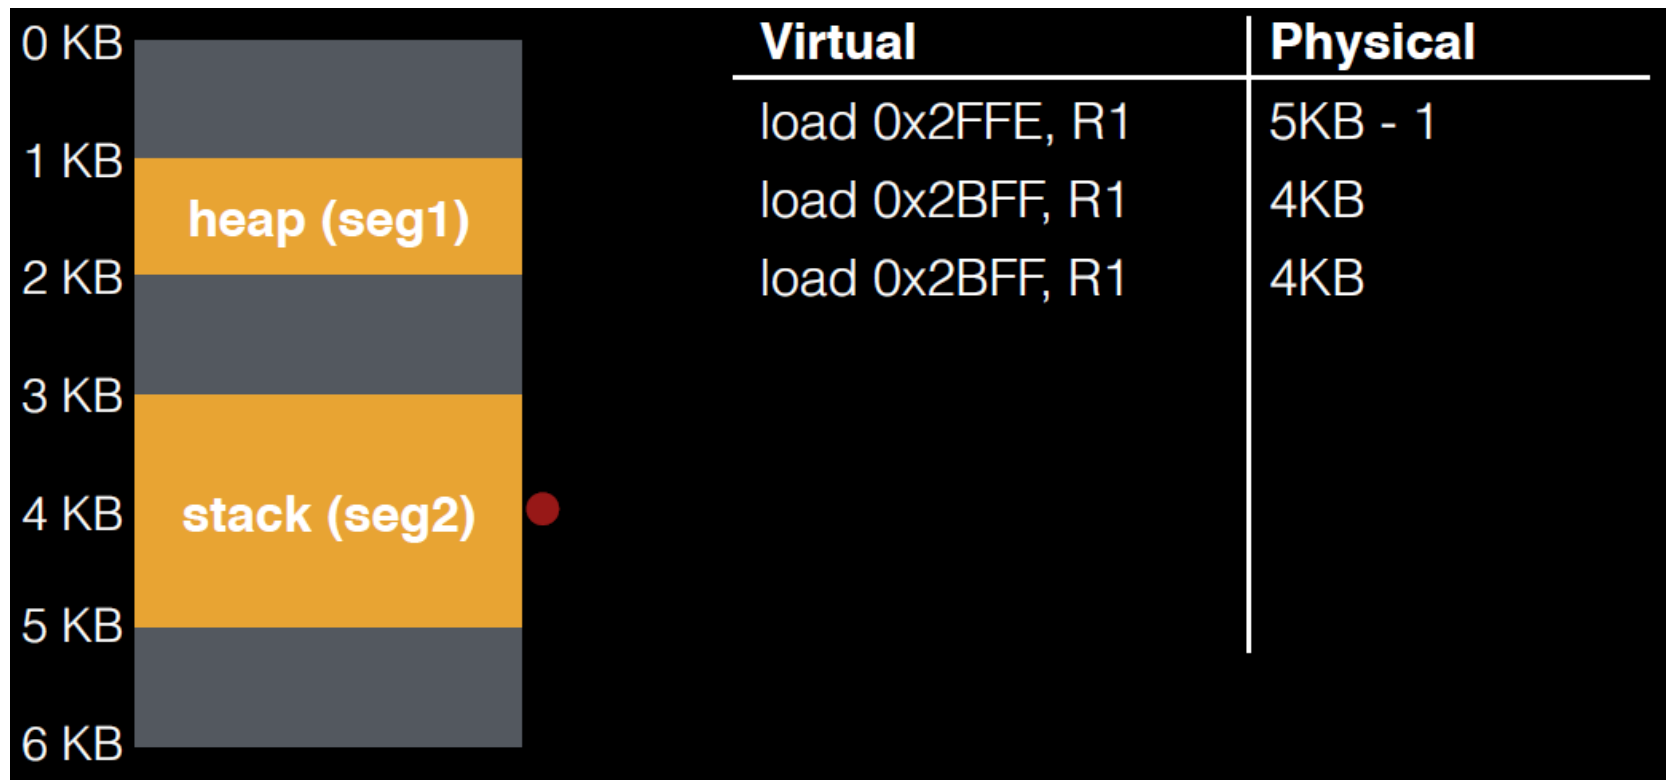








0 KB		<b>Virtual</b>	<b>Physical</b>
1 KB	heap (seg1)	load 0x2FFE, R1	5KB - 1
2 KB		load 0x2BFF, R1	4KB
3 KB		load 0x2BFF, R1	
4 KB	stack (seg2)		
5 KB			
6 KB			



# Translation Summary

Heap:  $\text{phys} = \text{base\_reg} + \text{virt\_offset}$

Stack:  $\text{phys} = \text{bounds\_reg} - (\text{max\_offset} - \text{virt\_offset})$

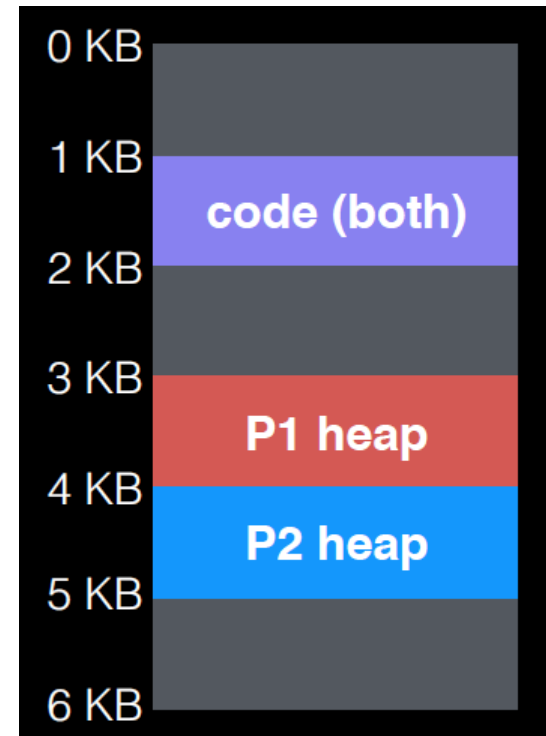
Anchors:

- for heap, anchor **smallest address to base register**
- for stack, anchor **biggest address to bounds register**

# Code Sharing

Idea: make base/bounds for  
the code of several processes  
point to the same physical mem

Careful: need extra protection!



# Segmentation Pros/Cons

## Pros?

- supports sparse address space
- code sharing
- fine grained protection

## Cons?

- external fragmentation

# Conclusion

HW+OS work together to virtualize memory

- Give illusion of private address space to each process

Add MMU registers for base+bounds so translation is fast

- OS not involved with every address translation, only on context switch or errors

Dynamic relocation with segments is good building block

- Next lecture: Solve fragmentation with paging