

# OSTEP

## Concurrency: Threads

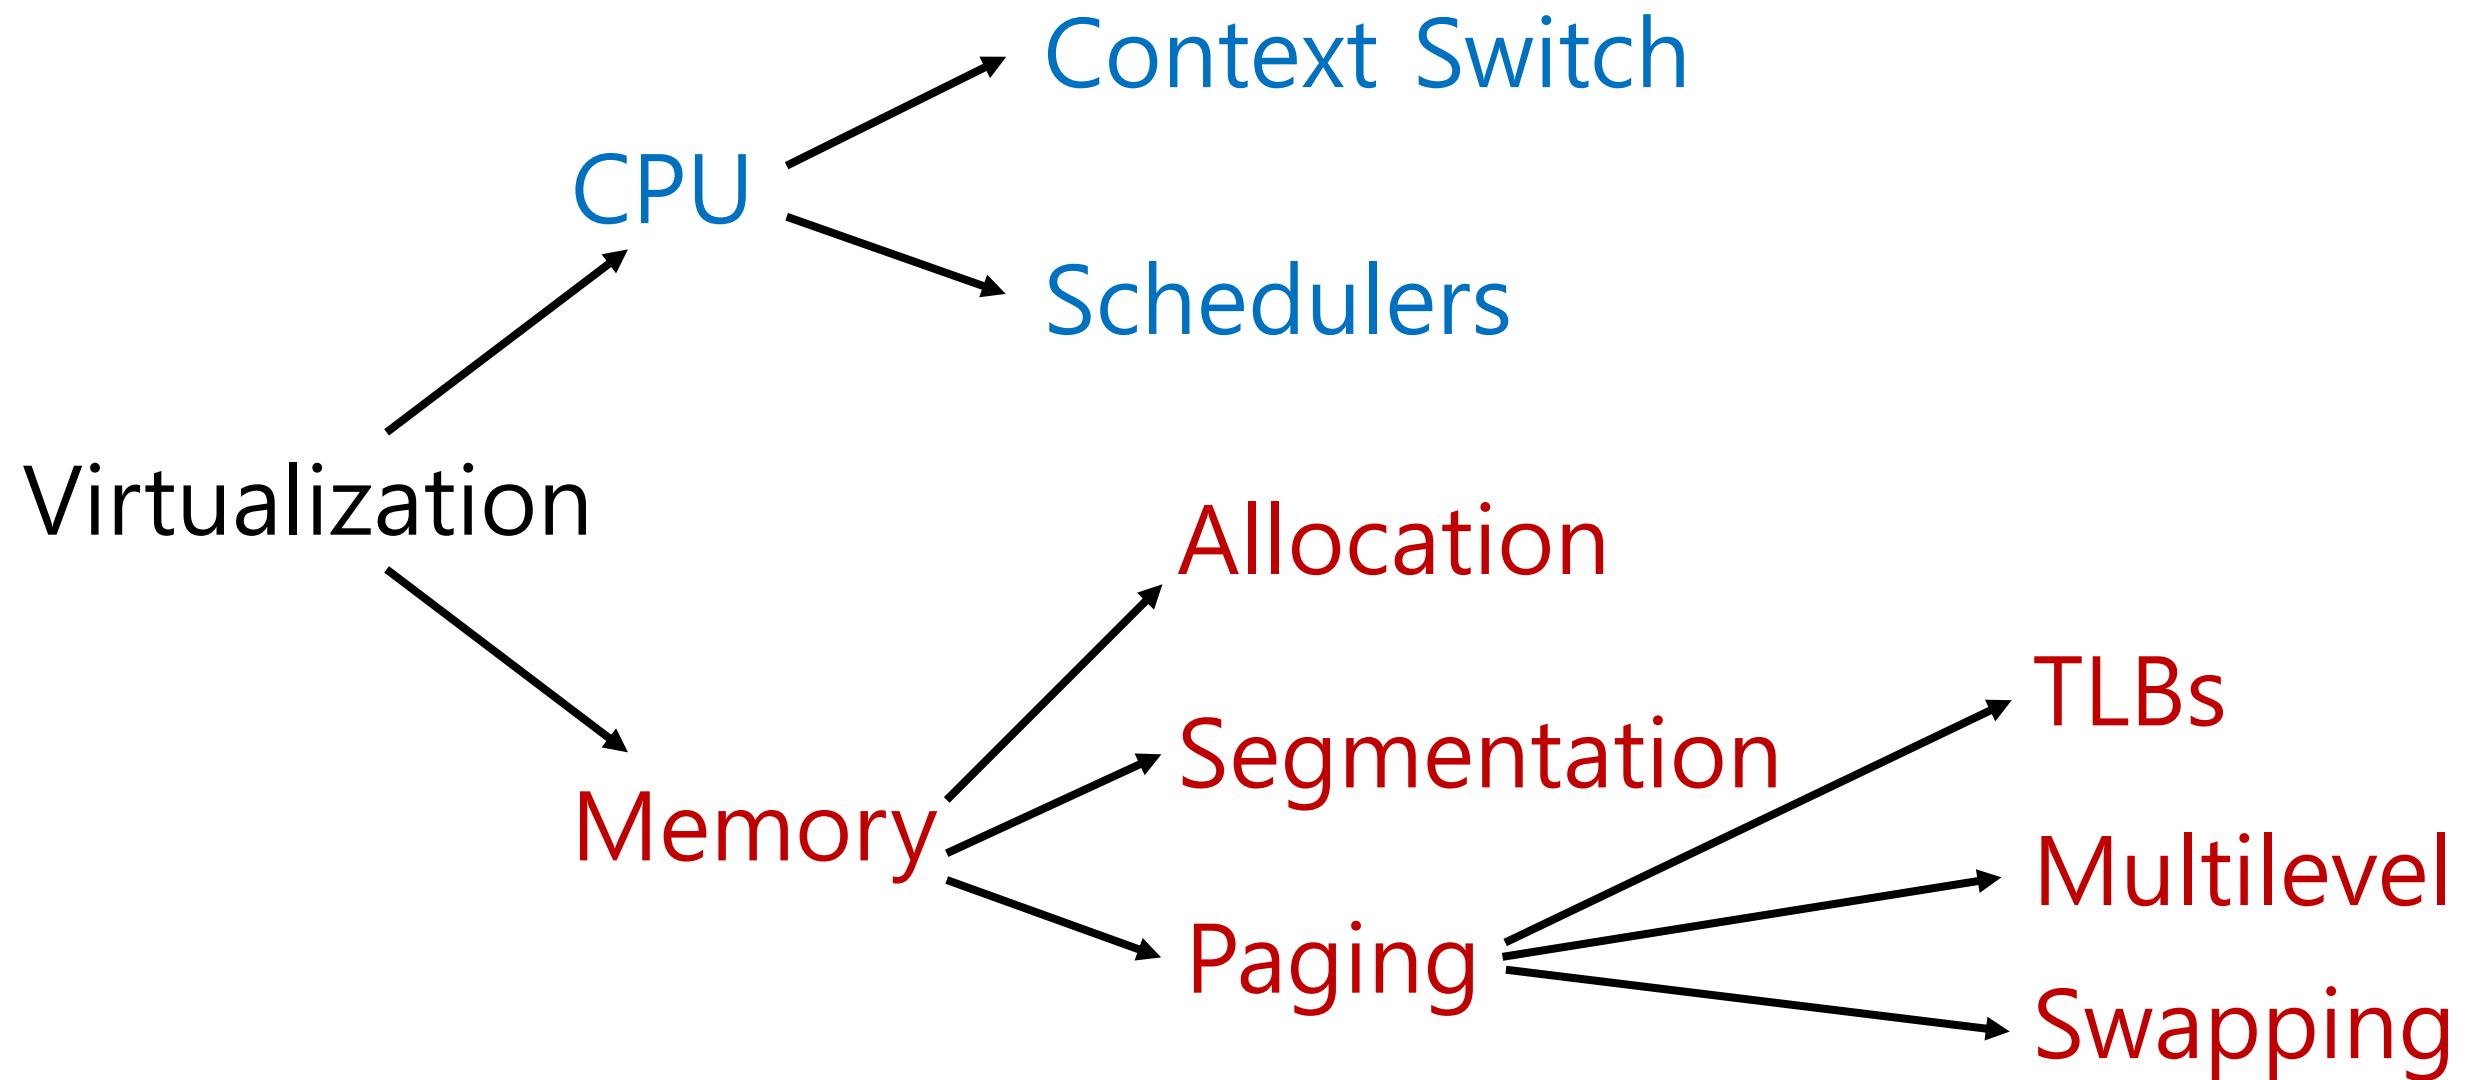
### **Questions answered in this lecture:**

Why is concurrency useful?


What is a thread and how does it differ from processes?

What can go wrong if scheduling of critical sections is not atomic?

# Review: Easy Piece 1



# Motivation for Concurrency



# Motivation

CPU Trend: Same speed, but multiple cores

Goal: Write applications that fully utilize many cores

**Option 1:** Build apps from many communicating **processes**

- Example: Chrome (process per tab)
- Communicate via pipe() or similar

Pros?

- Don't need new abstractions; good for security

Cons?

- Cumbersome programming
- High communication overheads
- Expensive context switching (why expensive?)

# CONCURRENCY: Option 2

**Thread** is a new abstraction for a single running process:

Threads are **like processes**, except:  
multiple threads of same process **share an address space**  
(e.g., using same PT).

Divide large task across several cooperative threads

Communicate through shared address space

# Context switch between threads

Each thread has its own program counter and set of registers.


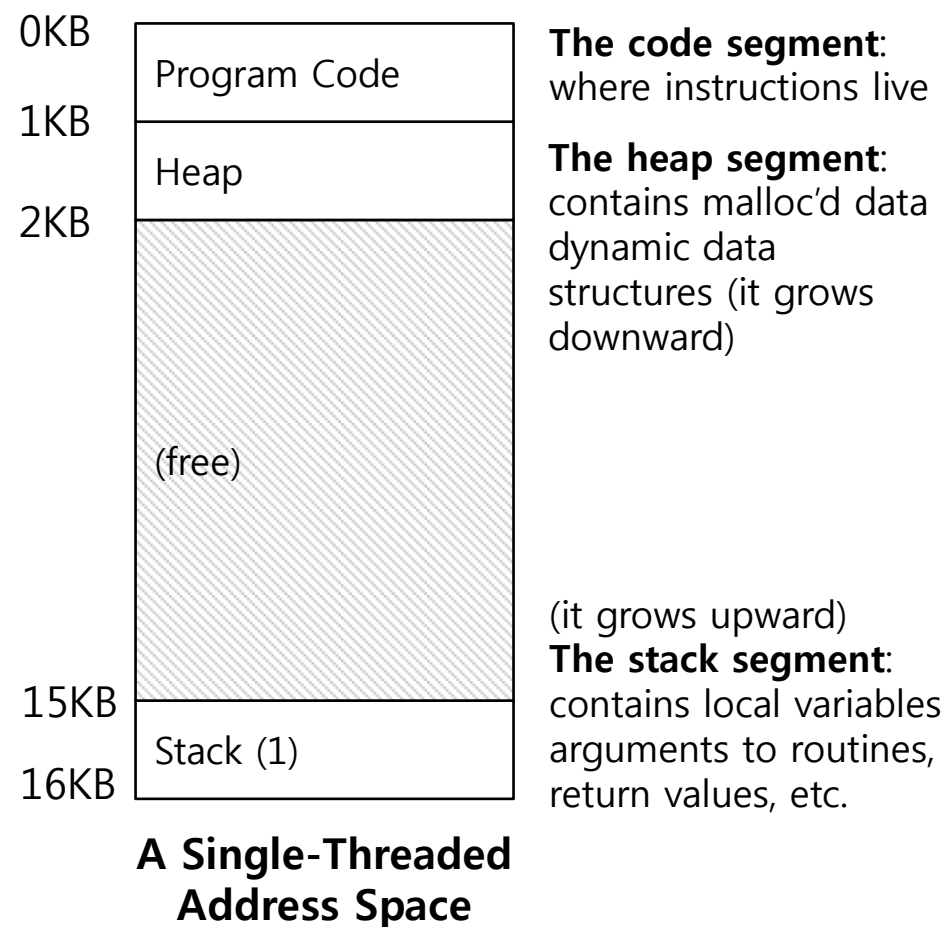
- One or more **thread control blocks(TCBs)** are needed to store the state of each thread.

When switching from running one (T1) to running the other (T2),

- The register state of T1 be saved.
- The register state of T2 restored.
- The **address space remains** the same.

# The stack of the relevant thread

There will be **one stack per thread**.



# Race condition

## Example with two threads

- counter = counter + 1 (default is 50)
- We expect the result is 52. However,

OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	before critical section		100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
interrupt	save T1's state				
	restore T2's state		100	0	50
		mov 0x8049a1c, %eax	105	50	50
		add \$0x1, %eax	108	51	50
		mov %eax, 0x8049a1c	113	51	51
interrupt	save T2's state				
	restore T1's state		108	51	50
	mov %eax, 0x8049a1c		113	51	51



# Critical section

A piece of code that **accesses a shared variable** and must not be concurrently executed by more than one thread.

- Multiple threads executing critical section can result in a race condition.
- Need to support **atomicity** for critical sections (**mutual exclusion**)

# Locks

Ensure that any such critical section executes as if it were a single atomic instruction (**execute a series of instructions atomically**).

```
1  lock_t mutex;  
2  . . .  
3  lock(&mutex);  
4  balance = balance + 1;  
5  unlock(&mutex);
```

→ Critical section

# Common Programming Models

Multi-threaded programs tend to be structured as:

- **Producer/consumer**


Multiple producer threads create data (or work) that is handled by one of the multiple consumer threads

- **Pipeline**


Task is divided into series of subtasks, each of which is handled in series by a different thread

- **Defer work with background thread**

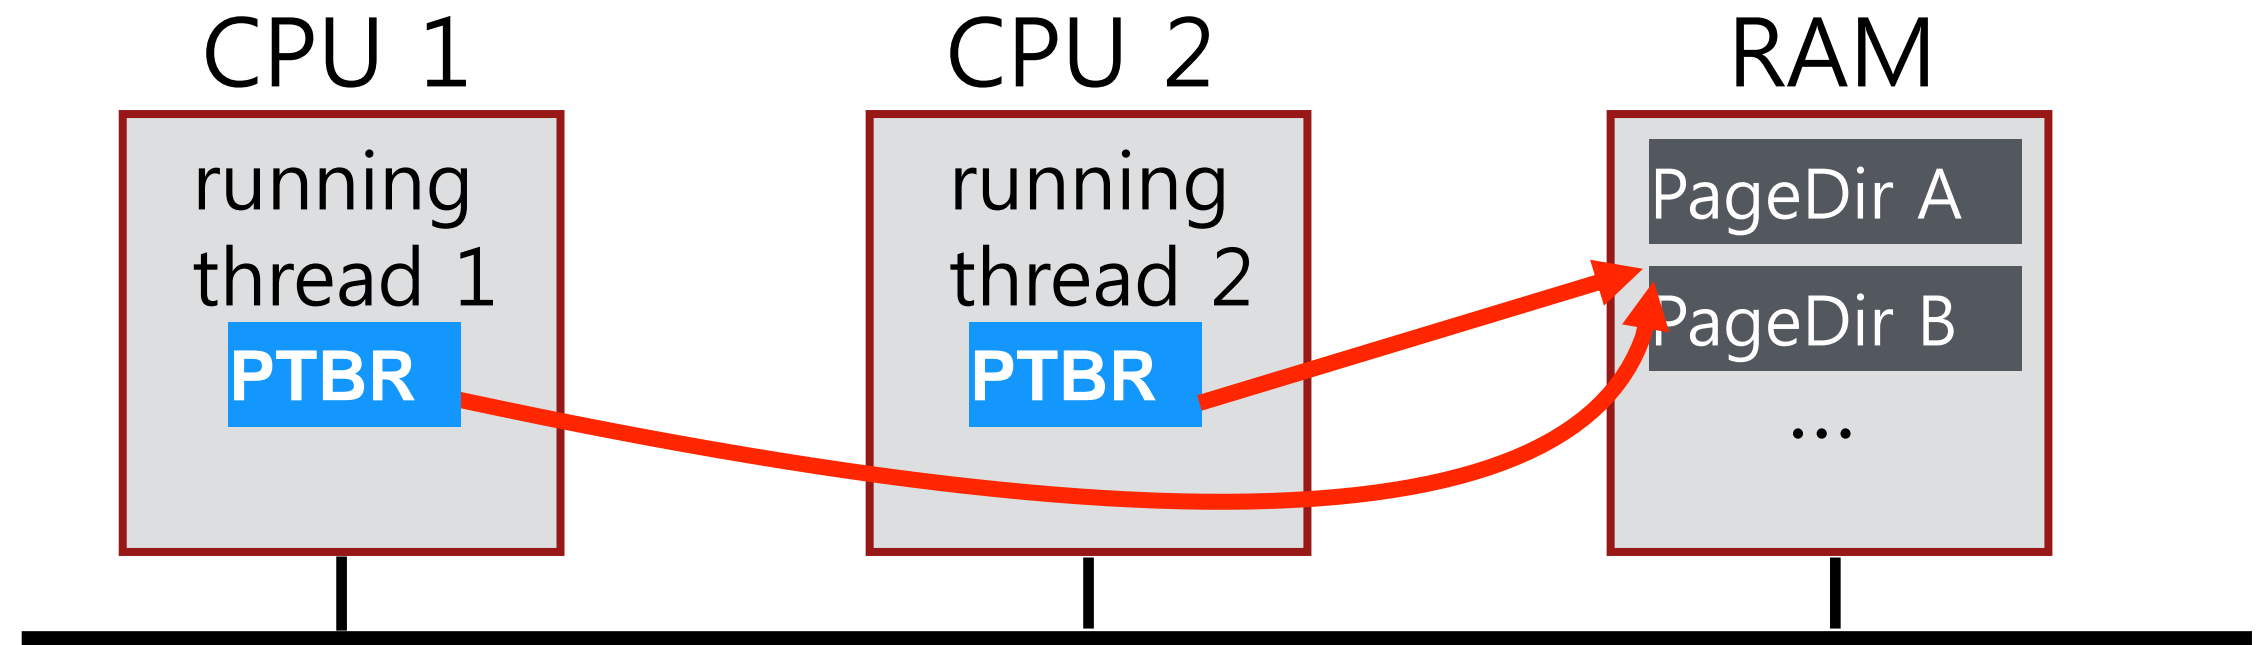
One thread performs non-critical work in the background (when CPU idle)




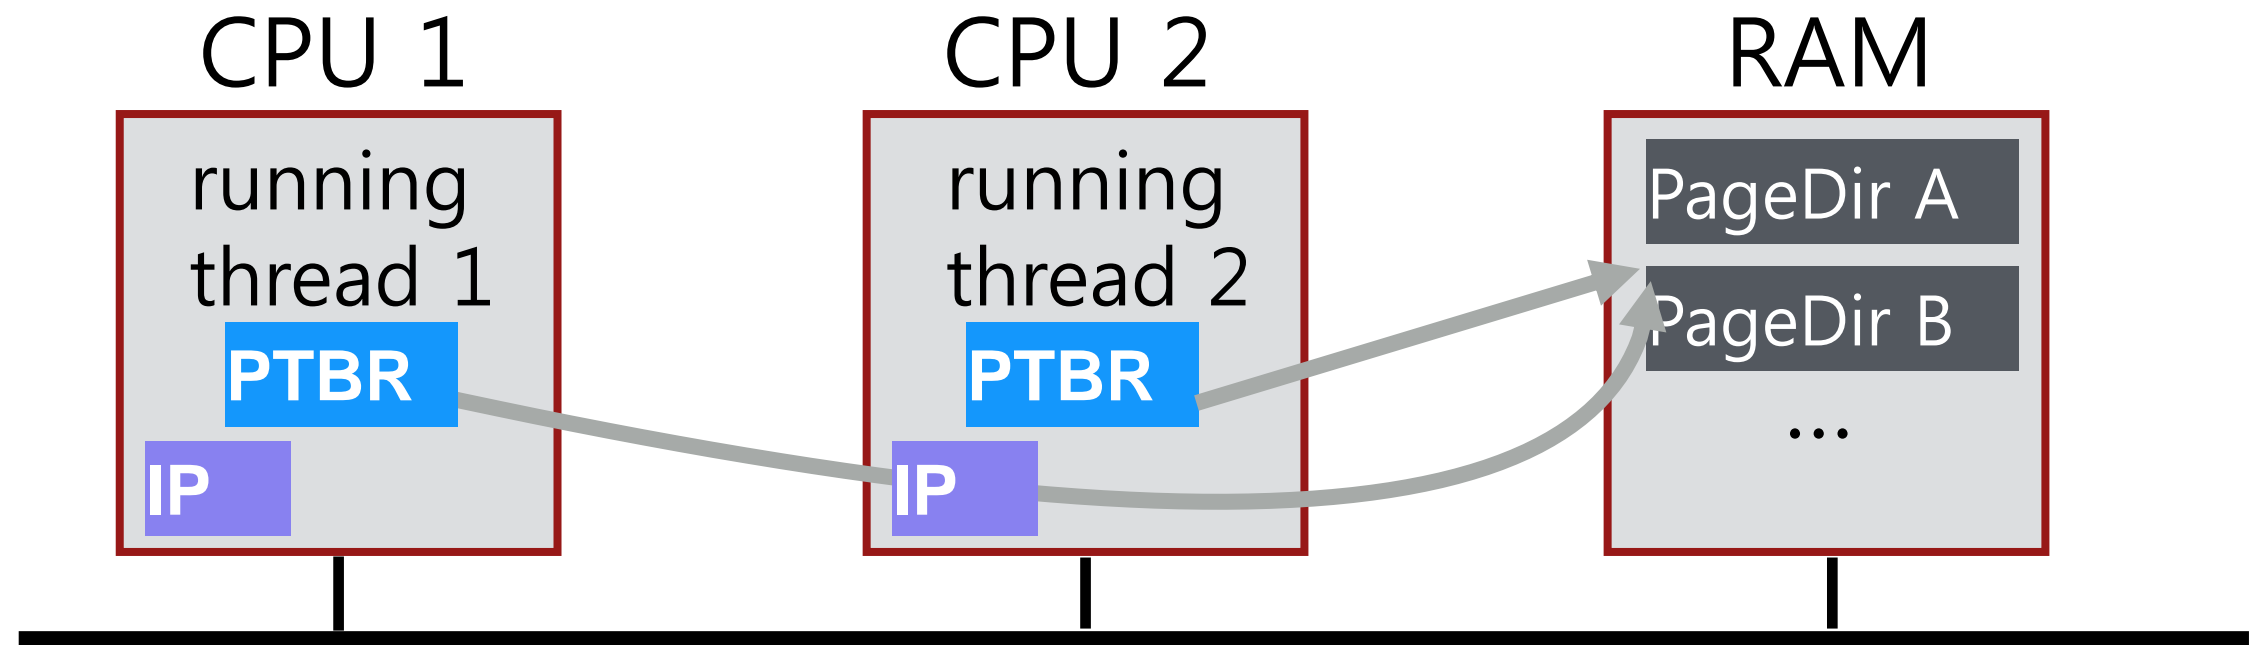
What state do threads share?



What threads share page directories?

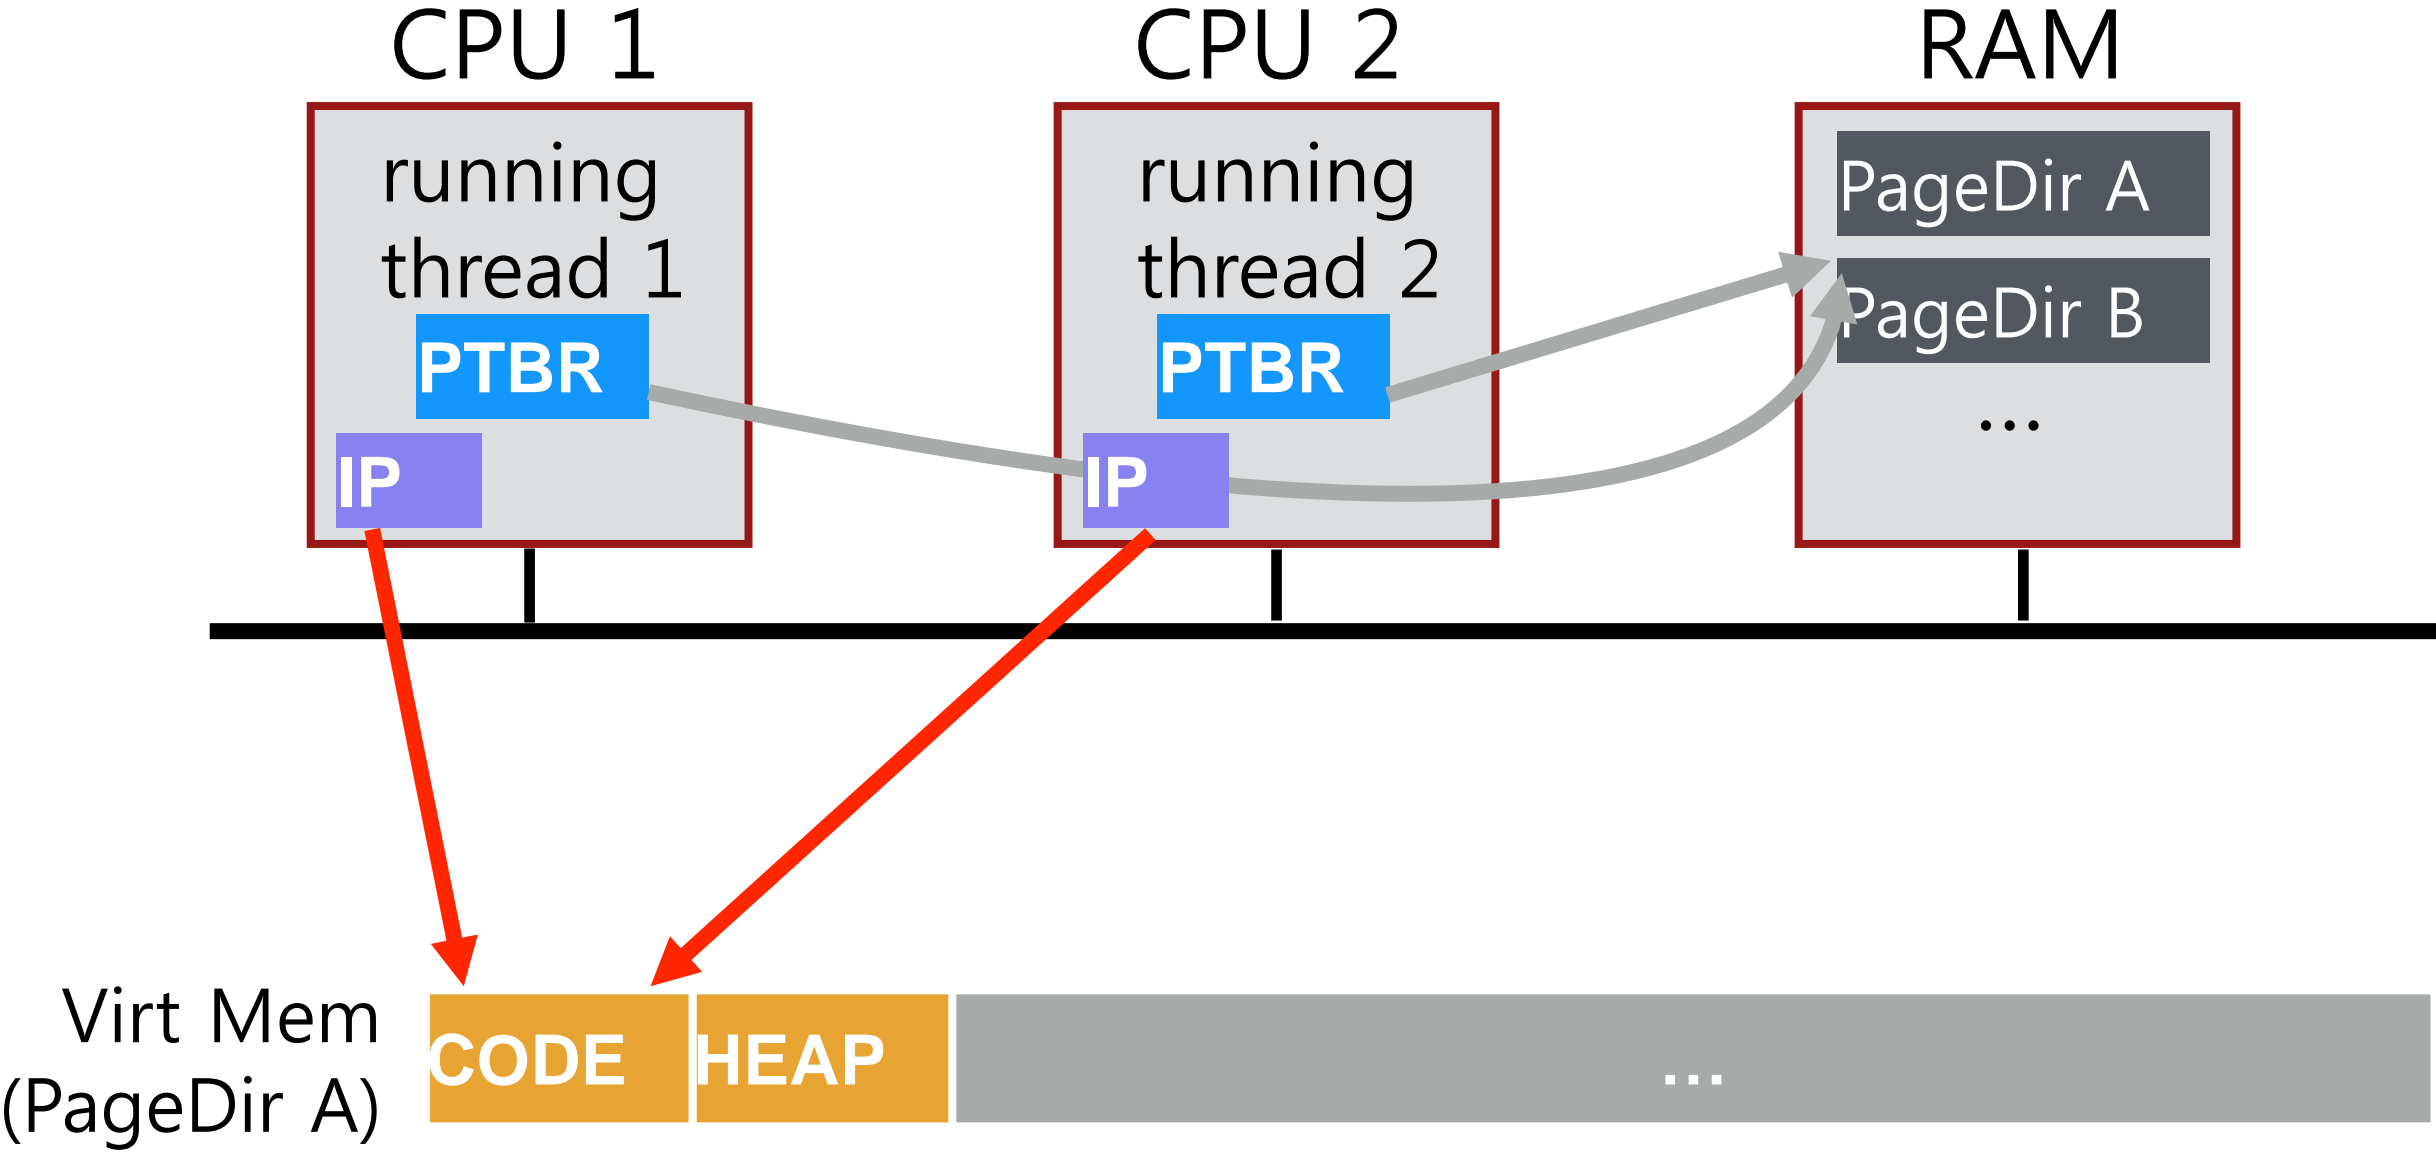


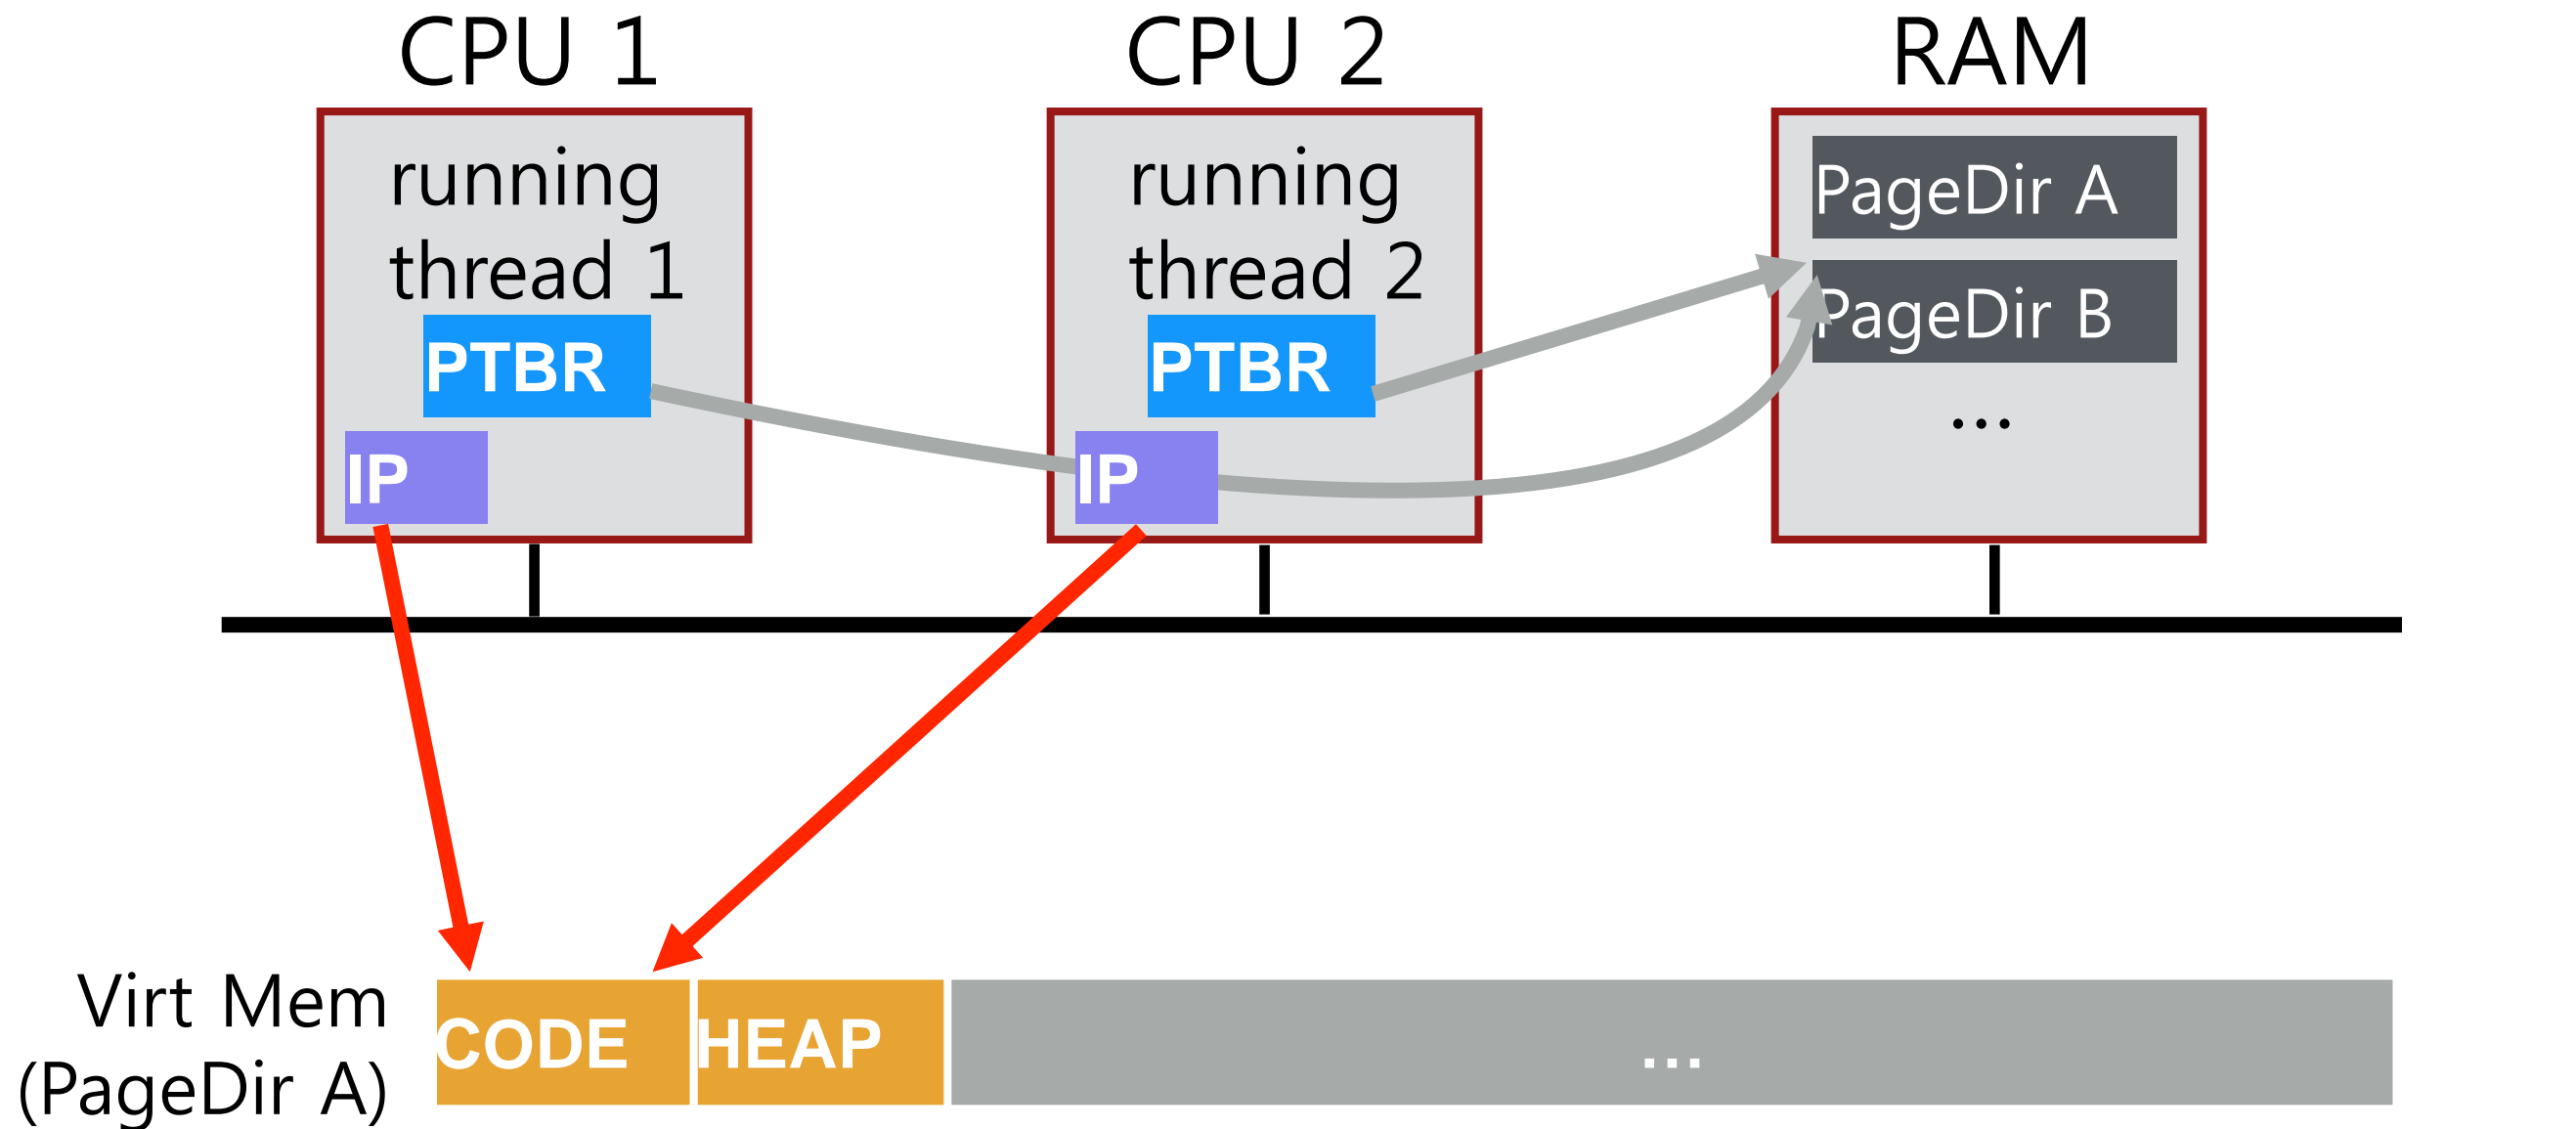




Do threads share Instruction Pointer?

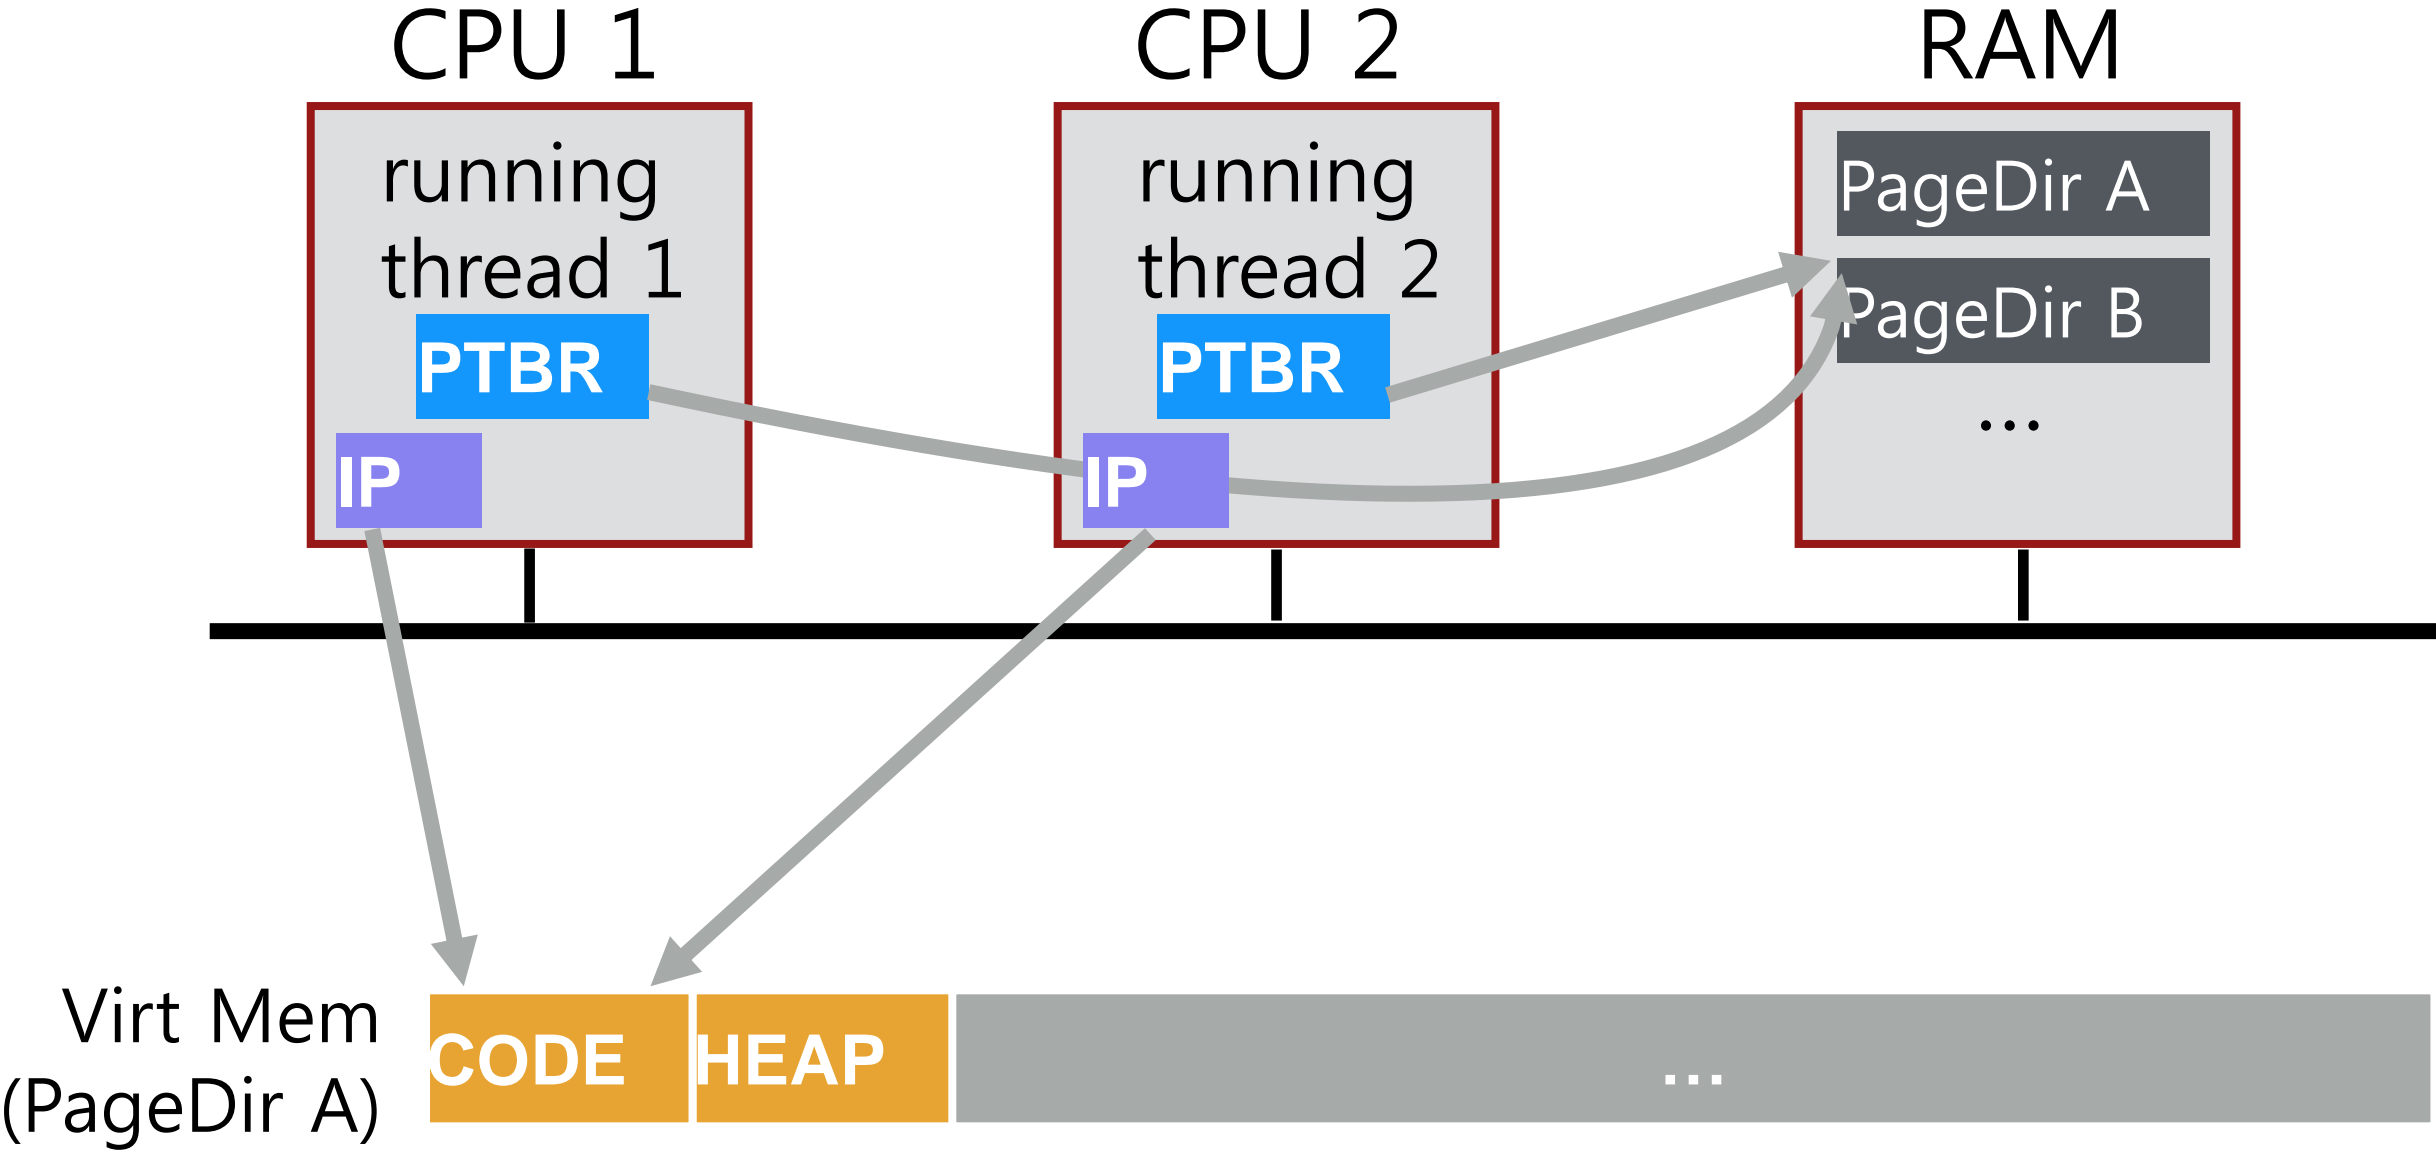


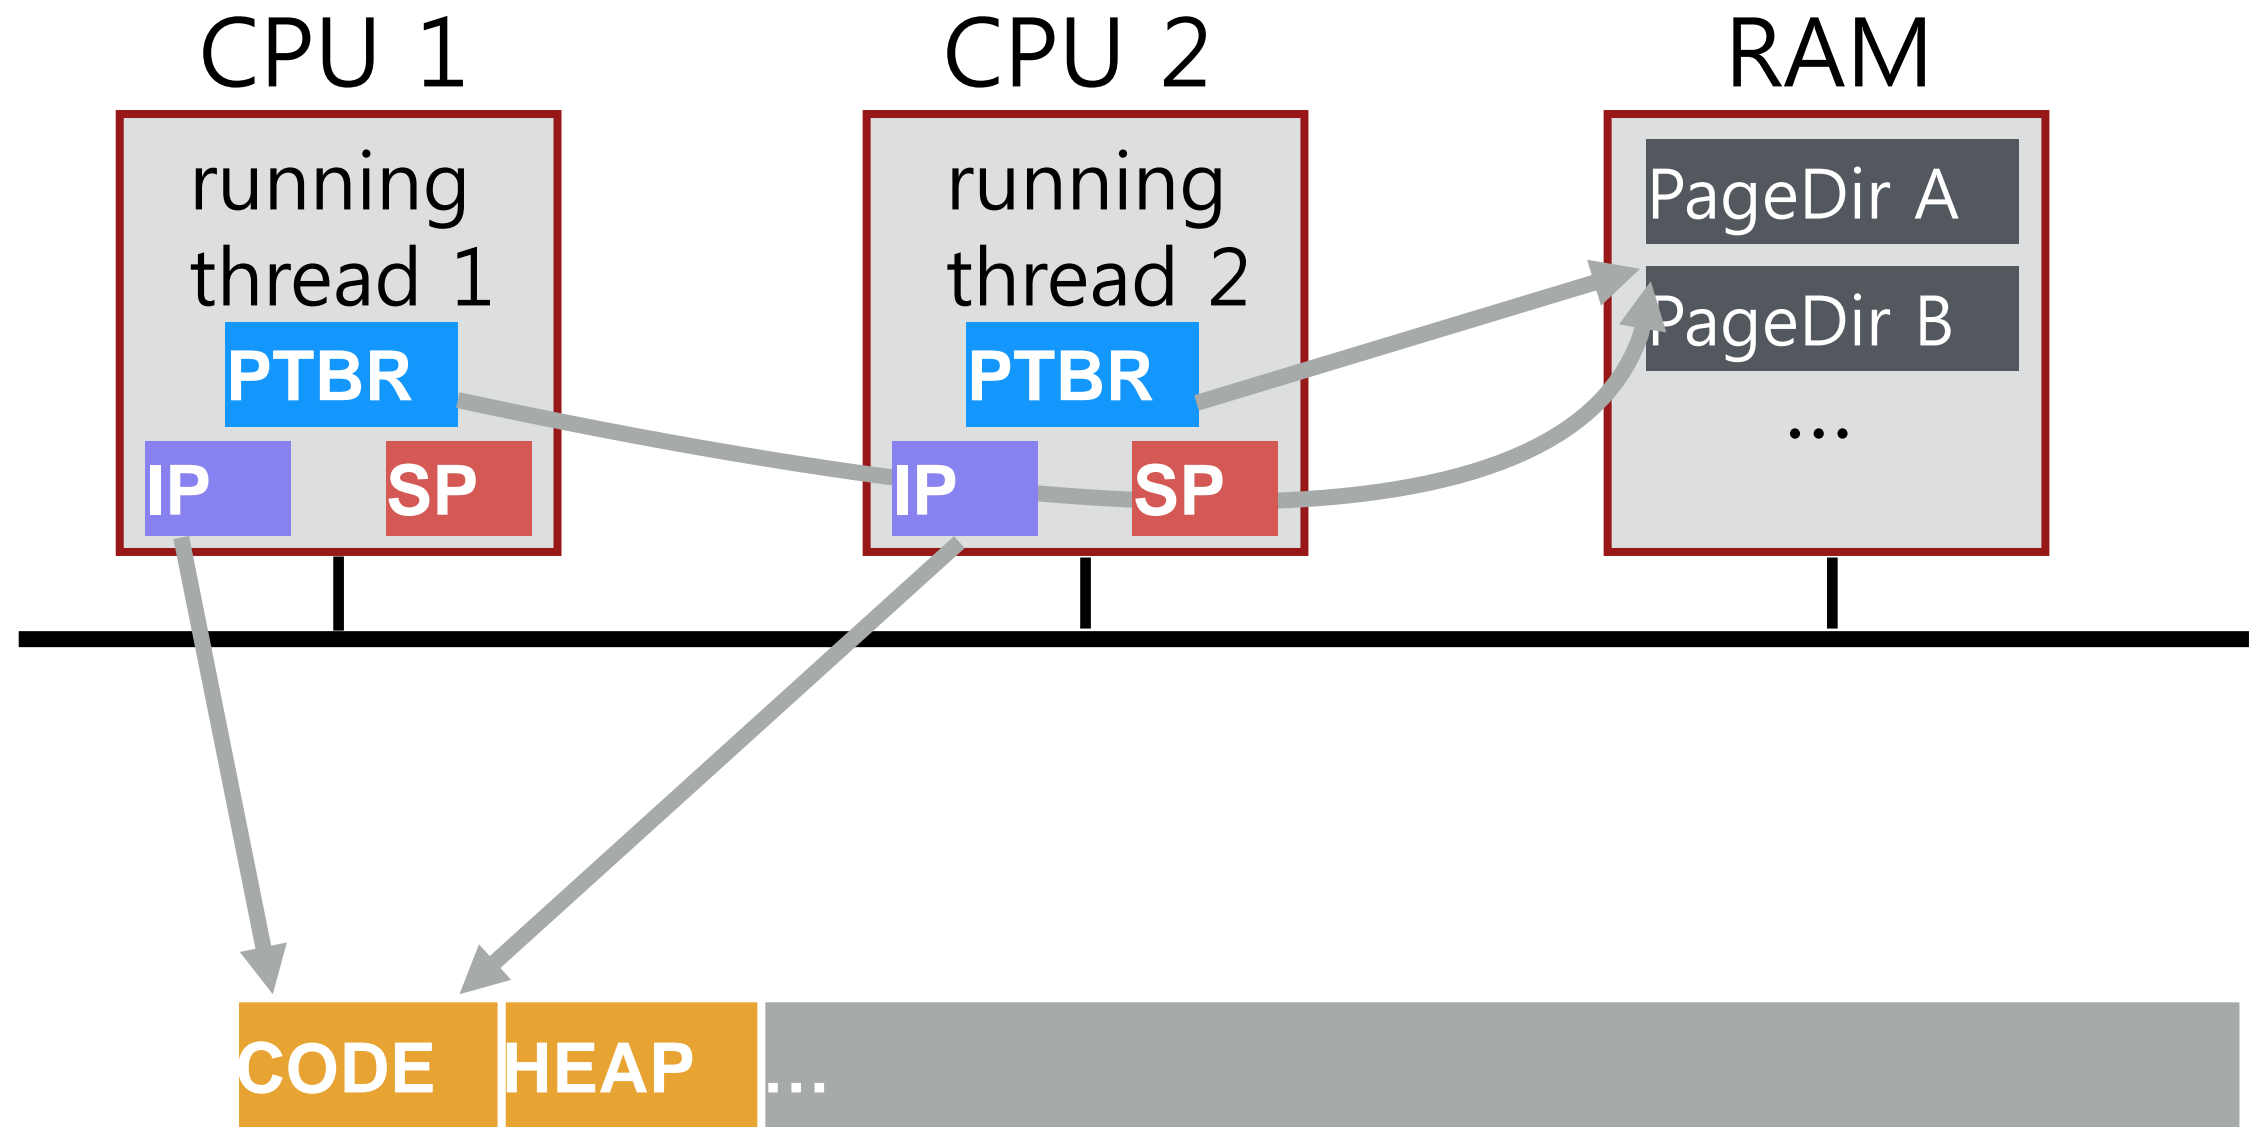




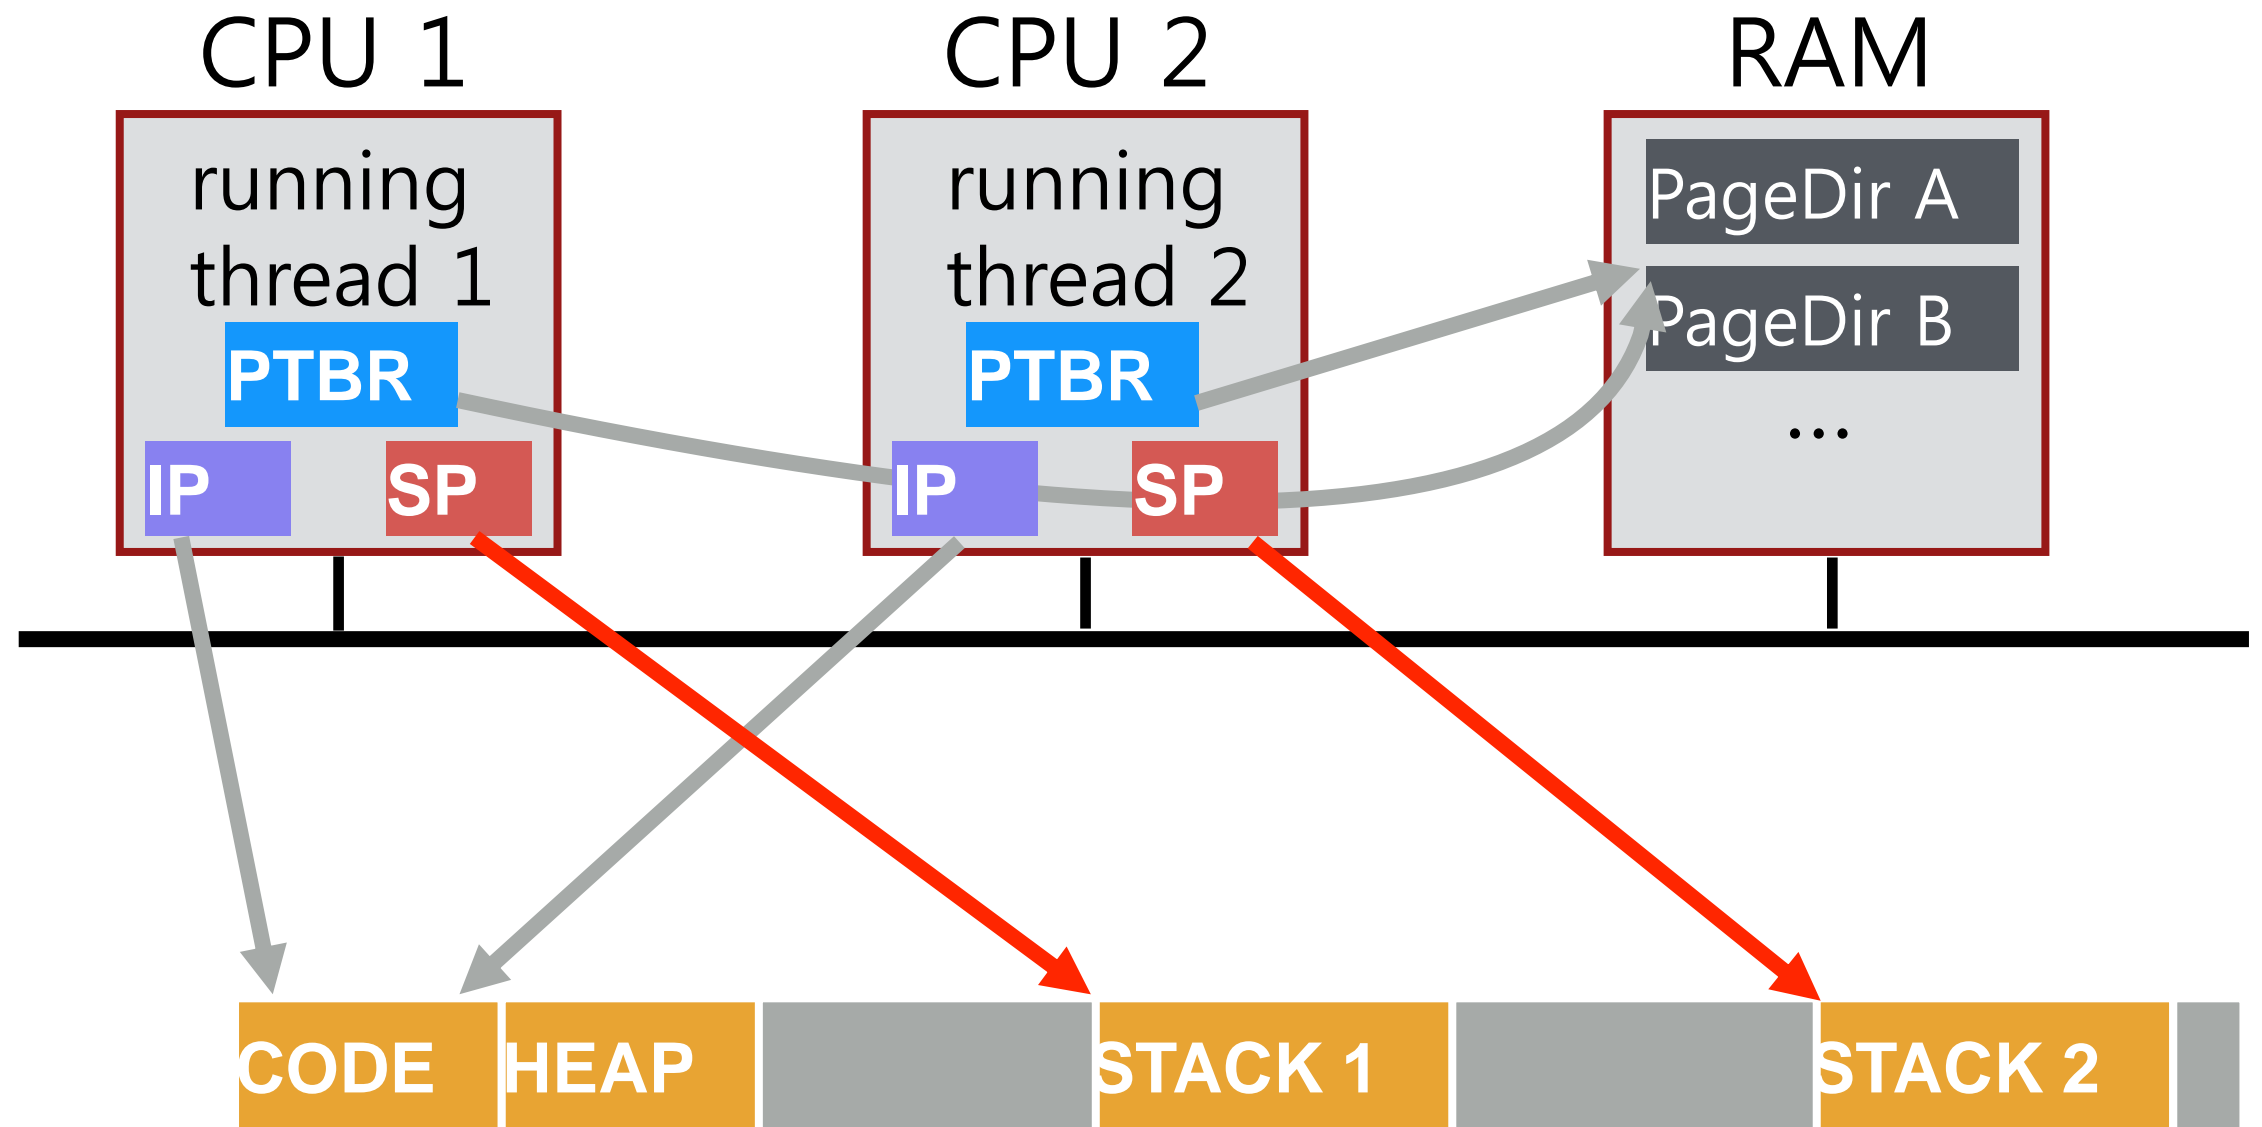
Share code, but each thread may be executing  
different code at the same time

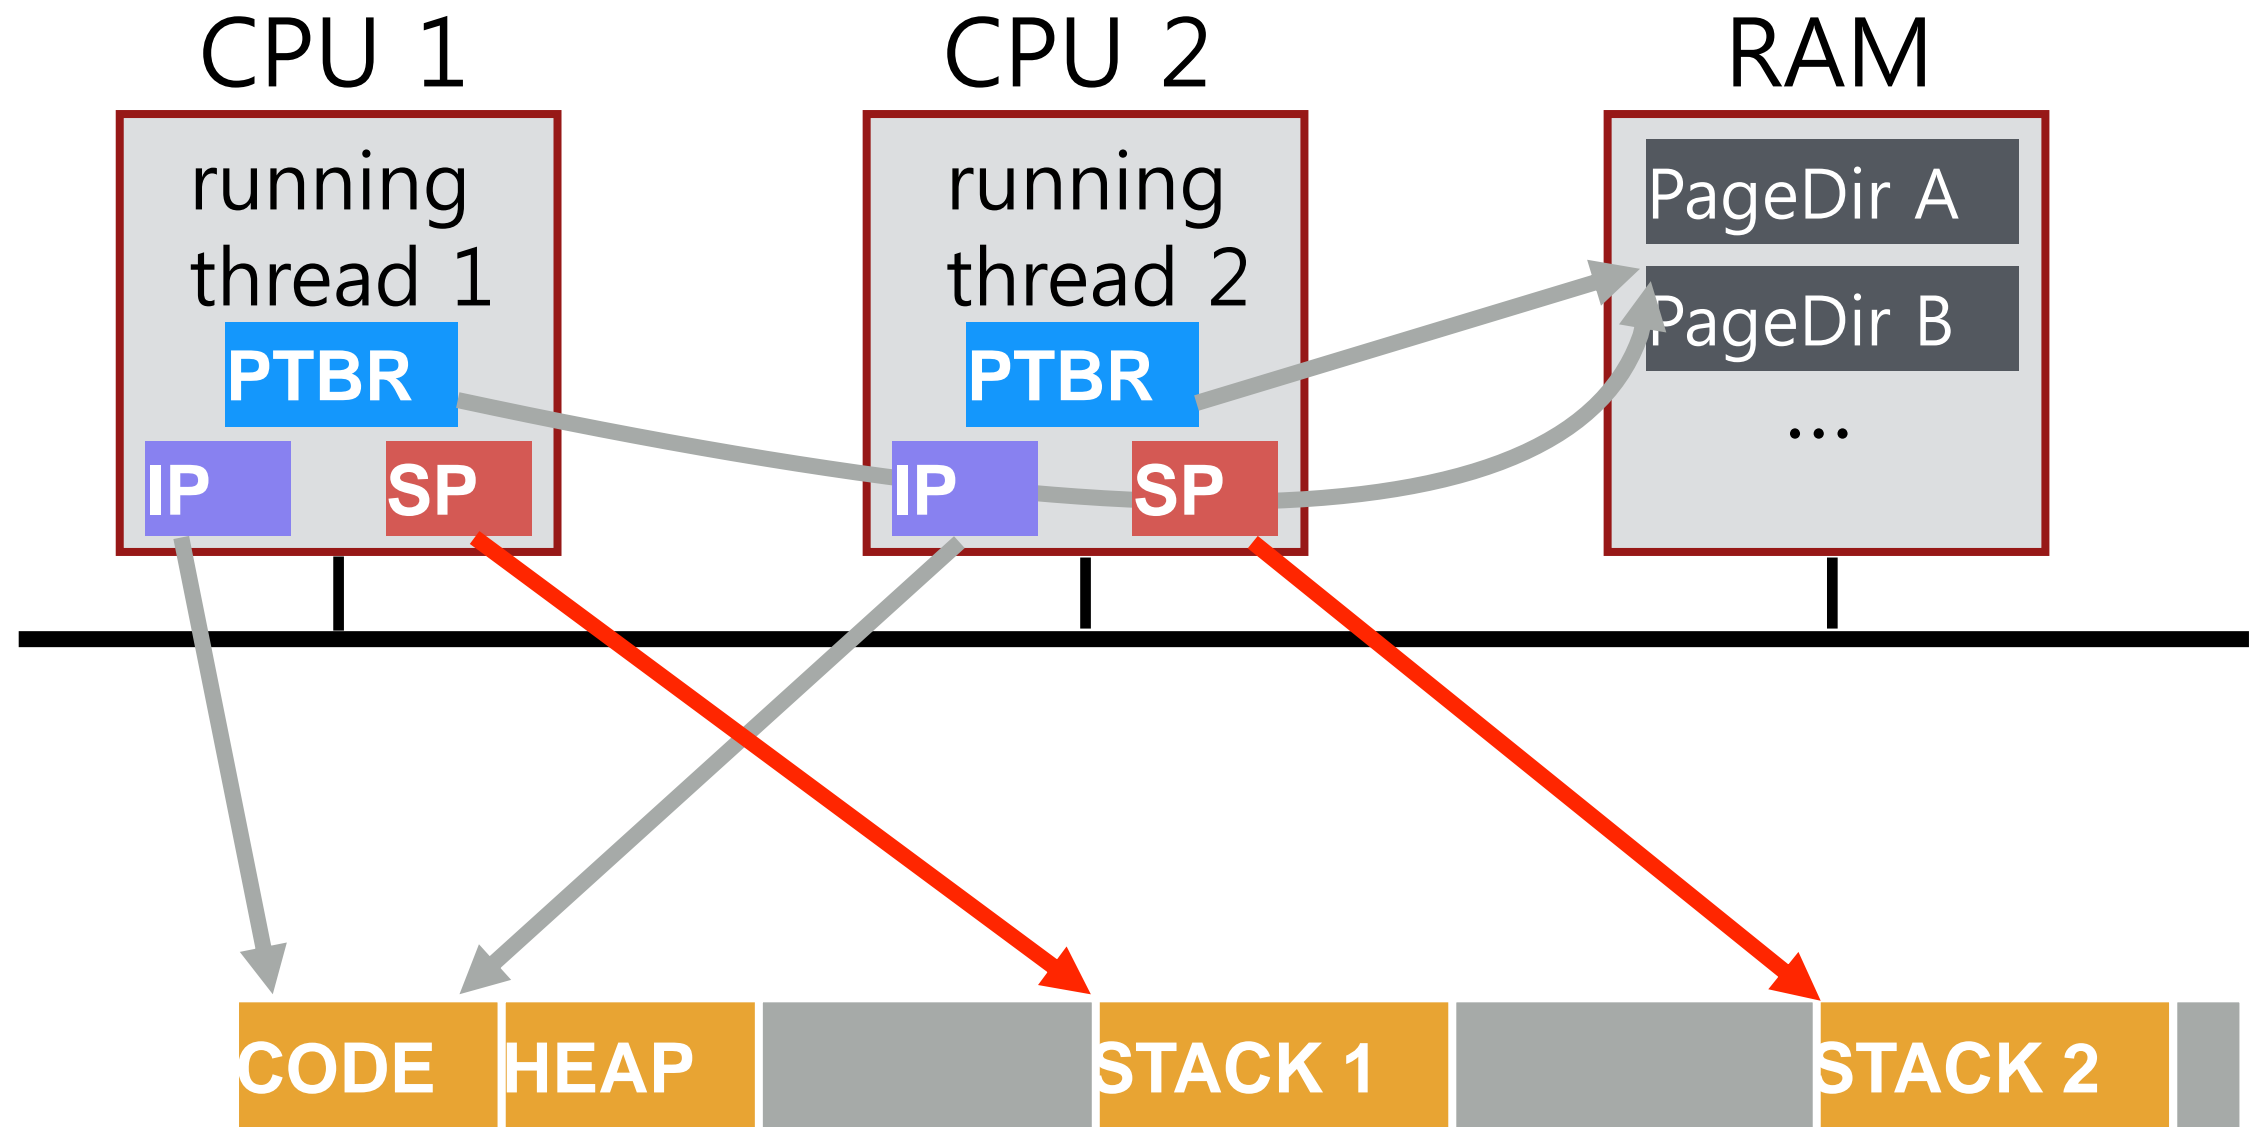
→ Different Instruction Pointers





Do threads share stack pointer?





threads executing different functions need different stacks

# THREAD VS. Process

Multiple threads within a single process share:

- Process ID (PID)
- Address space
  - Code (instructions)
  - Most data (heap)
- Open file descriptors
- Current working directory
- User and group id

Each thread has its own

- Thread ID (TID)
- Set of registers, including Program counter and Stack pointer
- Stack for local variables and return addresses  
(in same address space)

# THREAD API

Variety of thread systems exist

- POSIX Pthreads

Common thread operations

- Create
- Exit
- Join (instead of wait() for processes)



# Thread Creation

## How to create and control threads?

```
#include <pthread.h>

int
pthread_create(      pthread_t*      thread,
                    const pthread_attr_t* attr,
                    void*             (*start_routine) (void*),
                    void*             arg);
```

- `thread`: Used to interact with this thread.
- `attr`: Used to specify any attributes this thread might have.
  - Stack size, Scheduling priority, ...
- `start_routine`: the function this thread start running in.
- `arg`: the argument to be passed to the function (`start routine`)
  - *a void pointer* allows us to pass in *any type of* argument.

# Thread Creation (Cont.)

If `start_routine` instead required another type argument, the declaration would look like this:

- An integer argument:

```
int  
pthread_create(..., // first two args are the same  
                void*  (*start_routine)(int),  
                int    arg);
```

- Return an integer:

```
int  
pthread_create(..., // first two args are the same  
                int   (*start_routine)(void*),  
                void* arg);
```

# Example: Creating a Thread

```
#include <pthread.h>

typedef struct __myarg_t {
    int a;
    int b;
} myarg_t;

void *mythread(void *arg) {
    myarg_t *m = (myarg_t *) arg;
    printf("%d %d\n", m->a, m->b);
    return NULL;
}

int main(int argc, char *argv[]) {
    pthread_t p;
    int rc;

    myarg_t args;
    args.a = 10;
    args.b = 20;
    rc = pthread_create(&p, NULL, mythread, &args);
    ...
}
```

# Wait for a thread to complete

```
int pthread_join(pthread_t thread, void **value_ptr);
```

- thread: Specify which thread *to wait for*
- value\_ptr: A pointer to the return value
  - Because `pthread_join()` routine changes the value, you need to **pass in a pointer** to that value.

# Example: Waiting for Thread Completion

```
1  #include <stdio.h>
2  #include <pthread.h>
3  #include <assert.h>
4  #include <stdlib.h>
5
6  typedef struct __myarg_t {
7      int a;
8      int b;
9  } myarg_t;
10
11 typedef struct __myret_t {
12     int x;
13     int y;
14 } myret_t;
15
16 void *mythread(void *arg) {
17     myarg_t *m = (myarg_t *) arg;
18     printf("%d %d\n", m->a, m->b);
19     myret_t *r = malloc(sizeof(myret_t));
20     r->x = 1;
21     r->y = 2;
22     return (void *) r;
23 }
24
```

# Example: Waiting for Thread Completion (Cont.)

```
25  int main(int argc, char *argv[]) {
26      int rc;
27      pthread_t p;
28      myret_t *m;
29
30      myarg_t args;
31      args.a = 10;
32      args.b = 20;
33      pthread_create(&p, NULL, mythread, &args);
34      pthread_join(p, (void **) &m); // this thread has been
                                     // waiting inside of the
                                     // pthread_join() routine.
35      printf("returned %d %d\n", m->x, m->y);
36      return 0;
37 }
```

# Example: Dangerous code

Be careful with how values are returned from a thread.

```
1  void *mythread(void *arg) {
2      myarg_t *m = (myarg_t *) arg;
3      printf("%d %d\n", m->a, m->b);
4      myret_t r; // ALLOCATED ON STACK: BAD!
5      r.x = 1;
6      r.y = 2;
7      return (void *) &r;
8  }
```

- When the variable `r` returns, it is automatically **de-allocated**.

# Example: Simpler Argument Passing to a Thread

Just passing in a single value

```
1  void *mythread(void *arg) {
2      int m = (int) arg;
3      printf("%d\n", m);
4      return (void *) (arg + 1);
5  }
6
7  int main(int argc, char *argv[]) {
8      pthread_t p;
9      int rc, m;
10     pthread_create(&p, NULL, mythread, (void *) 100);
11     pthread_join(p, (void **) &m);
12     printf("returned %d\n", m);
13     return 0;
14 }
```



# Locks

Provide **mutual exclusion** to a critical section

- Interface

```
int pthread_mutex_lock(pthread_mutex_t *mutex);  
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- Usage (w/o *lock initialization* and *error check*)

```
pthread_mutex_t lock;  
pthread_mutex_lock(&lock);  
x = x + 1; // or whatever your critical section is  
pthread_mutex_unlock(&lock);
```

- No other thread holds the lock → the thread will acquire the lock and **enter the critical section**.
- If another thread hold the lock → the thread will **not return from the call** until it has acquired the lock.

# Locks (Cont.)

All locks must be properly initialized.

- One way: using `PTHREAD_MUTEX_INITIALIZER`

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
```

- The dynamic way: using `pthread_mutex_init()`

```
int rc = pthread_mutex_init(&lock, NULL);  
assert(rc == 0); // always check success!
```

# Locks (Cont.)

Check errors code when calling lock and unlock

- An example wrapper

```
// Use this to keep your code clean but check for failures
// Only use if exiting program is OK upon failure
void Pthread_mutex_lock(pthread_mutex_t *mutex) {
    int rc = pthread_mutex_lock(mutex);
    assert(rc == 0);
}
```

These two calls are used in lock acquisition

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);
int pthread_mutex_timelock(pthread_mutex_t *mutex,
                           struct timespec *abs_timeout);
```

- trylock: return failure if the lock is already held
- timelock: return after a timeout

# Locks (Cont.)

These two calls are also used in **lock acquisition**

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);  
int pthread_mutex_timelock(pthread_mutex_t *mutex,  
                           struct timespec *abs_timeout);
```

- `trylock`: return failure if the lock is already held
- `timelock`: return after a timeout or after acquiring the lock

# Condition Variables

**Condition variables** are useful when some kind of **signaling** must take place between threads.

```
int pthread_cond_wait(pthread_cond_t *cond,  
                      pthread_mutex_t *mutex);  
int pthread_cond_signal(pthread_cond_t *cond);
```

- `pthread_cond_wait`:
  - Put the calling thread to sleep.
  - Wait for some other thread to signal it.
- `pthread_cond_signal`:
  - Unblock at least one of the threads that are blocked on the condition variable

# Condition Variables (Cont.)

A thread calling wait routine:

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t init = PTHREAD_COND_INITIALIZER;

pthread_mutex_lock(&lock);
while (initialized == 0)
    pthread_cond_wait(&init, &lock);
pthread_mutex_unlock(&lock);
```

- The wait call **releases the lock** when putting said caller to sleep.
- Before returning after being woken, the wait call **re-acquire the lock**.

A thread calling signal routine:

```
pthread_mutex_lock(&lock);
initialized = 1;
pthread_cond_signal(&init);
pthread_mutex_unlock(&lock);
```

# Condition Variables (Cont.)

The waiting thread **re-checks** the condition **in a while loop**, instead of a simple if statement.

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t init = PTHREAD_COND_INITIALIZER;

pthread_mutex_lock(&lock);
while (initialized == 0)
    pthread_cond_wait(&init, &lock);
pthread_mutex_unlock(&lock);
```

- Without rechecking, the waiting thread will continue thinking that the condition has changed *even though it has not*.

# Condition Variables (Cont.)

Don't ever to this.

- A thread calling wait routine:

```
while(initialized == 0)  
    ; // spin
```

- A thread calling signal routine:

```
initialized = 1;
```

- It performs poorly in many cases. → just wastes CPU cycles.
- It is error prone.



# Compiling and Running

To compile them, you must include the header `pthread.h`

- Explicitly link with the **pthread library**, by adding the `-pthread` flag.

```
prompt> gcc -o main main.c -Wall -pthread
```

- For more information,

```
man -k pthread
```

# OS Support: Approach 1

## **User-level threads: Many-to-one thread mapping**

- Implemented by user-level runtime libraries
  - Create, schedule, synchronize threads at user-level
- OS is not aware of user-level threads
  - OS thinks each process contains only a single thread of control

## Advantages

- Does not require OS support; Portable
- Can tune scheduling policy to meet application demands
- Lower overhead thread operations since no system call

## Disadvantages?

- Cannot leverage multiprocessors
- Entire process blocks when one thread blocks

# OS Support: Approach 2

## **Kernel-level threads: One-to-one thread mapping**

- OS provides each user-level thread with a kernel thread
- Each kernel thread scheduled independently
- Thread operations (creation, scheduling, synchronization) performed by OS

### Advantages

- Each kernel-level thread can run in parallel on a multiprocessor
- When one thread blocks, other threads from process can be scheduled

### Disadvantages

- Higher overhead for thread operations
- OS must scale well with increasing number of threads

# Demo: Basic threads

# Thread Schedule #1

balance = balance + 1; balance at 0x9cd4

State:

0x9cd4: 100

%eax: ?

%rip = 0x195

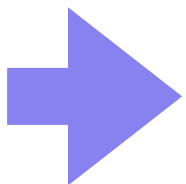
process  
control  
blocks:

Thread 1

%eax: ?  
%rip: 0x195

Thread 2

%eax: ?  
%rip: 0x195

T1 

- 0x195 mov 0x9cd4 %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax 0x9cd4A

# Thread Schedule #1

State:

0x9cd4: 100

%eax: 100

%rip = 0x19a

process  
control  
blocks:

Thread 1

%eax: ?  
%rip: 0x195

Thread 2

%eax: ?  
%rip: 0x195

T1 →

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

# Thread Schedule #1

State:

0x9cd4: 100

%eax: 101

%rip = 0x19d

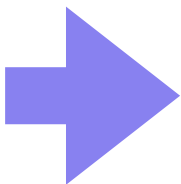
process  
control  
blocks:

Thread 1

%eax: ?  
%rip: 0x195

Thread 2

%eax: ?  
%rip: 0x195

T1 

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

# Thread Schedule #1

State:

0x9cd4: 101

%eax: 101

%rip = 0x1a2

process  
control  
blocks:

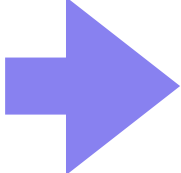
Thread 1

%eax: ?  
%rip: 0x195

Thread 2

%eax: ?  
%rip: 0x195

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

T1 



# Thread Schedule #1

State:

0x9cd4: 101

%eax: 101

%rip = 0x1a2

process  
control  
blocks:

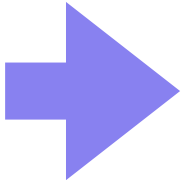
Thread 1

%eax: ?  
%rip: 0x195

Thread 2

%eax: ?  
%rip: 0x195

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

T1 

## Thread Context Switch

# Thread Schedule #1

State:

0x9cd4: 101

%eax: ?

%rip = 0x195

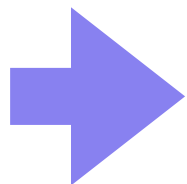
process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x1a2

Thread 2

%eax: ?  
%rip: 0x195

T2 

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

# Thread Schedule #1

State:

0x9cd4: 101

%eax: 101

%rip = 0x19a

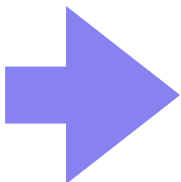
process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x1a2

Thread 2

%eax: ?  
%rip: 0x195

T2 

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

# Thread Schedule #1

State:

0x9cd4: 101

%eax: 102

%rip = 0x19d

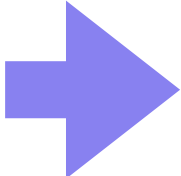
process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x1a2

Thread 2

%eax: ?  
%rip: 0x195

T2 

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

# Thread Schedule #1

State:

0x9cd4: 102

%eax: 102

%rip = 0x1a2

process  
control  
blocks:

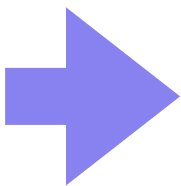
Thread 1

%eax: 101  
%rip: 0x1a2

Thread 2

%eax: ?  
%rip: 0x195

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

T2 

# Thread Schedule #1

State:

0x9cd4: 102

%eax: 102

%rip = 0x1a2

process  
control  
blocks:

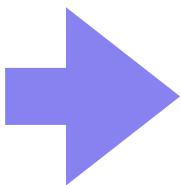
Thread 1

%eax: 101  
%rip: 0x1a2

Thread 2

%eax: ?  
%rip: 0x195

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

T2 

Desired Result!

Another schedule

# Thread Schedule #2

State:

0x9cd4: 100

%eax: ?

%rip = 0x195

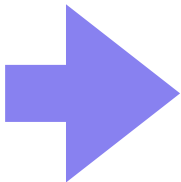
process  
control  
blocks:

Thread 1

%eax: ?  
%rip: 0x195

Thread 2

%eax: ?  
%rip: 0x195

T1 

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4



# Thread Schedule #2

State:

0x9cd4: 100

%eax: 100

%rip = 0x19a

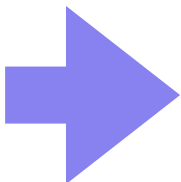
process  
control  
blocks:

Thread 1

%eax: ?  
%rip: 0x195

Thread 2

%eax: ?  
%rip: 0x195

T1 

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

# Thread Schedule #2

State:

0x9cd4: 100

%eax: 101

%rip = 0x19d

process  
control  
blocks:


Thread 1

%eax: ?  
%rip: 0x195

Thread 2

%eax: ?  
%rip: 0x195

T1



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

## Thread Context Switch

# Thread Schedule #2

State:

0x9cd4: 100

%eax: ?

%rip = 0x195

process  
control  
blocks:


Thread 1

%eax: 101  
%rip: 0x19d

Thread 2

%eax: ?  
%rip: 0x195

T2



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

# Thread Schedule #2

State:

0x9cd4: 100

%eax: 100

%rip = 0x19a

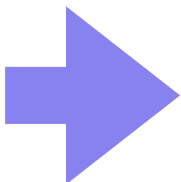
process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x19d

Thread 2

%eax: ?  
%rip: 0x195

T2 

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

# Thread Schedule #2

State:

0x9cd4: 100

%eax: 101

%rip = 0x19d

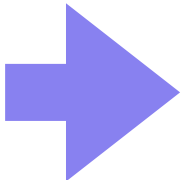
process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x19d

Thread 2

%eax: ?  
%rip: 0x195

T2 

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

# Thread Schedule #2

State:

0x9cd4: 101

%eax: 101

%rip = 0x1a2

process  
control  
blocks:

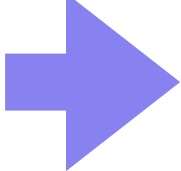
Thread 1

%eax: 101  
%rip: 0x19d

Thread 2

%eax: ?  
%rip: 0x195

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4A

T2 

# Thread Schedule #2

State:

0x9cd4: 101

%eax: 101

%rip = 0x1a2

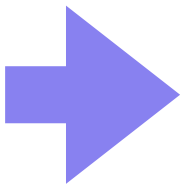
process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x19d

Thread 2

%eax: ?  
%rip: 0x195

T2 

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

## Thread Context Switch

# Thread Schedule #2

State:

0x9cd4: 101

%eax: 101

%rip = 0x19d

process  
control  
blocks:


Thread 1

%eax: 101  
%rip: 0x19d

Thread 2

%eax: 101  
%rip: 0x1a2

T1



- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

## Thread Context Switch



# Thread Schedule #2

State:

0x9cd4: 101

%eax: 101

%rip = 0x19d

process  
control  
blocks:

Thread 1

%eax: 101  
%rip: 0x19d

Thread 2

%eax: 101  
%rip: 0x1a2

T1 →

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

# Thread Schedule #2

State:

0x9cd4: 101

%eax: 101

%rip = 0x1a2

process  
control  
blocks:

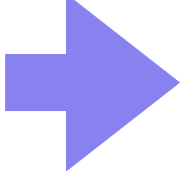
Thread 1

%eax: 101  
%rip: 0x1a2

Thread 2

%eax: 101  
%rip: 0x1a2

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

T1 

# Thread Schedule #2

State:

0x9cd4: 101

%eax: 101

%rip = 0x1a2

process  
control  
blocks:

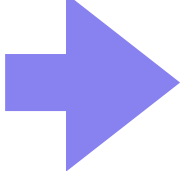
Thread 1

%eax: 101  
%rip: 0x1a2

Thread 2

%eax: 101  
%rip: 0x1a2

- 0x195 mov 0x9cd4, %eax
- 0x19a add \$0x1, %eax
- 0x19d mov %eax, 0x9cd4

T1 

WRONG Result! Final value of balance is 101

# Timeline View

## Thread 1

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

## Thread 2

```
mov 0x123, %eax  
  
add %0x2, %eax  
  
mov %eax, 0x123
```

How much is added to shared variable? 3: correct!

# Timeline View

## Thread 1

`mov 0x123, %eax`

`add %0x1, %eax`

`mov %eax, 0x123`

## Thread 2

`mov 0x123, %eax`

`add %0x2, %eax`

`mov %eax, 0x123`

How much is added?

2: incorrect!

# Timeline View

## Thread 1

`mov 0x123, %eax`

`add %0x1, %eax`

`mov %eax, 0x123`

## Thread 2

`mov 0x123, %eax`

`add %0x2, %eax`

`mov %eax, 0x123`

How much is added? 1: incorrect!

## Thread 1

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

## Thread 2

```
mov 0x123, %eax  
add %0x2, %eax  
mov %eax, 0x123
```

# Timeline View

## Thread 1

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

## Thread 2

```
mov 0x123, %eax  
add %0x2, %eax
```

```
mov %eax, 0x123
```

How much is added? 2: incorrect!



# Non-Determinism

Concurrency leads to **non-deterministic** results

- Not deterministic result: different results even with same inputs
- **race conditions**

Whether bug manifests depends on **CPU schedule!**

**Passing tests** means little

How to program: imagine **scheduler is malicious**

Assume scheduler will pick bad ordering at some point...

# What do we want?

Want 3 instructions to execute as an **uninterruptable** group  
That is, we want them to be **atomic**

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

—critical section

More general:

Need **mutual exclusion** for critical sections

- if process A is in critical section C, process B can't  
(okay if other processes do unrelated work)

# Synchronization

Build higher-level synchronization primitives in OS

- Operations that ensure correct ordering of instructions across threads

Motivation: Build them once and get them right

Monitors      Locks      Semaphores  
Condition Variables

Loads      Stores      Test&Set  
Disable Interrupts

# Locks

Goal: Provide mutual exclusion (mutex)

Three common operations:

- Allocate and Initialize
  - `Pthread_mutex_t mylock = PTHREAD_MUTEX_INITIALIZER;`
- Acquire
  - Acquire exclusion access to lock;
  - Wait if lock is not available (some other process in critical section)
  - Spin or block (relinquish CPU) while waiting
  - `Pthread_mutex_lock(&mylock);`
- Release
  - Release exclusive access to lock; let another process enter critical section
  - `Pthread_mutex_unlock(&mylock);`

More Demos

# Conclusions

Concurrency is needed to obtain high performance by utilizing multiple cores

Threads are multiple execution streams within a single process or address space (share PID and address space, own registers and stack)

Context switches within a critical section can lead to non-deterministic bugs (race conditions)

Use locks to provide mutual exclusion

# Implementing Synchronization

To implement, need atomic operations

**Atomic operation:** No other instructions can be interleaved

Examples of atomic operations

- Code between interrupts on uniprocessors
  - Disable timer interrupts, don't do any I/O
- Loads and stores of words
  - Load r1, B
  - Store r1, A
- **Special hw instructions**
  - **Test&Set**
  - **Compare&Swap**

# Implementing Locks: Attempt #1

Turn off interrupts for critical sections

Prevent dispatcher from running another thread

Code executes atomically

```
Void acquire(lockT *l) {  
    disableInterrupts();  
}
```

```
Void release(lockT *l) {  
    enableInterrupts();  
}
```

Disadvantages??



# Implementing LOCKS: Attempt #2

Code uses a single shared lock variable

```
Boolean lock = false; // shared variable
```

```
Void acquire() {  
    while (lock) /* wait */ ;  
    lock = true;  
}
```

```
Void release() {  
    lock = false;  
}
```

Why doesn't this work?