

# OSTEP

## Concurrency:

### Locks

#### Questions answered in this lecture:

Review: Why threads and mutual exclusion for critical sections?


How can locks be used to protect shared data structures such as **linked lists**?

Can locks be implemented by **disabling interrupts**?

Can locks be implemented with **loads and stores**?

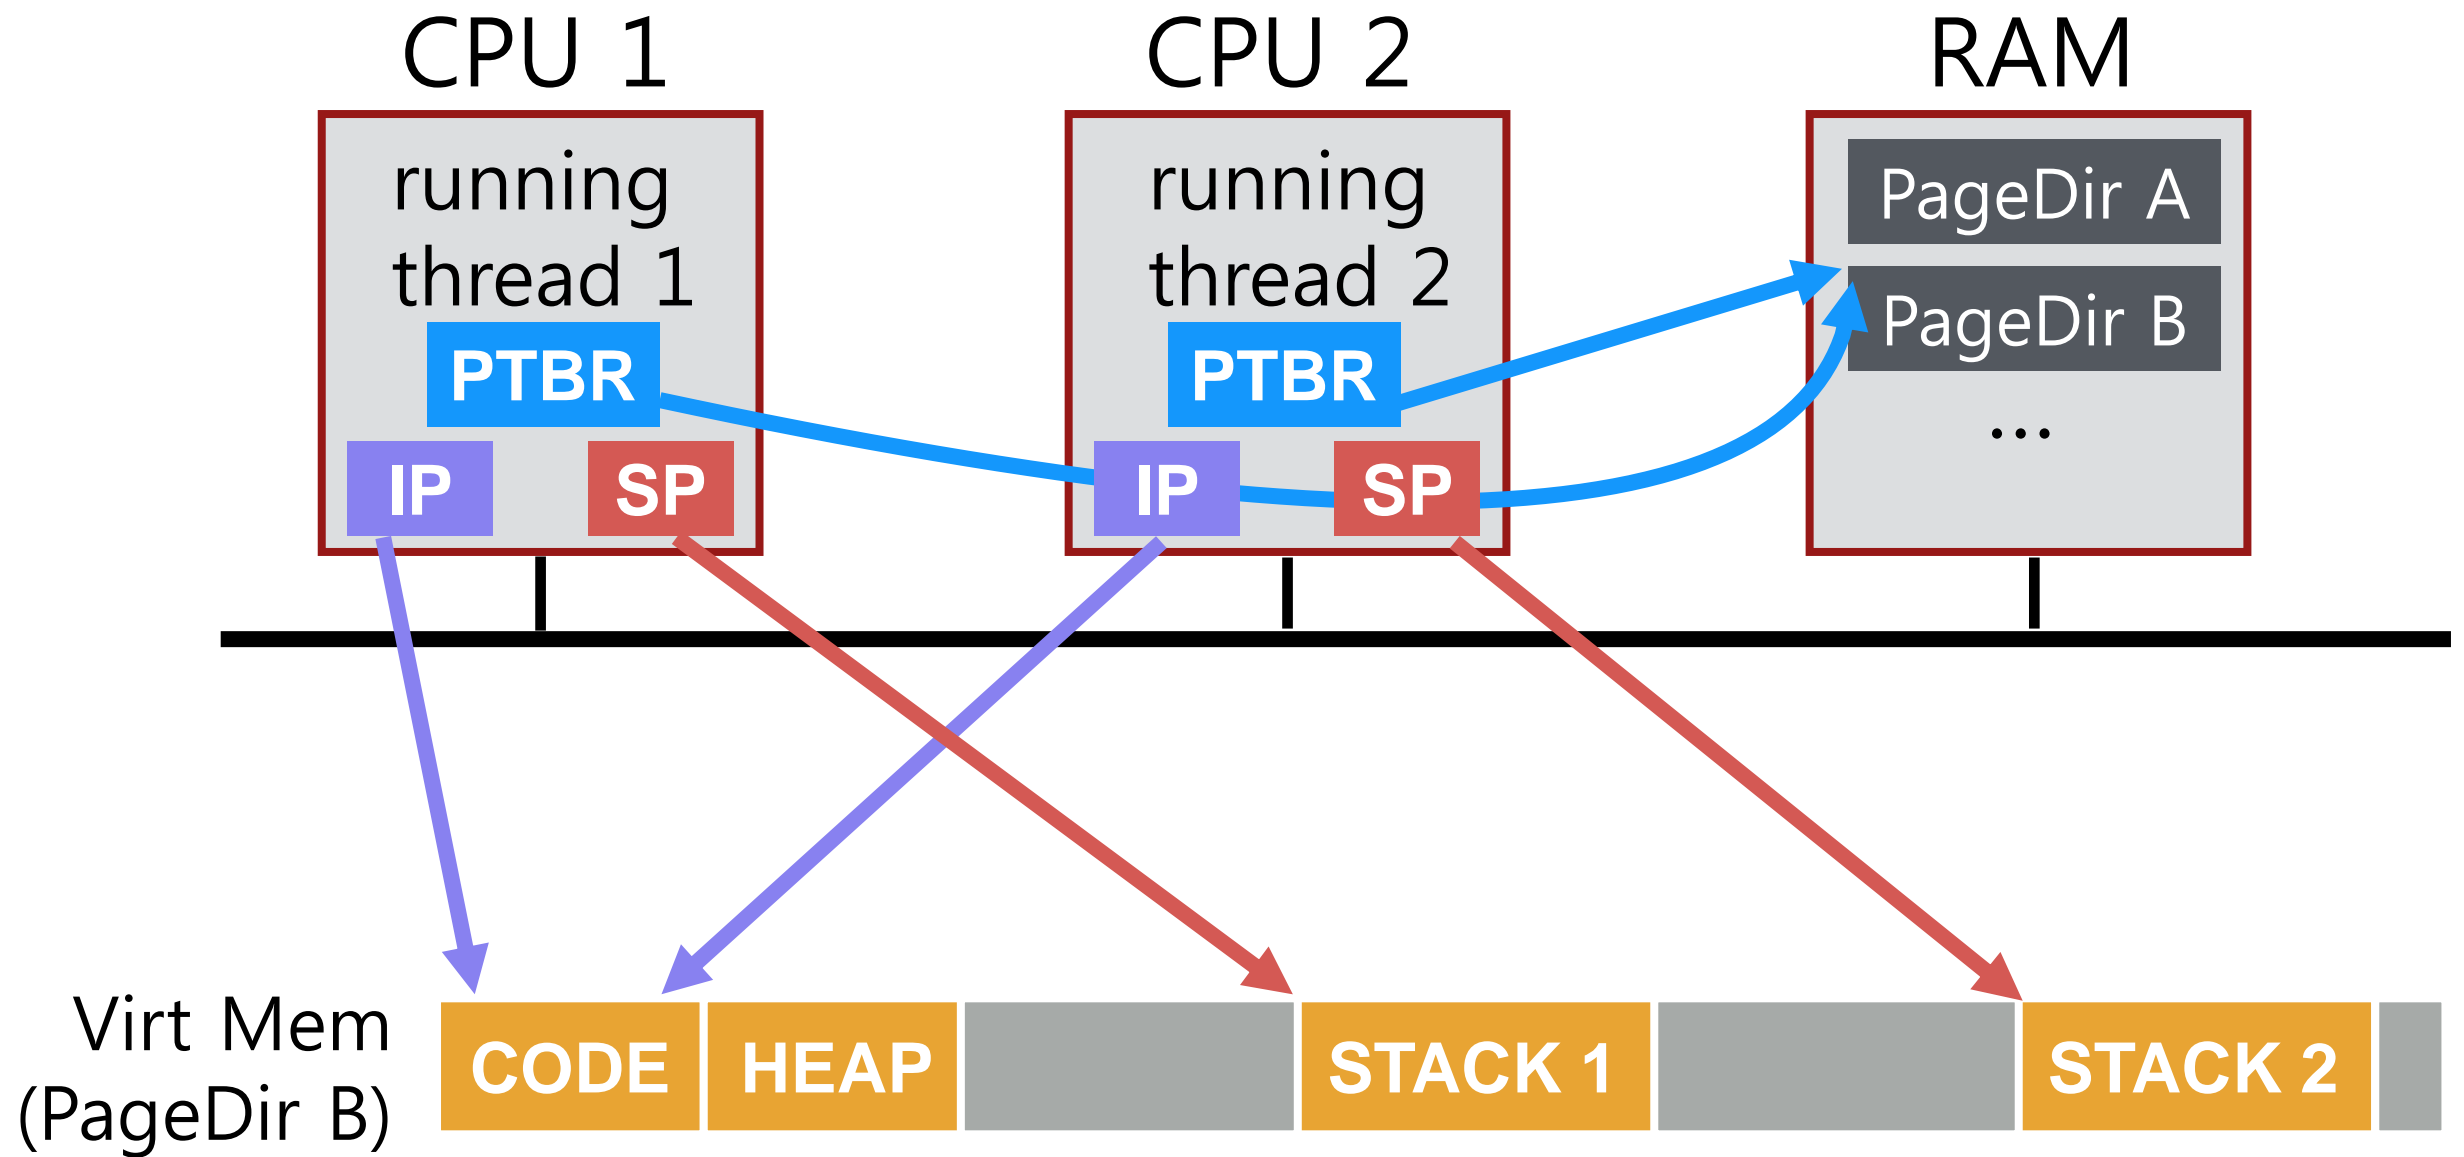
Can locks be implemented with **atomic hardware instructions**?

When are **spinlocks** a good idea?



Review:

Which registers store the same/different values across threads?



# Locks: The Basic Idea

Ensure that any **critical section** executes as if it were **a single atomic instruction**.

- An example: the canonical update of a shared variable

```
balance = balance + 1;
```

- Add some code around the critical section

```
1  lock_t mutex; // some globally-allocated lock 'mutex'
2  ...
3  lock(&mutex);
4  balance = balance + 1;
5  unlock(&mutex);
```

# Locks: The Basic Idea

Lock variable holds the state of the lock.

- **available** (or **unlocked** or **free**)
  - No thread holds the lock.
- **acquired** (or **locked** or **held**)
  - Exactly one thread holds the lock and presumably is in a critical section.

# The semantics of the lock()

lock()

- **Try to** acquire the lock.
- If no other thread holds the lock, the thread will **acquire** the lock.
- **Enter** the *critical section*.
  - This thread is said to be the owner of the lock.
- Other threads are *prevented from* entering the critical section while the first thread that holds the lock is in there.

# Pthread Locks - mutex

The name that the POSIX library uses for a lock.

- Used to provide **mutual exclusion** between threads.

```
1  pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
2
3  Pthread_mutex_lock(&lock); // wrapper for pthread_mutex_lock()
4  balance = balance + 1;
5  Pthread_mutex_unlock(&lock);
```

- We may be using *different locks* to protect *different variables* → Increase **concurrency** (a more **fine-grained** approach).

# Building A Lock

Efficient locks provided mutual exclusion at **low cost**.

Building a lock need some help from the **hardware** and the **OS**.



# Evaluating locks – Basic criteria

## **Mutual exclusion**

- Does the lock work, preventing multiple threads from entering *a critical section*?

## **Fairness**

- Does each thread contending for the lock get a fair shot at acquiring it once it is free? (Starvation)

## **Performance**

- The time overheads added by using the lock

# Controlling Interrupts

## Disable Interrupts for critical sections

- One of the earliest solutions used to provide mutual exclusion
- Invented for single-processor systems.

```
1  void lock() {  
2      DisableInterrupts();  
3  }  
4  void unlock() {  
5      EnableInterrupts();  
6  }
```

- Problem:
  - Require too much *trust* in applications
    - Greedy (or malicious) program could monopolize the processor.
  - Do not work on **multiprocessors**
  - Code that masks or unmask interrupts be executed *slowly* by modern CPUs

# Why hardware support needed?

**First attempt:** Using a *flag* denoting whether the lock is held or not.

- The code below has problems.

```
1  typedef struct __lock_t { int flag; } lock_t;
2
3  void init(lock_t *mutex) {
4      // 0 → lock is available, 1 → held
5      mutex->flag = 0;
6  }
7
8  void lock(lock_t *mutex) {
9      while (mutex->flag == 1) // TEST the flag
10         ; // spin-wait (do nothing)
11     mutex->flag = 1; // now SET it !
12 }
13
14 void unlock(lock_t *mutex) {
15     mutex->flag = 0;
16 }
```

# Why hardware support needed? (Cont.)

- **Problem 1:** No Mutual Exclusion (assume `flag=0` to begin)

Thread1	Thread2
<pre>call lock() while (flag == 1) interrupt: switch to Thread 2</pre>	<pre>call lock() while (flag == 1) flag = 1; interrupt: switch to Thread 1</pre>
<pre>flag = 1; // set flag to 1 (too!)</pre>	

- **Problem 2:** Spin-waiting wastes time waiting for another thread.
- So, we need an atomic instruction supported by **Hardware!**
  - *test-and-set* instruction, also known as *atomic exchange*

# Test And Set (Atomic Exchange)

An instruction to support the creation of simple locks

```
1  int TestAndSet(int *ptr, int new) {  
2      int old = *ptr;    // fetch old value at ptr  
3      *ptr = new; // store 'new' into ptr  
4      return old; // return the old value  
5  }
```

- **return**(testing) old value pointed to by the `ptr`.
- *Simultaneously* **update**(setting) said value to `new`.
- This sequence of operations is **performed atomically**.

# Evaluating Spin Locks

## **Correctness:** yes

- The spin lock only allows a single thread to entry the critical section.

## **Fairness:** no

- Spin locks don't provide any fairness guarantees.
- Indeed, a thread spinning may spin *forever*.

## **Performance:**

- In the single CPU, performance overheads can be quite *painful*.
- If the number of threads roughly equals the number of CPUs, spin locks work *reasonably well*.

# Compare-And-Swap

Test whether the value at the address(`ptr`) is equal to expected.

- *If so, update* the memory location pointed to by `ptr` with the new value.
- *In either case, return* the actual value at that memory location.

```
1  int CompareAndSwap(int *ptr, int expected, int new) {
2      int actual = *ptr;
3      if (actual == expected)
4          *ptr = new;
5      return actual;
6  }
```

**Compare-and-Swap hardware atomic instruction (C-style)**

```
1  void lock(lock_t *lock) {
2      while (CompareAndSwap(&lock->flag, 0, 1) == 1)
3          ; // spin
4  }
```

**Spin lock with compare-and-swap**

# Compare-And-Swap (Cont.)

## C-callable x86-version of compare-and-swap

```
1  char CompareAndSwap(int *ptr, int old, int new) {
2      unsigned char ret;
3
4      // Note that sete sets a 'byte' not the word
5      __asm__ __volatile__ (
6          " lock\n"
7          " cmpxchgl %2,%1\n"
8          " sete %0\n"
9          : "=q" (ret), "=m" (*ptr)
10         : "r" (new), "m" (*ptr), "a" (old)
11         : "memory");
12     return ret;
13 }
```



# Load-Linked and Store-Conditional

```
1  int LoadLinked(int *ptr) {
2      return *ptr;
3  }
4
5  int StoreConditional(int *ptr, int value) {
6      if (no one has updated *ptr since the LoadLinked to this address) {
7          *ptr = value;
8          return 1; // success!
9      } else {
10         return 0; // failed to update
11     }
12 }
```

## Load-linked And Store-conditional

The store-conditional *only succeeds* if **no intermittent store** to the address has taken place.

- **success:** return 1 and update the value at `ptr` to `value`.
- **fail:** the value at `ptr` is not updates and 0 is returned.

# Load-Linked and Store-Conditional (Cont.)

```
1  void lock(lock_t *lock) {
2      while (1) {
3          while (LoadLinked(&lock->flag) == 1)
4              ; // spin until it's zero
5          if (StoreConditional(&lock->flag, 1) == 1)
6              return; // if set-it-to-1 was a success: all done
7                      // otherwise: try it all over again
8      }
9  }
10
11 void unlock(lock_t *lock) {
12     lock->flag = 0;
13 }
```

## Using LL/SC To Build A Lock

```
1  void lock(lock_t *lock) {
2      while (LoadLinked(&lock->flag) || !StoreConditional(&lock->flag, 1))
3          ; // spin
4  }
```

## A more concise form of the lock() using LL/SC

# Fetch-And-Add

**Atomically increment** a value while returning the old value at a particular address.

```
1  int FetchAndAdd(int *ptr) {  
2      int old = *ptr;  
3      *ptr = old + 1;  
4      return old;  
5  }
```

**Fetch-And-Add Hardware atomic instruction (C-style)**

# Ticket Lock (1)

**Ticket lock** can be built with fetch-and add.

- Ensure progress for all threads. → **fairness**

```
1  typedef struct __lock_t {
2      int ticket;
3      int turn;
4  } lock_t;
5
6  void lock_init(lock_t *lock) {
7      lock->ticket = 0;
8      lock->turn = 0;
9  }
10
11 void lock(lock_t *lock) {
12     int myturn = FetchAndAdd(&lock->ticket);
13     while (lock->turn < myturn)
14         ; // spin
15 }
16 void unlock(lock_t *lock) {
17     FetchAndAdd(&lock->turn);
18 }
```

# Ticket Lock (2)

**Ticket lock** can be built with fetch-and add.

- Ensure progress for all threads. → **fairness**

```
1  typedef struct __lock_t {
2      int ticket;
3      int turn;
4  } lock_t;
5
6  void lock_init(lock_t *lock) {
7      lock->ticket = 0;
8      lock->turn = 0;
9  }
10
11 void lock(lock_t *lock) {
12     int myturn = FetchAndAdd(&lock->ticket);
13     while (lock->turn != myturn)
14         ; // spin
15 }
16 void unlock(lock_t *lock) {
17     FetchAndAdd(&lock->turn);
18 }
```

# So Much Spinning

Hardware-based spin locks are **simple** and they work.

In some cases, these solutions can be quite **inefficient**.

- Any time a thread gets caught *spinning*, it **wastes an entire time slice** doing nothing but checking a value.

How To Avoid *Spinning*?  
We'll need **OS Support** too!

# A Simple Approach: Just Yield

When you are going to spin, **give up the CPU** to another thread.

- OS system call moves the caller from the *running state* to the *ready state*.
- The cost of a **context switch** can be substantial and the **starvation** problem still exists.

```
1  void init() {  
2      flag = 0;  
3  }  
4  
5  void lock() {  
6      while (TestAndSet(&flag, 1) == 1)  
7          yield(); // give up the CPU  
8  }  
9  
10 void unlock() {  
11     flag = 0;  
12 }
```

**Lock with Test-and-set and Yield**

# Using Queues: Sleeping Instead of Spinning

- **Queue** to keep track of which threads are waiting to enter the lock.
- `park()`
  - Put a calling thread to sleep
- `unpark(threadID)`
  - Wake a particular thread as designated by `threadID`.



# Using Queues: Sleeping Instead of Spinning

```
1  typedef struct __lock_t { int flag; int guard; queue_t *q; } lock_t;
2
3  void lock_init(lock_t *m) {
4      m->flag = 0;
5      m->guard = 0;
6      queue_init(m->q);
7  }
8
9  void lock(lock_t *m) {
10     while (TestAndSet(&m->guard, 1) == 1)
11         ; // acquire guard lock by spinning
12     if (m->flag == 0) {
13         m->flag = 1; // lock is acquired
14         m->guard = 0;
15     } else {
16         queue_add(m->q, gettid());
17         m->guard = 0;
18         park();
19     }
20 }
21 ...
```

**Lock With Queues, Test-and-set, Yield, And Wakeup**

# Using Queues: Sleeping Instead of Spinning

```
22 void unlock(lock_t *m) {
23     while (TestAndSet(&m->guard, 1) == 1)
24         ; // acquire guard lock by spinning
25     if (queue_empty(m->q))
26         m->flag = 0; // let go of lock; no one wants it
27     else
28         unpark(queue_remove(m->q)); // hold lock (for next thread!)
29     m->guard = 0;
30 }
```

**Lock With Queues, Test-and-set, Yield, And Wakeup (Cont.)**

# Wakeup/waiting race

- In case of releasing the lock (*thread A*) just before the call to `park()` (*thread B*) → Thread B would **sleep forever** (potentially).
- **Solaris** solves this problem by adding a third system call: `setpark()`.
  - By calling this routine, a thread can indicate it *is about to* `park`.
  - If it happens to be interrupted and another thread calls `unpark` before `park` is actually called, the subsequent `park` returns immediately instead of sleeping.

```
1      queue_add(m->q, gettid());  
2      setpark(); // new code  
3      m->guard = 0;  
4      park();
```

**Code modification inside of `lock()`**

# Futex

Linux provides a **futex** (is similar to Solaris's `park` and `unpark`).

- `futex_wait(address, expected)`
  - Put the calling thread to sleep
  - If the value at `address` is not equal to `expected`, the call returns immediately.
- `futex_wake(address)`
  - Wake one thread that is waiting on the queue.

# Futex (Cont.)

Snippet from `lowlevellock.h` in the **nptl** library

- The high bit of the integer `v`: track whether the lock is held or not
- All the other bits : the number of waiters

```
1  void mutex_lock(int *mutex) {
2      int v;
3      /* Bit 31 was clear, we got the mutex (this is the fastpath) */
4      if (atomic_bit_test_set(mutex, 31) == 0)
5          return;
6      atomic_increment(mutex);
7      while (1) {
8          if (atomic_bit_test_set(mutex, 31) == 0) {
9              atomic_decrement(mutex);
10             return;
11         }
12         /* We have to wait now. First make sure the futex value
13         we are monitoring is truly negative (i.e. locked). */
14         v = *mutex;
15         ...
```

## Linux-based Futex Locks

# Futex (Cont.)

```
16         if (v >= 0)
17             continue;
18         futex_wait(mutex, v);
19     }
20 }
21
22 void mutex_unlock(int *mutex) {
23     /* Adding 0x80000000 to the counter results in 0 if and only if
24        there are not other interested threads */
25     if (atomic_add_zero(mutex, 0x80000000))
26         return;
27     /* There are other threads waiting for this mutex,
28        wake one of them up */
29     futex_wake(mutex);
30 }
```

## Linux-based Futex Locks (Cont.)

# Two-Phase Locks

A two-phase lock realizes that **spinning can be useful** if the lock *is about to* be released.

- **First phase**

- The lock spins for a while, *hoping that* it can acquire the lock.
- If the lock is not acquired during the first spin phase, a second phase is entered,

- **Second phase**

- The caller is put to sleep.
- The caller is only woken up when the lock becomes free later.

# Lock-based Concurrent Data Structures



# Lock-based Concurrent Data structure

Adding locks to a data structure makes the structure **thread safe**.

- How locks are added determine both the **correctness** and **performance** of the data structure.

# Example: Concurrent Counters without Locks

Simple but not scalable

```
1  typedef struct __counter_t {
2      int value;
3  } counter_t;
4
5  void init(counter_t *c) {
6      c->value = 0;
7  }
8
9  void increment(counter_t *c) {
10     c->value++;
11 }
12
13 void decrement(counter_t *c) {
14     c->value--;
15 }
16
17 int get(counter_t *c) {
18     return c->value;
19 }
```

# Example: Concurrent Counters with Locks

## Add a **single lock**.

- The lock is acquired when calling a routine that manipulates the data structure.

```
1     typedef struct __counter_t {
2         int value;
3         pthread_lock_t lock;
4     } counter_t;
5
6     void init(counter_t *c) {
7         c->value = 0;
8         Pthread_mutex_init(&c->lock, NULL);
9     }
10
11    void increment(counter_t *c) {
12        Pthread_mutex_lock(&c->lock);
13        c->value++;
14        Pthread_mutex_unlock(&c->lock);
15    }
16
```


# Example: Concurrent Counters with Locks (Cont.)

```
(Cont.)
17  void decrement(counter_t *c) {
18      pthread_mutex_lock(&c->lock);
19      c->value--;
20      pthread_mutex_unlock(&c->lock);
21  }
22
23  int get(counter_t *c) {
24      pthread_mutex_lock(&c->lock);
25      int rc = c->value;
26      pthread_mutex_unlock(&c->lock);
27      return rc;
28  }
```

# The performance costs of the simple approach

Each thread updates a single shared counter.

- Each thread updates the counter one million times.
- iMac with four Intel 2.7GHz i5 CPUs.



**Performance of Traditional vs. Sloppy Counters**  
(Threshold of Sloppy,  $s$ , is set to 1024)

Synchronized counter **scales poorly.**

# Perfect Scaling

- Even though more work is done, it is **done in parallel**.
- The time taken to complete the task is *not increased*.

# Sloppy counter

- The sloppy counter works by representing ...
  - A single **logical counter** via numerous local physical counters, on per CPU core
  - A single **global counter**
  - There are **locks**:
    - One for each local counter and one for the global counter
- Example: on a machine with four CPUs
  - Four local counters
  - One global counter

# The basic idea of sloppy counting

- When a thread running on a core wishes to increment the counter.
  - It increment its local counter.
- Each CPU has its own local counter:
  - Threads across CPUs can update local counters *without contention*.
  - Thus counter updates are **scalable**.
- The local values are periodically transferred to the global counter.
  - Acquire the global lock
  - Increment it by the local counter's value
  - The local counter is then reset to zero.



# The basic idea of sloppy counting (Cont.)

- How often the local-to-global transfer occurs is determined by a threshold,  $S$  (sloppiness).
  - The smaller  $S$ :
    - The more the counter behaves like the *non-scalable counter*.
  - The bigger  $S$ :
    - The more scalable the counter.
    - The further off the global value might be from the *actual count*.

# Sloppy counter example

## Tracing the Sloppy Counters


- The threshold  $S$  is set to 5.
- There are threads on each of four CPUs
- Each thread updates their local counters  $L_1 \dots L_4$ .

Time	$L_1$	$L_2$	$L_3$	$L_4$	$G$
0	0	0	0	0	0
1	0	0	1	1	0
2	1	0	2	1	0
3	2	0	3	1	0
4	3	0	3	2	0
5	4	1	3	3	0
6	$5 \rightarrow 0$	1	3	4	5 (from $L_1$ )
7	0	2	4	$5 \rightarrow 0$	10 (from $L_4$ )

# Importance of the threshold value $S$

Each four threads increments a counter 1 million times on four CPUs.

- Low  $S \rightarrow$  Performance is **poor**, The global count is always quire **accurate**.
- High  $S \rightarrow$  Performance is **excellent**, The global count **lags**.



Scaling Sloppy Counters

# Sloppy Counter Implementation

```
1  typedef struct __counter_t {
2      int global;           // global count
3      pthread_mutex_t glock; // global lock
4      int local[NUMCPUS];   // local count (per cpu)
5      pthread_mutex_t llock[NUMCPUS]; // ... and locks
6      int threshold; // update frequency
7  } counter_t;
8
9  // init: record threshold, init locks, init values
10 //      of all local counts and global count
11 void init(counter_t *c, int threshold) {
12     c->threshold = threshold;
13
14     c->global = 0;
15     pthread_mutex_init(&c->glock, NULL);
16
17     int i;
18     for (i = 0; i < NUMCPUS; i++) {
19         c->local[i] = 0;
20         pthread_mutex_init(&c->llock[i], NULL);
21     }
22 }
23
```

# Sloppy Counter Implementation (Cont.)

```
(Cont.)
24 // update: usually, just grab local lock and update local amount
25 //           once local count has risen by 'threshold', grab global
26 //           lock and transfer local values to it
27 void update(counter_t *c, int threadID, int amt) {
28     pthread_mutex_lock(&c->llock[threadID]);
29     c->local[threadID] += amt;    // assumes amt > 0
30     if (c->local[threadID] >= c->threshold) { // transfer to global
31         pthread_mutex_lock(&c->glock);
32         c->global += c->local[threadID];
33         pthread_mutex_unlock(&c->glock);
34         c->local[threadID] = 0;
35     }
36     pthread_mutex_unlock(&c->llock[threadID]);
37 }
38
39 // get: just return global amount (which may not be perfect)
40 int get(counter_t *c) {
41     pthread_mutex_lock(&c->glock);
42     int val = c->global;
43     pthread_mutex_unlock(&c->glock);
44     return val;    // only approximate!
45 }
```

# Concurrent Linked Lists

```
1    // basic node structure
2    typedef struct __node_t {
3        int key;
4        struct __node_t *next;
5    } node_t;
6
7    // basic list structure (one used per list)
8    typedef struct __list_t {
9        node_t *head;
10       pthread_mutex_t lock;
11    } list_t;
12
13    void List_Init(list_t *L) {
14        L->head = NULL;
15        pthread_mutex_init(&L->lock, NULL);
16    }
17
18    (Cont.)
```

# Concurrent Linked Lists

(Cont.)

```
18     int List_Insert(list_t *L, int key) {
19         pthread_mutex_lock(&L->lock);
20         node_t *new = malloc(sizeof(node_t));
21         if (new == NULL) {
22             perror("malloc");
23             pthread_mutex_unlock(&L->lock);
24             return -1; // fail
26         new->key = key;
27         new->next = L->head;
28         L->head = new;
29         pthread_mutex_unlock(&L->lock);
30         return 0; // success
31     }
```

(Cont.)

# Concurrent Linked Lists (Cont.)

```
(Cont.)
32
32     int List_Lookup(list_t *L, int key) {
33         pthread_mutex_lock(&L->lock);
34         node_t *curr = L->head;
35         while (curr) {
36             if (curr->key == key) {
37                 pthread_mutex_unlock(&L->lock);
38                 return 0; // success
39             }
40             curr = curr->next;
41         }
42         pthread_mutex_unlock(&L->lock);
43         return -1; // failure
44     }
```



# Concurrent Linked Lists (Cont.)

- The code **acquires** a lock in the insert routine upon entry.
- The code **releases** the lock upon exit.
  - If `malloc()` happens to *fail*, the code must also release the lock before failing the insert.
  - This kind of exceptional control flow has been shown to be **quite error prone**.
  - **Solution:** The lock and release *only surround* the actual critical section in the insert code

# Concurrent Linked List: Rewritten

```
1  void List_Init(list_t *L) {
2      L->head = NULL;
3      pthread_mutex_init(&L->lock, NULL);
4  }
5
6  void List_Insert(list_t *L, int key) {
7      // synchronization not needed
8      node_t *new = malloc(sizeof(node_t));
9      if (new == NULL) {
10         perror("malloc");
11         return;
12     }
13     new->key = key;
14
15     // just lock critical section
16     pthread_mutex_lock(&L->lock);
17     new->next = L->head;
18     L->head = new;
19     pthread_mutex_unlock(&L->lock);
20 }
21
```

# Concurrent Linked List: Rewritten (Cont.)

```
(Cont.)
22  int List_Lookup(list_t *L, int key) {
23      int rv = -1;
24      pthread_mutex_lock(&L->lock);
25      node_t *curr = L->head;
26      while (curr) {
27          if (curr->key == key) {
28              rv = 0;
29              break;
30          }
31          curr = curr->next;
32      }
33      pthread_mutex_unlock(&L->lock);
34      return rv; // now both success and failure
35  }
```

# Scaling Linked List

## Hand-over-hand locking (lock coupling)

- Add **a lock per node** of the list instead of having a single lock for the entire list.
- When traversing the list,
  - First grabs the next node's lock.
  - And then releases the current node's lock.
- Enable a high degree of concurrency in list operations.
  - However, in practice, the overheads of acquiring and releasing locks for each node of a list traversal is *prohibitive*.

# Michael and Scott Concurrent Queues

There are two locks.

- One for the **head** of the queue.
- One for the **tail**.
- The goal of these two locks is to enable concurrency of *enqueue* and *dequeue* operations.

Add a dummy node

- Allocated in the queue initialization code
- Enable the separation of head and tail operations

# Concurrent Queues (Cont.)

```
1  typedef struct __node_t {
2      int value;
3      struct __node_t *next;
4  } node_t;
5
6  typedef struct __queue_t {
7      node_t *head;
8      node_t *tail;
9      pthread_mutex_t headLock;
10     pthread_mutex_t tailLock;
11 } queue_t;
12
13 void Queue_Init(queue_t *q) {
14     node_t *tmp = malloc(sizeof(node_t));
15     tmp->next = NULL;
16     q->head = q->tail = tmp;
17     pthread_mutex_init(&q->headLock, NULL);
18     pthread_mutex_init(&q->tailLock, NULL);
19 }
20
(Cont.)
```

# Concurrent Queues (Cont.)

(Cont.)

```
21     void Queue_Enqueue(queue_t *q, int value) {
22         node_t *tmp = malloc(sizeof(node_t));
23         assert(tmp != NULL);
24
25         tmp->value = value;
26         tmp->next = NULL;
27
28         pthread_mutex_lock(&q->tailLock);
29         q->tail->next = tmp;
30         q->tail = tmp;
31         pthread_mutex_unlock(&q->tailLock);
32     }
```

(Cont.)

# Concurrent Queues (Cont.)

```
(Cont.)
33  int Queue_Dequeue(queue_t *q, int *value) {
34      pthread_mutex_lock(&q->headLock);
35      node_t *tmp = q->head;
36      node_t *newHead = tmp->next;
37      if (newHead == NULL) {
38          pthread_mutex_unlock(&q->headLock);
39          return -1; // queue was empty
40      }
41      *value = newHead->value;
42      q->head = newHead;
43      pthread_mutex_unlock(&q->headLock);
44      free(tmp);
45      return 0;
46  }
```



# Concurrent Hash Table


Focus on a simple hash table

- The hash table does not resize.
- Built using the concurrent lists
- It uses a **lock per hash bucket** each of which is represented by *a list*.

# Performance of Concurrent Hash Table

From 10,000 to 50,000 concurrent updates from each of four threads.

- iMac with four Intel 2.7GHz i5 CPUs.



**The simple concurrent hash table *scales* *magnificently*.**

# Concurrent Hash Table

```
1      #define BUCKETS (101)
2
3      typedef struct __hash_t {
4          list_t lists[BUCKETS];
5      } hash_t;
6
7      void Hash_Init(hash_t *H) {
8          int i;
9          for (i = 0; i < BUCKETS; i++) {
10              List_Init(&H->lists[i]);
11          }
12      }
13
14      int Hash_Insert(hash_t *H, int key) {
15          int bucket = key % BUCKETS;
16          return List_Insert(&H->lists[bucket], key);
17      }
18
19      int Hash_Lookup(hash_t *H, int key) {
20          int bucket = key % BUCKETS;
21          return List_Lookup(&H->lists[bucket], key);
22      }
```

# Context Switch

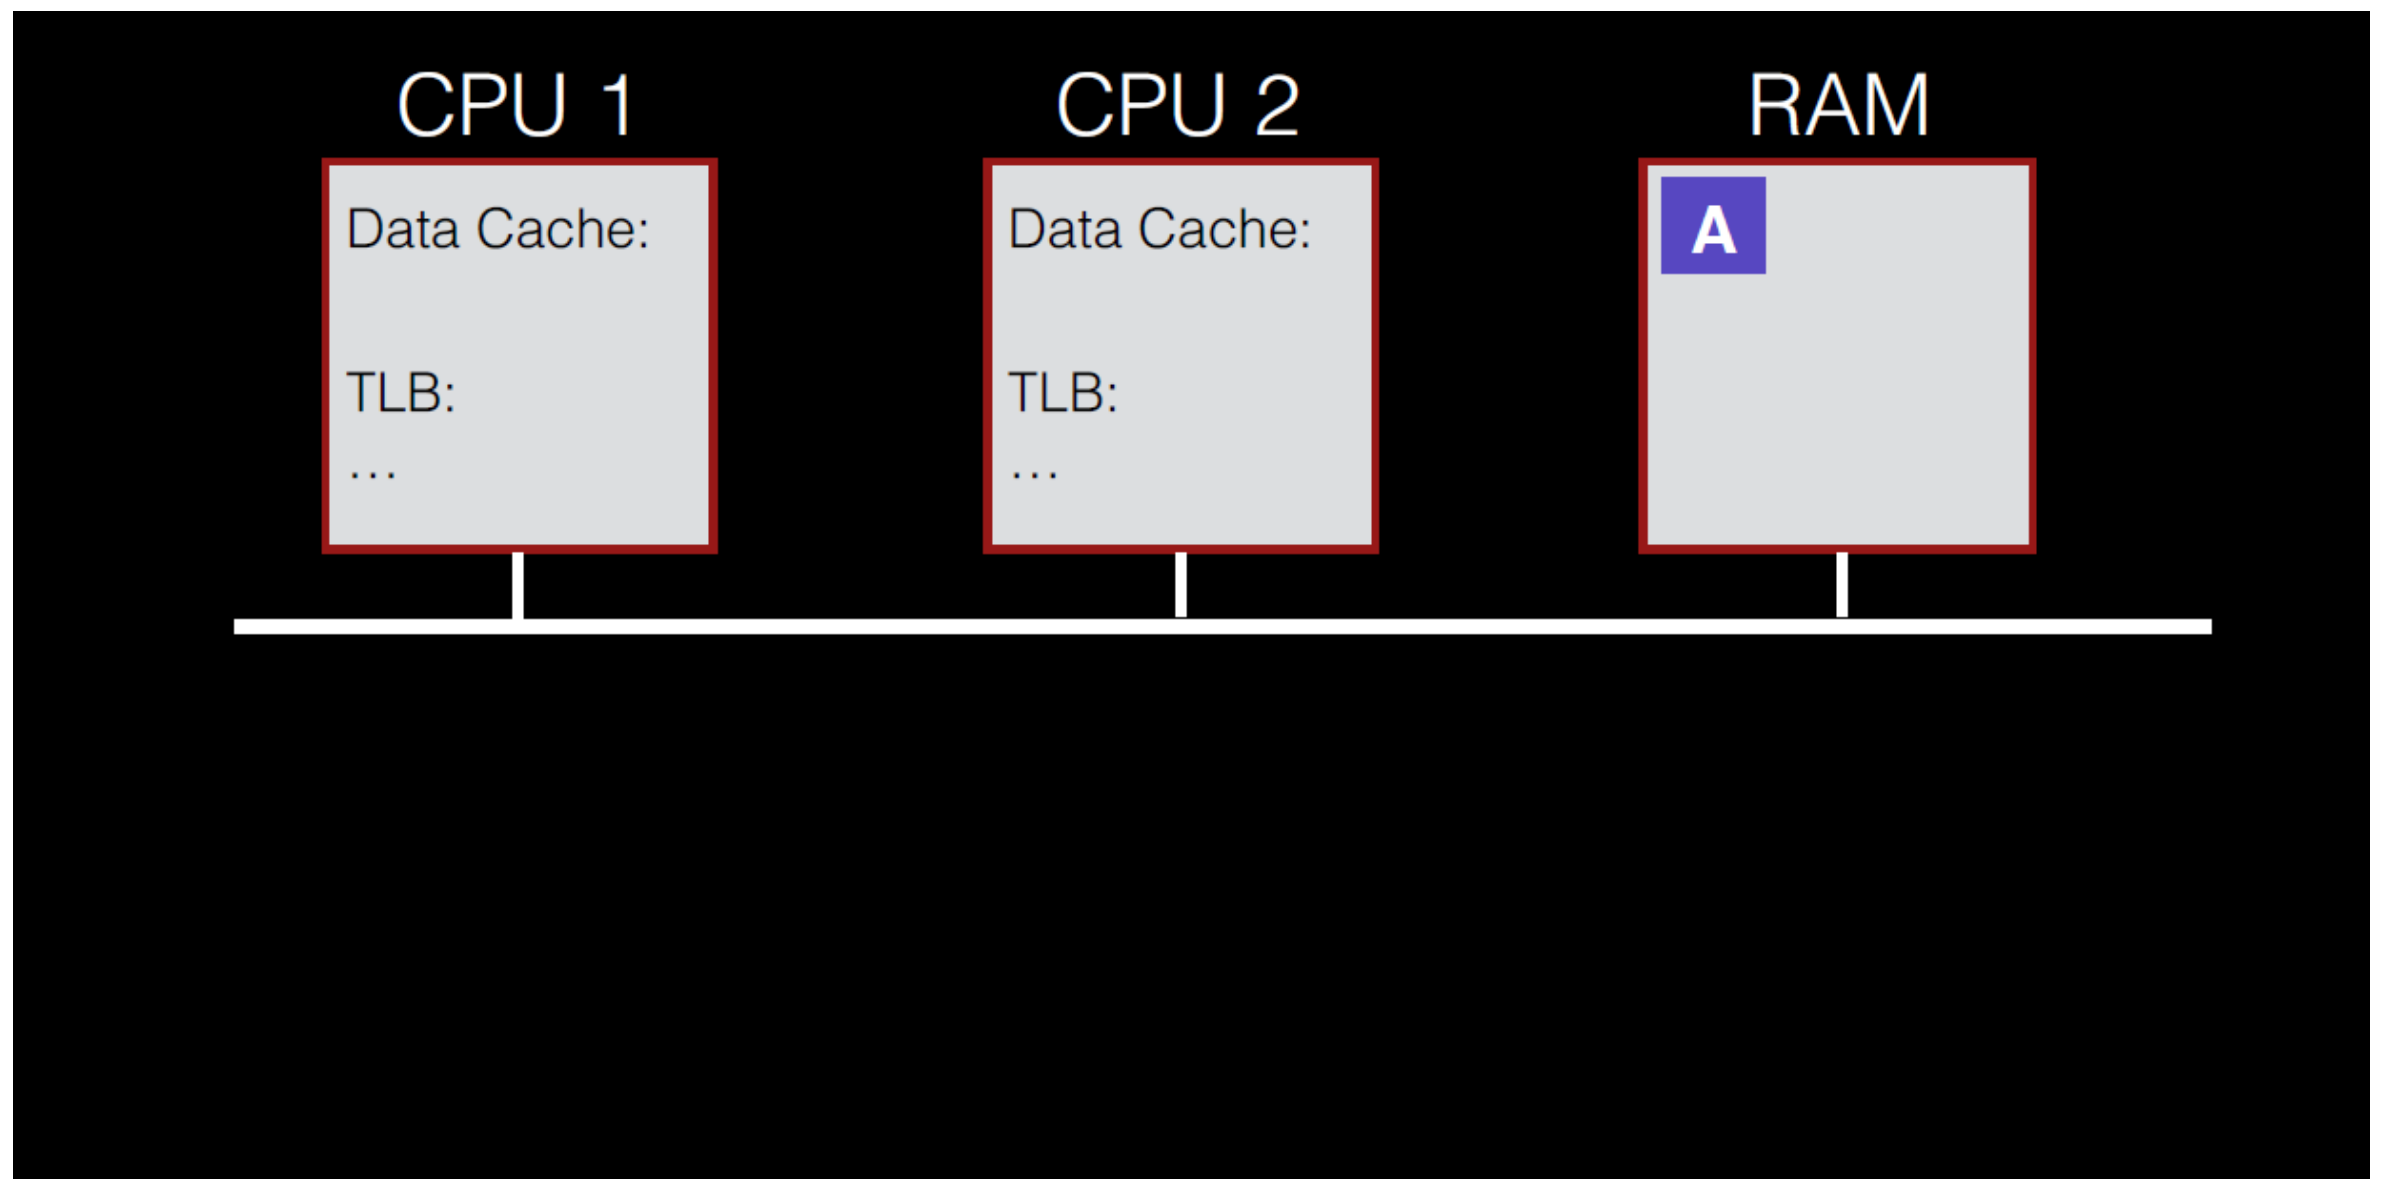
Why is switching between threads cheaper than switching between processes?

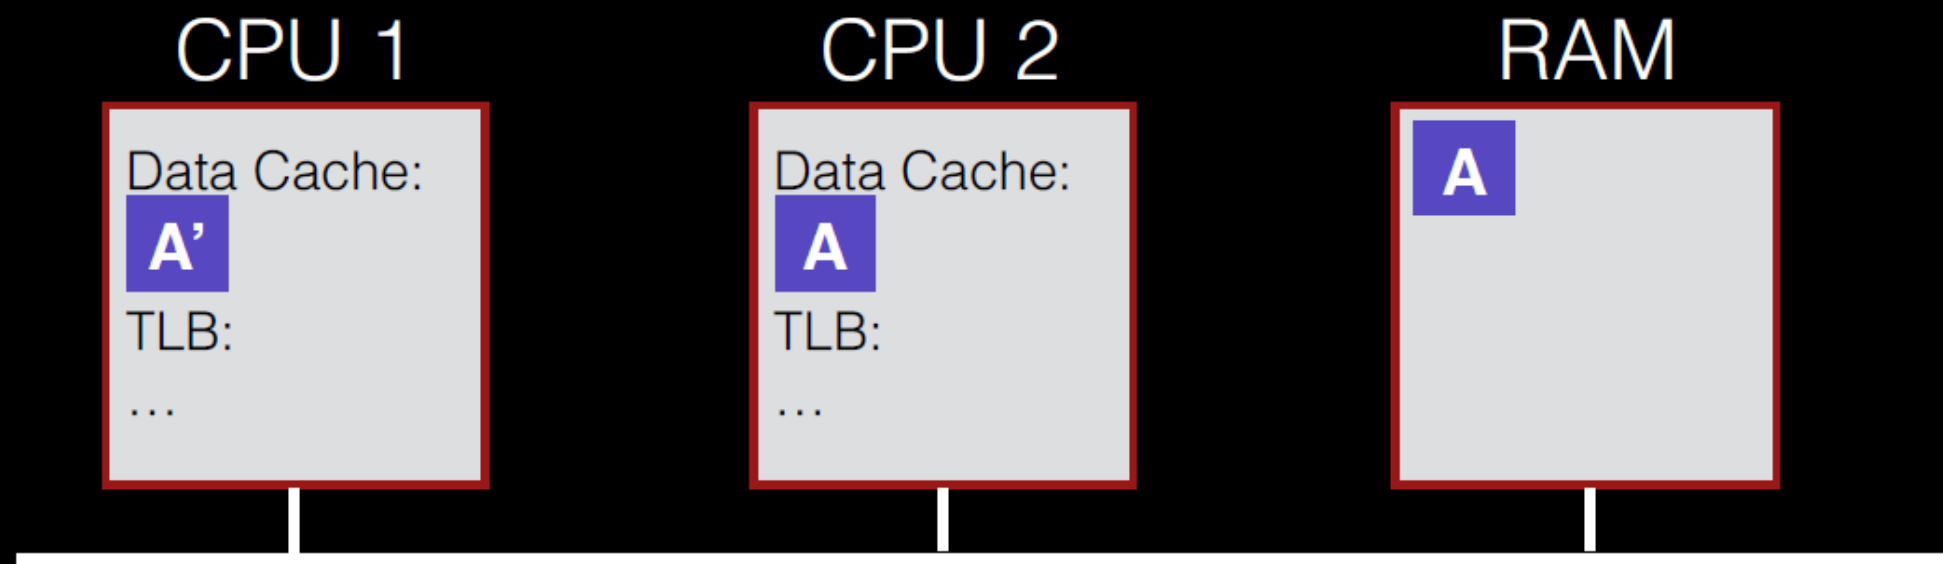
Why is switching between threads not free?

# Why is concurrency hard?

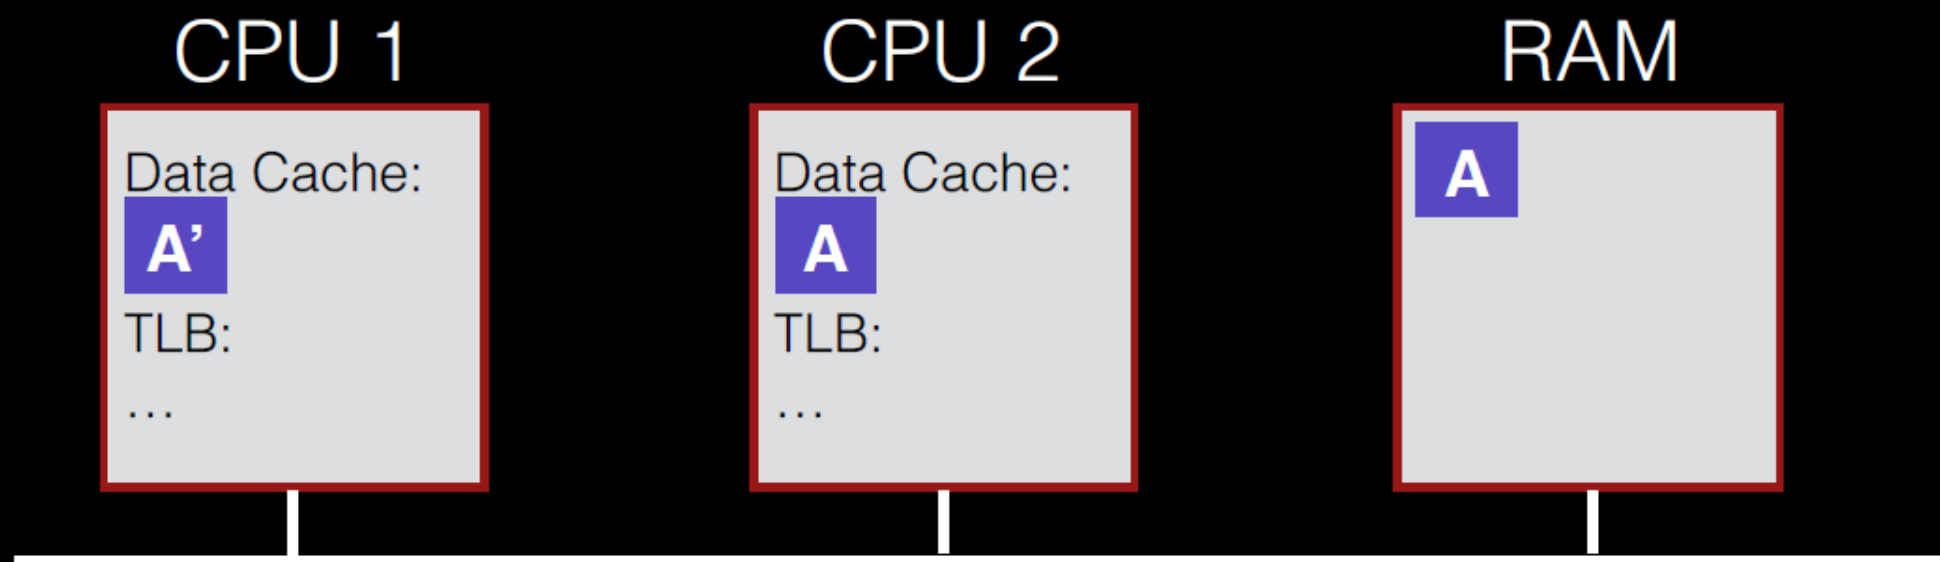
H/W caches

OS scheduler



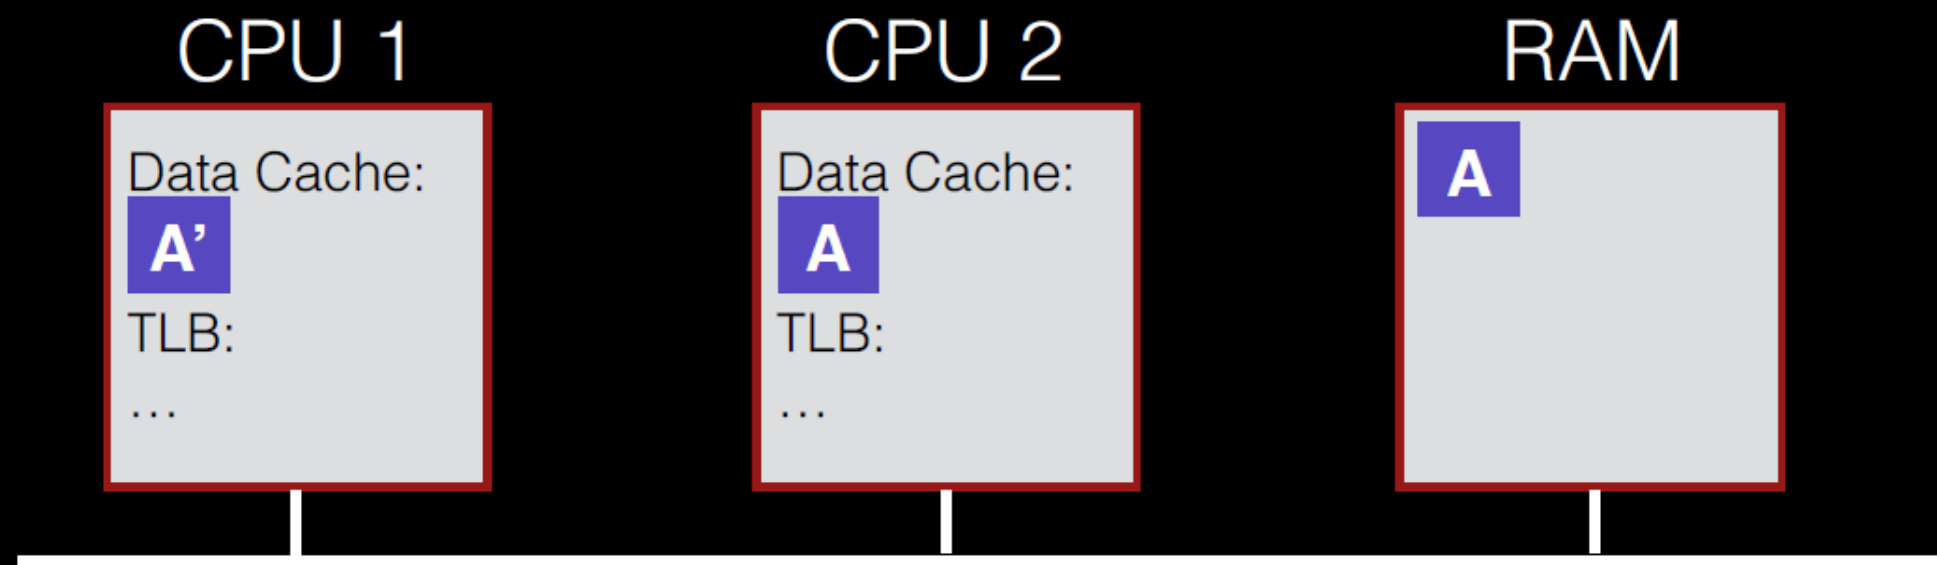


CPU 2: memory load returns **A**  
CPU 1: memory store of **A'**



CPU 2: memory load returns **A**  
CPU 1: memory store of **A'**  
CPU 2: memory load returns **A**

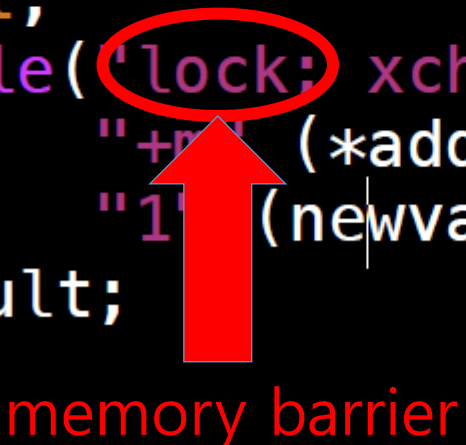




Updates from one critical section must be **visible** to others.  
CPU needs to know when to **flush caches** (or similar).

# xchg: atomic exchange, or test-and-set

```
//  
// xchg(int *addr, int newval)  
// return what is pointed to by addr  
// at the same time, store newval into addr  
//  
static inline uint  
xchg(volatile unsigned int *addr, unsigned int newval) {  
    uint result;  
    asm volatile("lock; xchgl %0, %1" :  
                 "+r" (*addr), "=a" (result) :  
                 "1" (newval) : "cc");  
    return result;  
}
```



memory barrier

# Test-and-set Spinlock

```
void SpinLock(volatile unsigned int *lock) {  
    while (xchg(lock, 1) == 1)  
        ; // spin  
  
void SpinUnlock(volatile unsigned int *lock) {  
    xchg(lock, 0);  
}
```

# Test-and-set Spinlock (optimized)

```
void SpinLock(volatile unsigned int *lock) {  
    while (xchg(lock, 1) == 1)  
        ; // spin  
  
void SpinUnlock(volatile unsigned int *lock) {  
    *lock = 0;  
}
```

Works on newer x86 processors.

Not on all CPUs (sometimes due to CPU bugs!)

# Why is concurrency hard?

H/W caches

OS scheduler

# What if multiple threads run this?

```
for (i = 0; i < max; i++) {  
    balance = balance + 1; // shared: only one  
}
```

# Balance Adder

## Thread 1

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

## Thread 2

```
mov 0x123, %eax  
add %0x1, %eax
```

```
mov %eax, 0x123
```

How much is added?

## Thread 1

```
mov 0x123, %eax (eax = 100)  
add %0x1, %eax (eax = 101)  
mov %eax, 0x123 (0x123 = 101)
```

## Thread 2

```
mov 0x123, %eax (eax = 100)  
add %0x1, %eax (eax = 101)
```

```
mov %eax, 0x123 (0x123 = 101)
```

How much is added?



## Thread 1

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

## Thread 2

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

How much is added?

## Thread 1

```
mov 0x123, %eax (eax = 101)
add %0x1, %eax (eax = 102)
mov %eax, 0x123 (0x123 = 102)
```

## Thread 2

```
mov 0x123, %eax (eax = 100)
add %0x1, %eax (eax = 101)
mov %eax, 0x123 (0x123 = 101)
```

How much is added?

## Thread 1

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

## Thread 2

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

Need **atomic sections** that don't run simultaneously, even on different CPUs!

# Review: What is needed for Correctness?

Balance = balance + 1;

Instructions accessing shared memory must execute as uninterruptable group

- Need instructions to be atomic

```
mov 0x123, %eax  
add %0x1, %eax  
mov %eax, 0x123
```

— critical section

More general:

Need **mutual exclusion** for critical sections

- if process A is in critical section C, process B can't  
(okay if other processes do unrelated work)

# Other Examples

Consider multi-threaded applications that do more than increment shared balance

Multi-threaded application with shared linked-list

- All concurrent:
  - Thread A inserting element a
  - Thread B inserting element b
  - Thread C looking up element c

# Shared Linked List

```
Void List_Insert(list_t *L,
                 int key) {
    node_t *new =
        malloc(sizeof(node_t));
    assert(new);
    new->key = key;
    new->next = L->head;
    L->head = new;
}
```

```
int List_Lookup(list_t *L,
                int key) {
    node_t *tmp = L->head;
    while (tmp) {
        if (tmp->key == key)
            return 1;
        tmp = tmp->next;
    }
    return 0;
}
```

```
typedef struct __node_t {
    int key;
    struct __node_t *next;
} node_t;
```

```
typedef struct __list_t {
    node_t *head;
} list_t;
```

```
Void List_Init(list_t *L) {
    L->head = NULL;
}
```

What can go wrong?  
Find schedule that leads to problem?

# Linked-List Race

Thread 1

Thread 2

---

`new->key = key`

`new->next = L->head`

`new->key = key`

`new->next = L->head`


`L->head = new`

`L->head = new`

**Both entries point to old head**

Only one entry (which one?) can be the new head.

# Resulting Linked List





# Locking Linked Lists

```
Void List_Insert(list_t *L,
    int key) {
    node_t *new =
        malloc(sizeof(node_t));
    assert(new);
    new->key = key;
    new->next = L->head;
    L->head = new;
}
int List_Lookup(list_t *L,
    int key) {
    node_t *tmp = L->head;
    while (tmp) {
        if (tmp->key == key)
            return 1;
        tmp = tmp->next;
    }
    return 0;
}
```

```
typedef struct __node_t {
    int key;
    struct __node_t *next;
} node_t;

typedef struct __list_t {
    node_t *head;
} list_t;

Void List_Init(list_t *L) {
    L->head = NULL;
}
```

How to add locks?

# Locking Linked Lists

```
typedef struct __node_t {  
    int key;  
    struct __node_t *next;  
} node_t;
```

```
typedef struct __list_t {  
    node_t *head;  
} list_t;
```

```
Void List_Init(list_t *L) {  
    L->head = NULL;  
}
```

How to add locks?

pthread\_mutex\_t lock;

One lock per list





```
typedef struct __node_t {  
    int key;  
    struct __node_t *next;  
} node_t;
```

```
typedef struct __list_t {  
    node_t *head;  
    pthread_mutex_t lock;  
} list_t;
```

```
Void List_Init(list_t *L) {  
    L->head = NULL;  
    pthread_mutex_init(&L->lock,  
        NULL);  
}
```

# Locking Linked Lists : Approach #1

Consider everything critical section  
Can critical section be smaller?

```
pthread_mutex_lock(&L->lock);  void List_Insert(list_t *L, int key) {  
    node_t *new =  
        malloc(sizeof(node_t));  
    assert(new);  
    new->key = key;  
    new->next = L->head;  
    L->head = new;  
}  
pthread_mutex_unlock(&L->lock);   
  
pthread_mutex_lock(&L->lock);  int List_Lookup(list_t *L, int key) {  
    node_t *tmp = L->head;  
    while (tmp) {  
        if (tmp->key == key)  
            return 1;  
        tmp = tmp->next;  
    }  
    return 0;  
}  
pthread_mutex_unlock(&L->lock);   
}
```

# Locking Linked Lists : Approach #2

Critical section small as possible

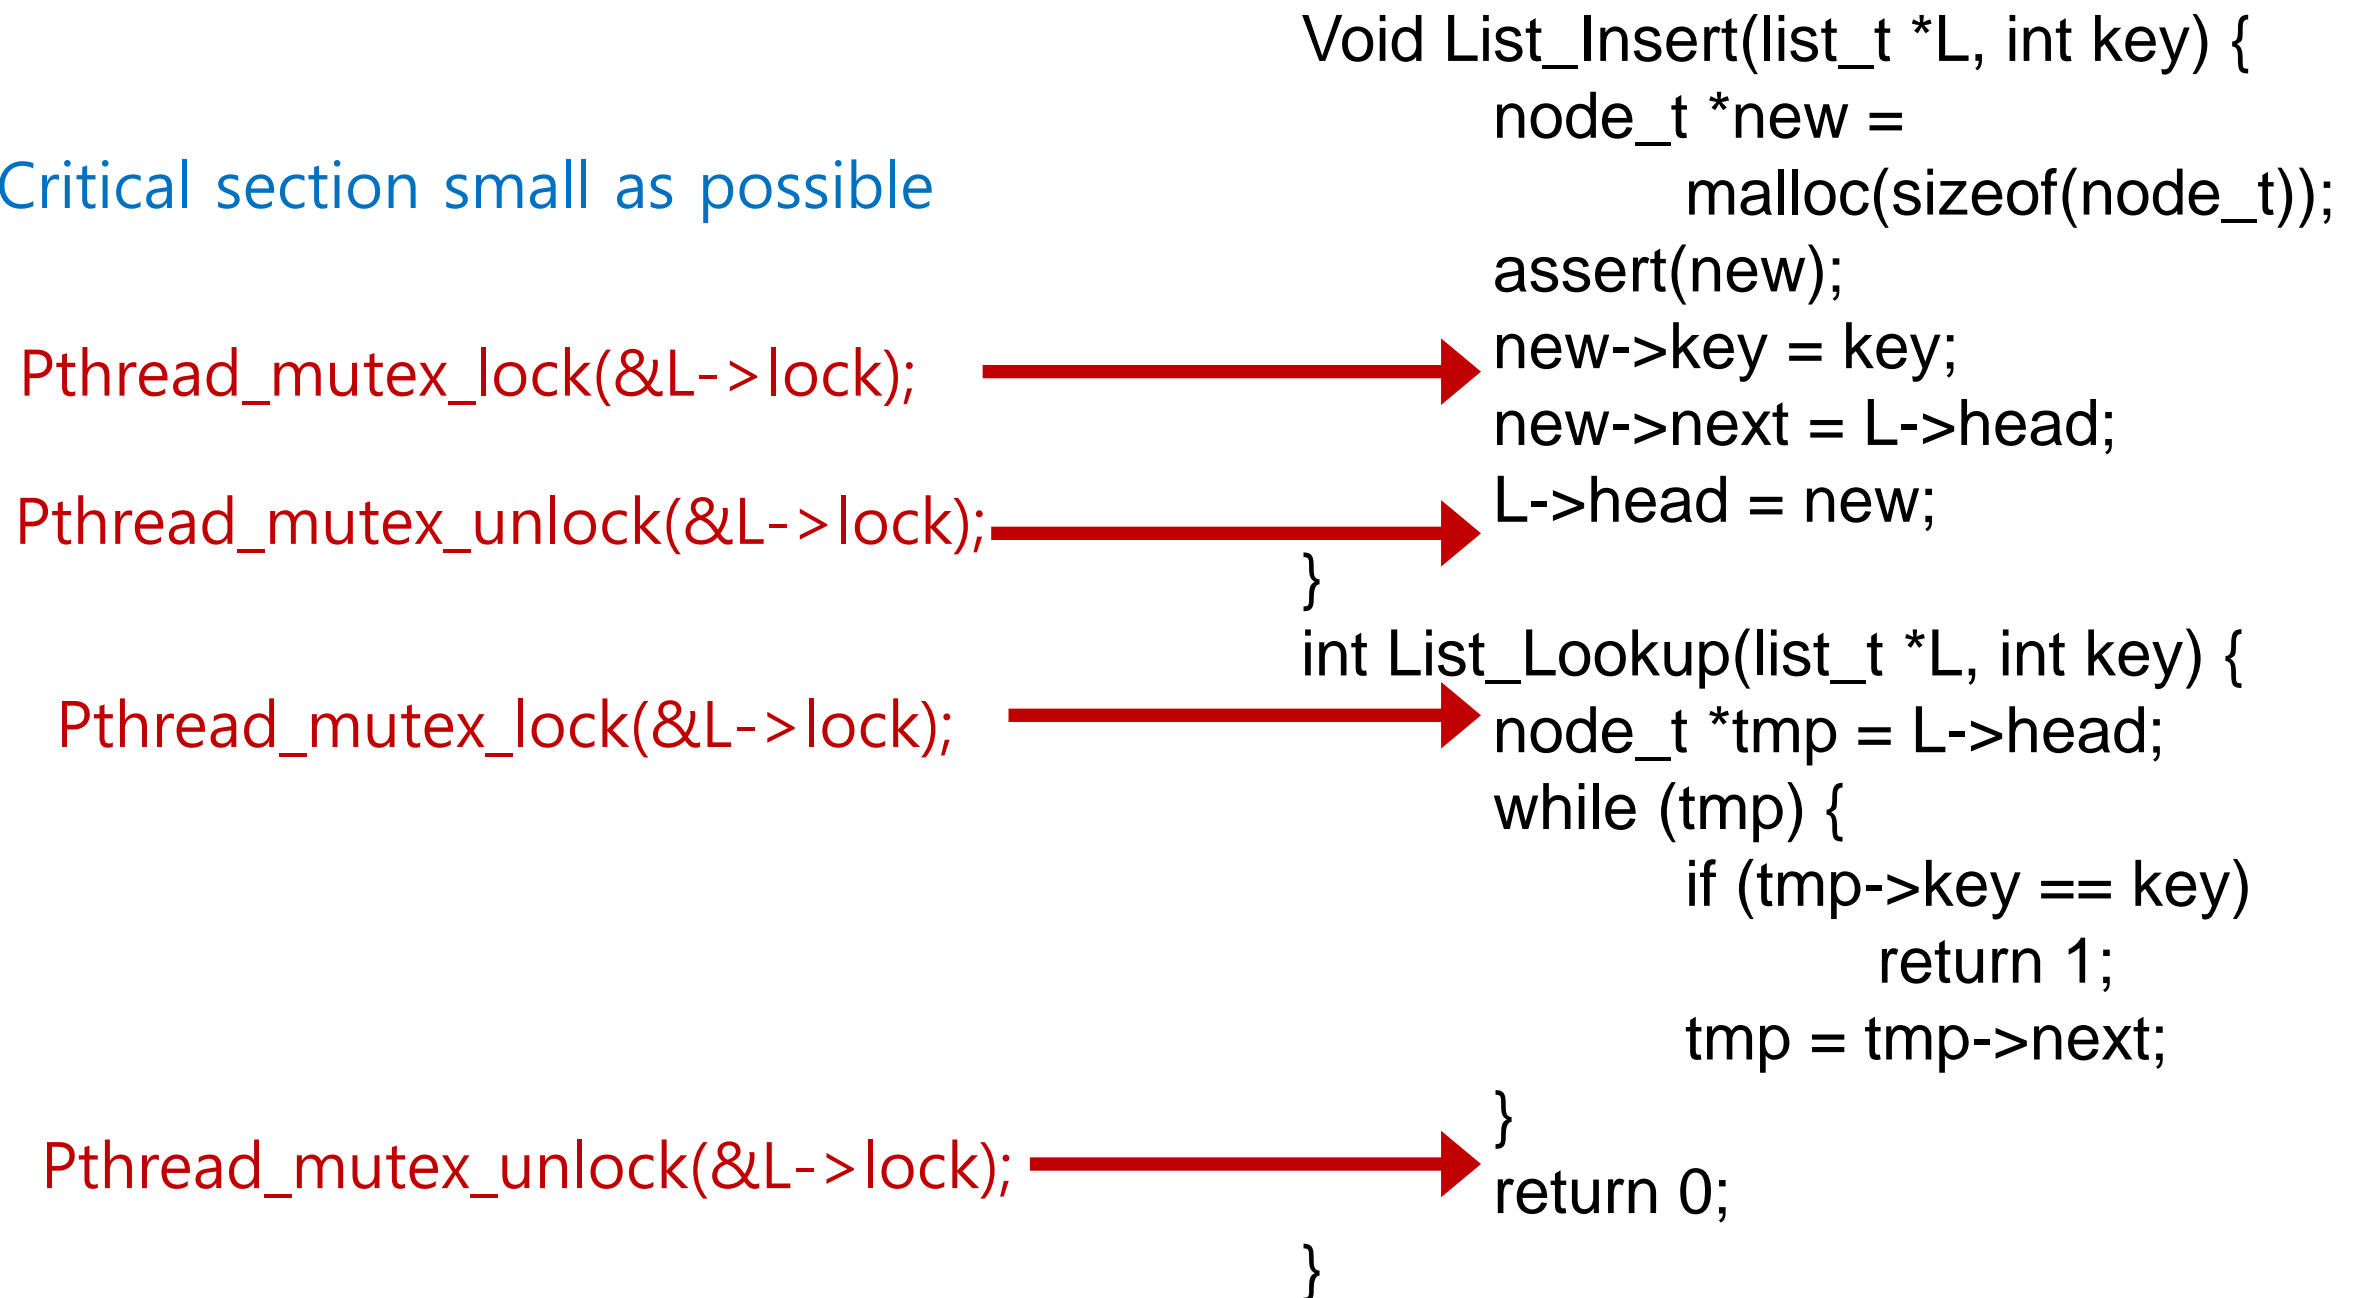
```
Void List_Insert(list_t *L, int key) {  
    node_t *new =  
        malloc(sizeof(node_t));  
    assert(new);  
    new->key = key;  
    new->next = L->head;  
    L->head = new;  
}  
int List_Lookup(list_t *L, int key) {  
    node_t *tmp = L->head;  
    while (tmp) {  
        if (tmp->key == key)  
            return 1;  
        tmp = tmp->next;  
    }  
    return 0;  
}
```

pthread\_mutex\_lock(&L->lock);

pthread\_mutex\_unlock(&L->lock);

pthread\_mutex\_lock(&L->lock);

pthread\_mutex\_unlock(&L->lock);



# Locking Linked Lists : Approach #3

What about Lookup()?

```
Void List_Insert(list_t *L, int key) {
    node_t *new =
        malloc(sizeof(node_t));
    assert(new);
    new->key = key;
    new->next = L->head;
    L->head = new;
}

int List_Lookup(list_t *L, int key) {
    node_t *tmp = L->head;
    while (tmp) {
        if (tmp->key == key)
            return 1;
        tmp = tmp->next;
    }
    return 0;
}
```

Diagram illustrating the locking of the List\_Insert and List\_Lookup functions:

- `Pthread_mutex_lock(&L->lock);` is shown with a red arrow pointing to the start of the `List_Insert` function body.
- `Pthread_mutex_unlock(&L->lock);` is shown with a red arrow pointing to the end of the `List_Insert` function body.
- `Pthread_mutex_lock(&L->lock);` is shown with a red arrow pointing to the start of the `List_Lookup` function body.
- `Pthread_mutex_unlock(&L->lock);` is shown with a red arrow pointing to the end of the `List_Lookup` function body.


If no List\_Delete(), locks not needed

# Implementing Synchronization

Build higher-level synchronization primitives in OS

- Operations that ensure correct ordering of instructions across threads

Motivation: Build them once and get them right



# Lock Implementation Goals

## Correctness

- Mutual exclusion
  - Only one thread in critical section at a time
- Progress (deadlock-free)
  - If several simultaneous requests, must allow one to proceed
- Bounded (starvation-free)
  - Must eventually allow each waiting thread to enter

## Fairness

Each thread waits for same amount of time

## Performance

CPU is not used unnecessarily (e.g., spinning)

# Implementing Synchronization

To implement, need atomic operations

**Atomic operation:** No other instructions can be interleaved

## Examples of atomic operations

- Code between interrupts on uniprocessors
  - Disable timer interrupts, don't do any I/O
- Loads and stores of words
  - Load r1, B
  - Store r1, A
- **Special hw instructions**
  - **Test&Set**
  - **Compare&Swap**



# Implementing Locks: W/ Interrupts

Turn off interrupts for critical sections

Prevent dispatcher from running another thread

Code between interrupts executes atomically

```
Void acquire(lockT *l) {  
    disableInterrupts();  
}
```

```
Void release(lockT *l) {  
    enableInterrupts();  
}
```

Disadvantages??

Only works on uniprocessors

Process can keep control of CPU for arbitrary length

Cannot perform other necessary work

# Implementing LOCKS: w/ Load+Store

Code uses a single **shared** lock variable

```
Boolean lock = false; // shared variable
```

```
Void acquire(Boolean *lock) {
```

```
    while (*lock) /* wait */ ;
```

```
    *lock = true;
```

```
}
```

```
Void release(Boolean *lock) {
```

```
    *lock = false;
```

```
}
```

Why doesn't this work?

Example schedule that fails with 2 threads?

# Race Condition with LOAD and STORE

\*lock == 0 initially

Thread 1

Thread 2

while(\*lock == 1)

while(\*lock == 1)

\*lock = 1

\*lock = 1

Both threads grab lock!

Problem: Testing lock and setting lock are not atomic

# Demo

Critical section not protected with faulty lock implementation

# Peterson's Algorithm

Assume only two threads ( $tid = 0, 1$ ) and use just loads and stores

```
int turn = 0; // shared
```

```
Boolean lock[2] = {false, false};
```

```
Void acquire() {
```

```
    lock[tid] = true;
```

```
    turn = 1-tid;
```

```
    while (lock[1-tid] && turn == 1-tid) /* wait */ ;
```

```
}
```

```
Void release() {
```

```
    lock[tid] = false;
```

```
}
```

# Different Cases: All Work

Only thread 0 wants lock

```
Lock[0] = true;
```

```
turn = 1;
```

```
while (lock[1] && turn == 1);
```

Thread 0 and thread 1 both want lock;

```
Lock[0] = true;
```

```
turn = 1;
```

```
while (lock[1] && turn == 1);
```

```
Lock[1] = true;
```

```
turn = 0;
```

```
while (lock[0] && turn == 0);
```

# Different Cases: All Work

Thread 0 and thread 1 both want lock

Lock[0] = true;

Lock[1] = true;

turn = 0;

turn = 1;

while (lock[1] && turn == 1);

while (lock[0] && turn == 0);

# Different Cases: All Work

Thread 0 and thread 1 both want lock;

Lock[0] = true;

turn = 1;

Lock[1] = true;

while (lock[1] && turn == 1);

turn = 0;

while (lock[0] && turn == 0);

while (lock[1] && turn == 1);



# Peterson's Algorithm: Intuition

Mutual exclusion: Enter critical section if and only if

- Other thread does not want to enter

- Other thread wants to enter, but your turn

Progress: Both threads cannot wait forever at while() loop

- Completes if other process does not want to enter

- Other process (matching turn) will eventually finish

Bounded waiting (not shown in examples)

- Each process waits at most one critical section

Problem: doesn't work on modern hardware  
(cache-consistency issues)

# xchg: atomic exchange, or test-and-set

```
// xchg(int *addr, int newval)  
// return what was pointed to by addr  
// at the same time, store newval into addr
```

```
int xchg(int *addr, int newval) {  
    int old = *addr;  
    *addr = newval;  
    return old;  
}
```

```
static inline uint  
xchg(volatile unsigned int *addr, unsigned int newval) {  
    uint result;  
    asm volatile("lock; xchgl %0, %1" :  
                 "+m" (*addr), "=a" (result) :  
                 "1" (newval) : "cc");  
    return result;  
}
```

# LOCK Implementation with XCHG

```
typedef struct __lock_t {  
    int flag;  
} lock_t;
```

```
void init(lock_t *lock) {  
    lock->flag = ??;  
}
```

```
void acquire(lock_t *lock) {  
    ????  
    // spin-wait (do nothing)  
}
```

```
void release(lock_t *lock) {  
    lock->flag = ??;  
}
```

```
int xchg(int *addr, int newval)
```

# XCHG Implementation

```
typedef struct __lock_t {  
    int flag;  
} lock_t;
```

```
void init(lock_t *lock) {  
    lock->flag = 0;  
}
```

```
void acquire(lock_t *lock) {  
    while(xchg(&lock->flag, 1) == 1) ;  
    // spin-wait (do nothing)  
}
```

```
void release(lock_t *lock) {  
    lock->flag = 0;  
}
```

# DEMO XCHG

Critical section protected with our lock implementation!!

# Other Atomic HW Instructions

```
int CompareAndSwap(int *addr, int expected, int new) {  
    int actual = *addr;  
    if (actual == expected)  
        *addr = new;  
    return actual;  
}
```

```
void acquire(lock_t *lock) {  
    while(CompareAndSwap(&lock->flag, ?, ?)  
        == ?) ;  
    // spin-wait (do nothing)  
}
```

# Other Atomic HW Instructions

```
int CompareAndSwap(int *addr, int expected, int new) {  
    int actual = *addr;  
    if (actual == expected)  
        *addr = new;  
    return actual;  
}
```

```
void acquire(lock_t *lock) {  
    while(CompareAndSwap(&lock->flag, 0, 1)  
        == 1) ;  
    // spin-wait (do nothing)  
}
```

# Lock Implementation Goals

## Correctness

- Mutual exclusion
  - Only one thread in critical section at a time
- Progress (deadlock-free)
  - If several simultaneous requests, must allow one to proceed
- Bounded (starvation-free)
  - Must eventually allow each waiting thread to enter

## Fairness

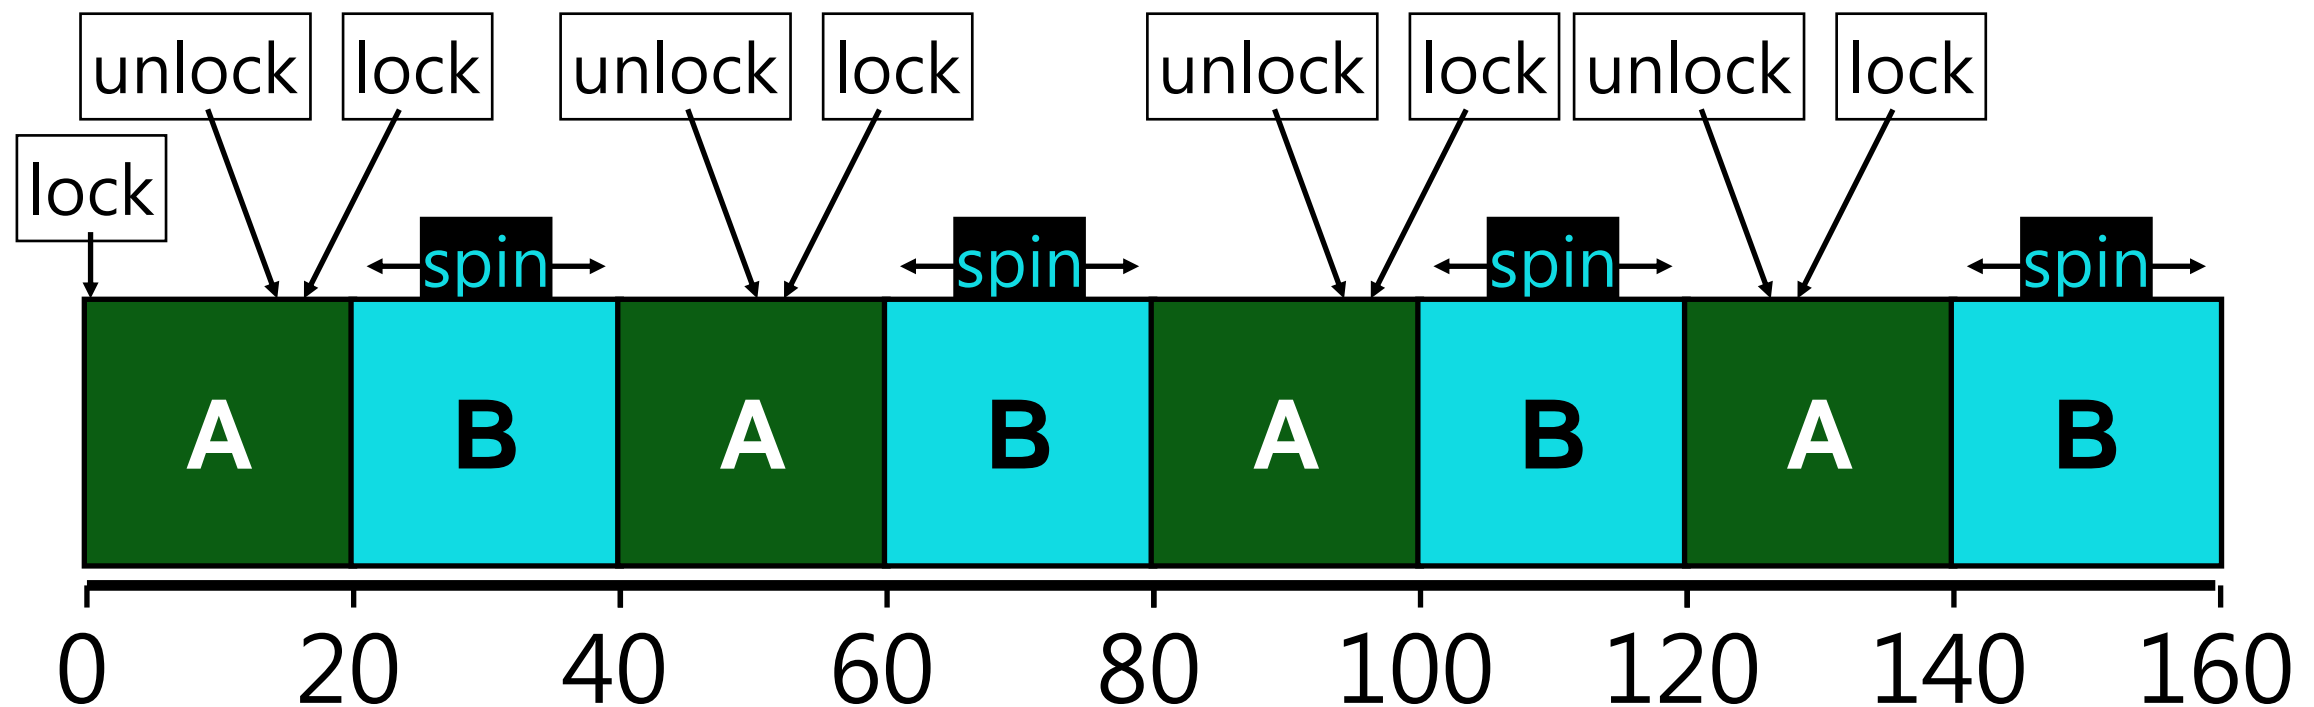
**Each thread waits for same amount of time**

## Performance

CPU is not used unnecessarily



# Basic Spinlocks are Unfair



Scheduler is independent of locks/unlocks

# Fairness: Ticket Locks

Idea: reserve each thread's turn to use a lock.

Each thread spins until their turn.


Use new atomic primitive, **fetch-and-add**:

```
int FetchAndAdd(int *ptr) {  
    int old = *ptr;  
    *ptr = old + 1;  
    return old;  
}
```


Acquire: Grab ticket;

Spin while not thread's ticket != turn

Release: Advance to next turn




A `lock()`: gets ticket 0, runs



A `lock()`: gets ticket 0, runs


B `lock()`: gets ticket 1, spins until `turn=1`




A `lock()`: gets ticket 0, runs

B `lock()`: gets ticket 1, spins until `turn=1`


C `lock()`: gets ticket 2, spins until `turn=2`



A `lock()`: gets ticket 0, runs  
B `lock()`: gets ticket 1, spins until `turn=1`  
C `lock()`: gets ticket 2, spins until `turn=2`  
A `unlock()`: `turn++`  
B runs




A **lock()**: gets ticket 0, runs  
B **lock()**: gets ticket 1, spins until turn=1  
C **lock()**: gets ticket 2, spins until turn=2  
A **unlock()**: turn++  
B runs  
A **lock()**: gets ticket 3, spins until turn=3







A **lock()**: gets ticket 0, runs  
B **lock()**: gets ticket 1, spins until turn=1  
C **lock()**: gets ticket 2, spins until turn=2  
A **unlock()**: turn++  
B runs  
A **lock()**: gets ticket 3, spins until turn=3  
B **unlock()**: turn++  
C runs




A **lock()**: gets ticket 0, runs  
B **lock()**: gets ticket 1, spins until turn=1  
C **lock()**: gets ticket 2, spins until turn=2  
A **unlock()**: turn++  
B runs  
A **lock()**: gets ticket 3, spins until turn=3  
B **unlock()**: turn++  
C runs  
C **unlock()**: turn++  
A runs



A **lock()**: gets ticket 0, runs  
B **lock()**: gets ticket 1, spins until turn=1  
C **lock()**: gets ticket 2, spins until turn=2  
A **unlock()**: turn++  
B runs  
A **lock()**: gets ticket 3, spins until turn=3  
B **unlock()**: turn++  
C runs  
C **unlock()**: turn++  
A runs  
A **unlock()**: turn++



A **lock()**: gets ticket 0, runs  
B **lock()**: gets ticket 1, spins until turn=1  
C **lock()**: gets ticket 2, spins until turn=2  
A **unlock()**: turn++  
B runs  
A **lock()**: gets ticket 3, spins until turn=3  
B **unlock()**: turn++  
C runs  
C **unlock()**: turn++  
A runs  
A **unlock()**: turn++  
C **lock()**: gets ticket 4, runs



# Ticket Lock Implementation

```
typedef struct __lock_t {  
    int ticket;  
    int turn;  
}
```

```
void lock_init(lock_t *lock) {  
    lock->ticket = 0;  
    lock->turn = 0;  
}
```

```
void acquire(lock_t *lock) {  
    int myturn = FAA(&lock->ticket);  
    while (lock->turn != myturn); // spin  
}
```

```
void release (lock_t *lock) {  
    FAA(&lock->turn);  
}
```

# Lock Implementation Goals

## Correctness

- Mutual exclusion
  - Only one thread in critical section at a time
- Progress (deadlock-free)
  - If several simultaneous requests, must allow one to proceed
- Bounded (starvation-free)
  - Must eventually allow each waiting thread to enter

## Fairness

Each thread waits for same amount of time

# Spinlock Performance

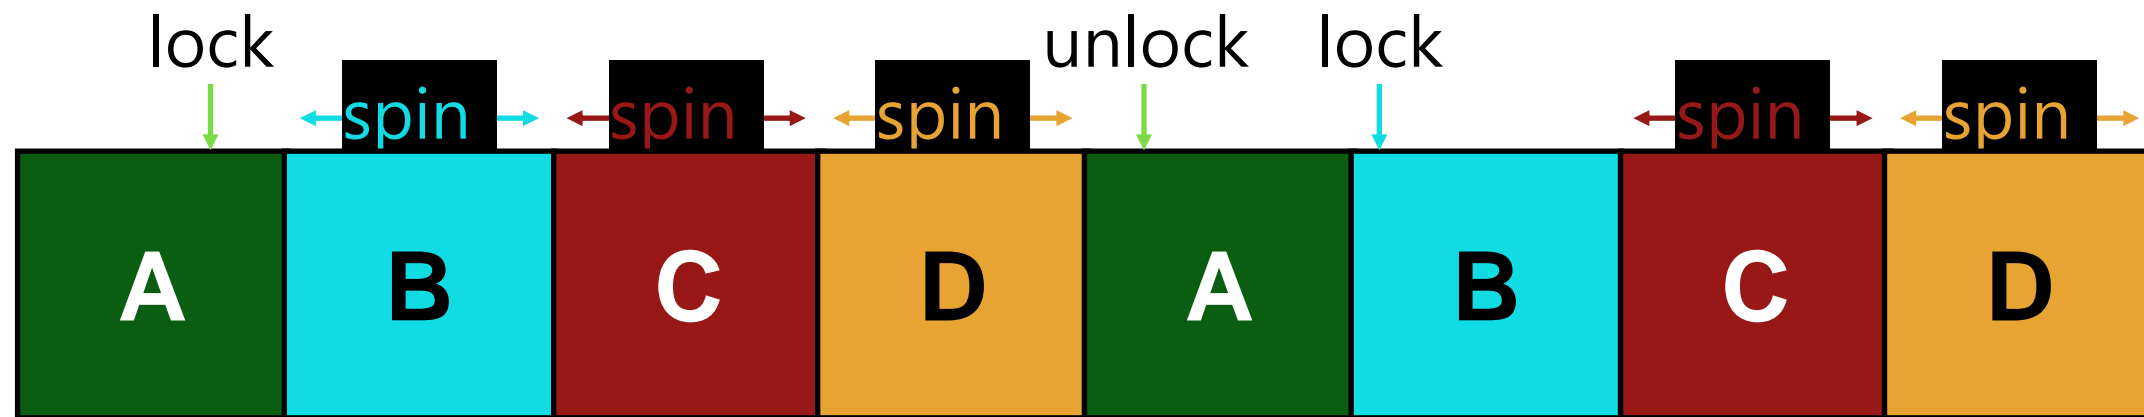
## Fast when...

- many CPUs
- locks held a short time
- advantage: avoid context switch

## Slow when...

- one CPU
- locks held a long time
- disadvantage: spinning is wasteful

# CPU Scheduler is Ignorant



CPU scheduler may run **B** instead of **A**  
even though **B** is waiting for **A**



# Test-and-set Spinlock

```
void SpinLock(volatile unsigned int *lock) {  
    while (xchg(lock, 1) == 1)  
        ; // spin  
  
void SpinUnlock(volatile unsigned int *lock) {  
    *lock = 0;  
}
```

```
void SpinLock(volatile unsigned int *lock) {  
    while (xchg(lock, 1) == 1)  
        yield(); // spin  
  
void SpinUnlock(volatile unsigned int *lock) {  
    *lock = 0;  
}
```

**Pro:** we won't waste cycles on spin now

**Con:** we may have to context switch many times to get the right thread

# Queue Locks

Idea: put threads on **queue**.

Tell kernel **don't schedule** queued threads.

Upon unlock, tell kernel it can run thread(s) again.

**Hybrid approach**: spin a while, then queue self  
- called "two-phase locks"

# In-Kernel locking

Sometimes interrupt handlers have **no context!**

Queue locks cannot work. Why?

Approach: cooperative scheduling.  
- **use spin locks, disable interrupts**

# Ticket Lock with Yield()

```
typedef struct __lock_t {  
    int ticket;  
    int turn;  
}
```

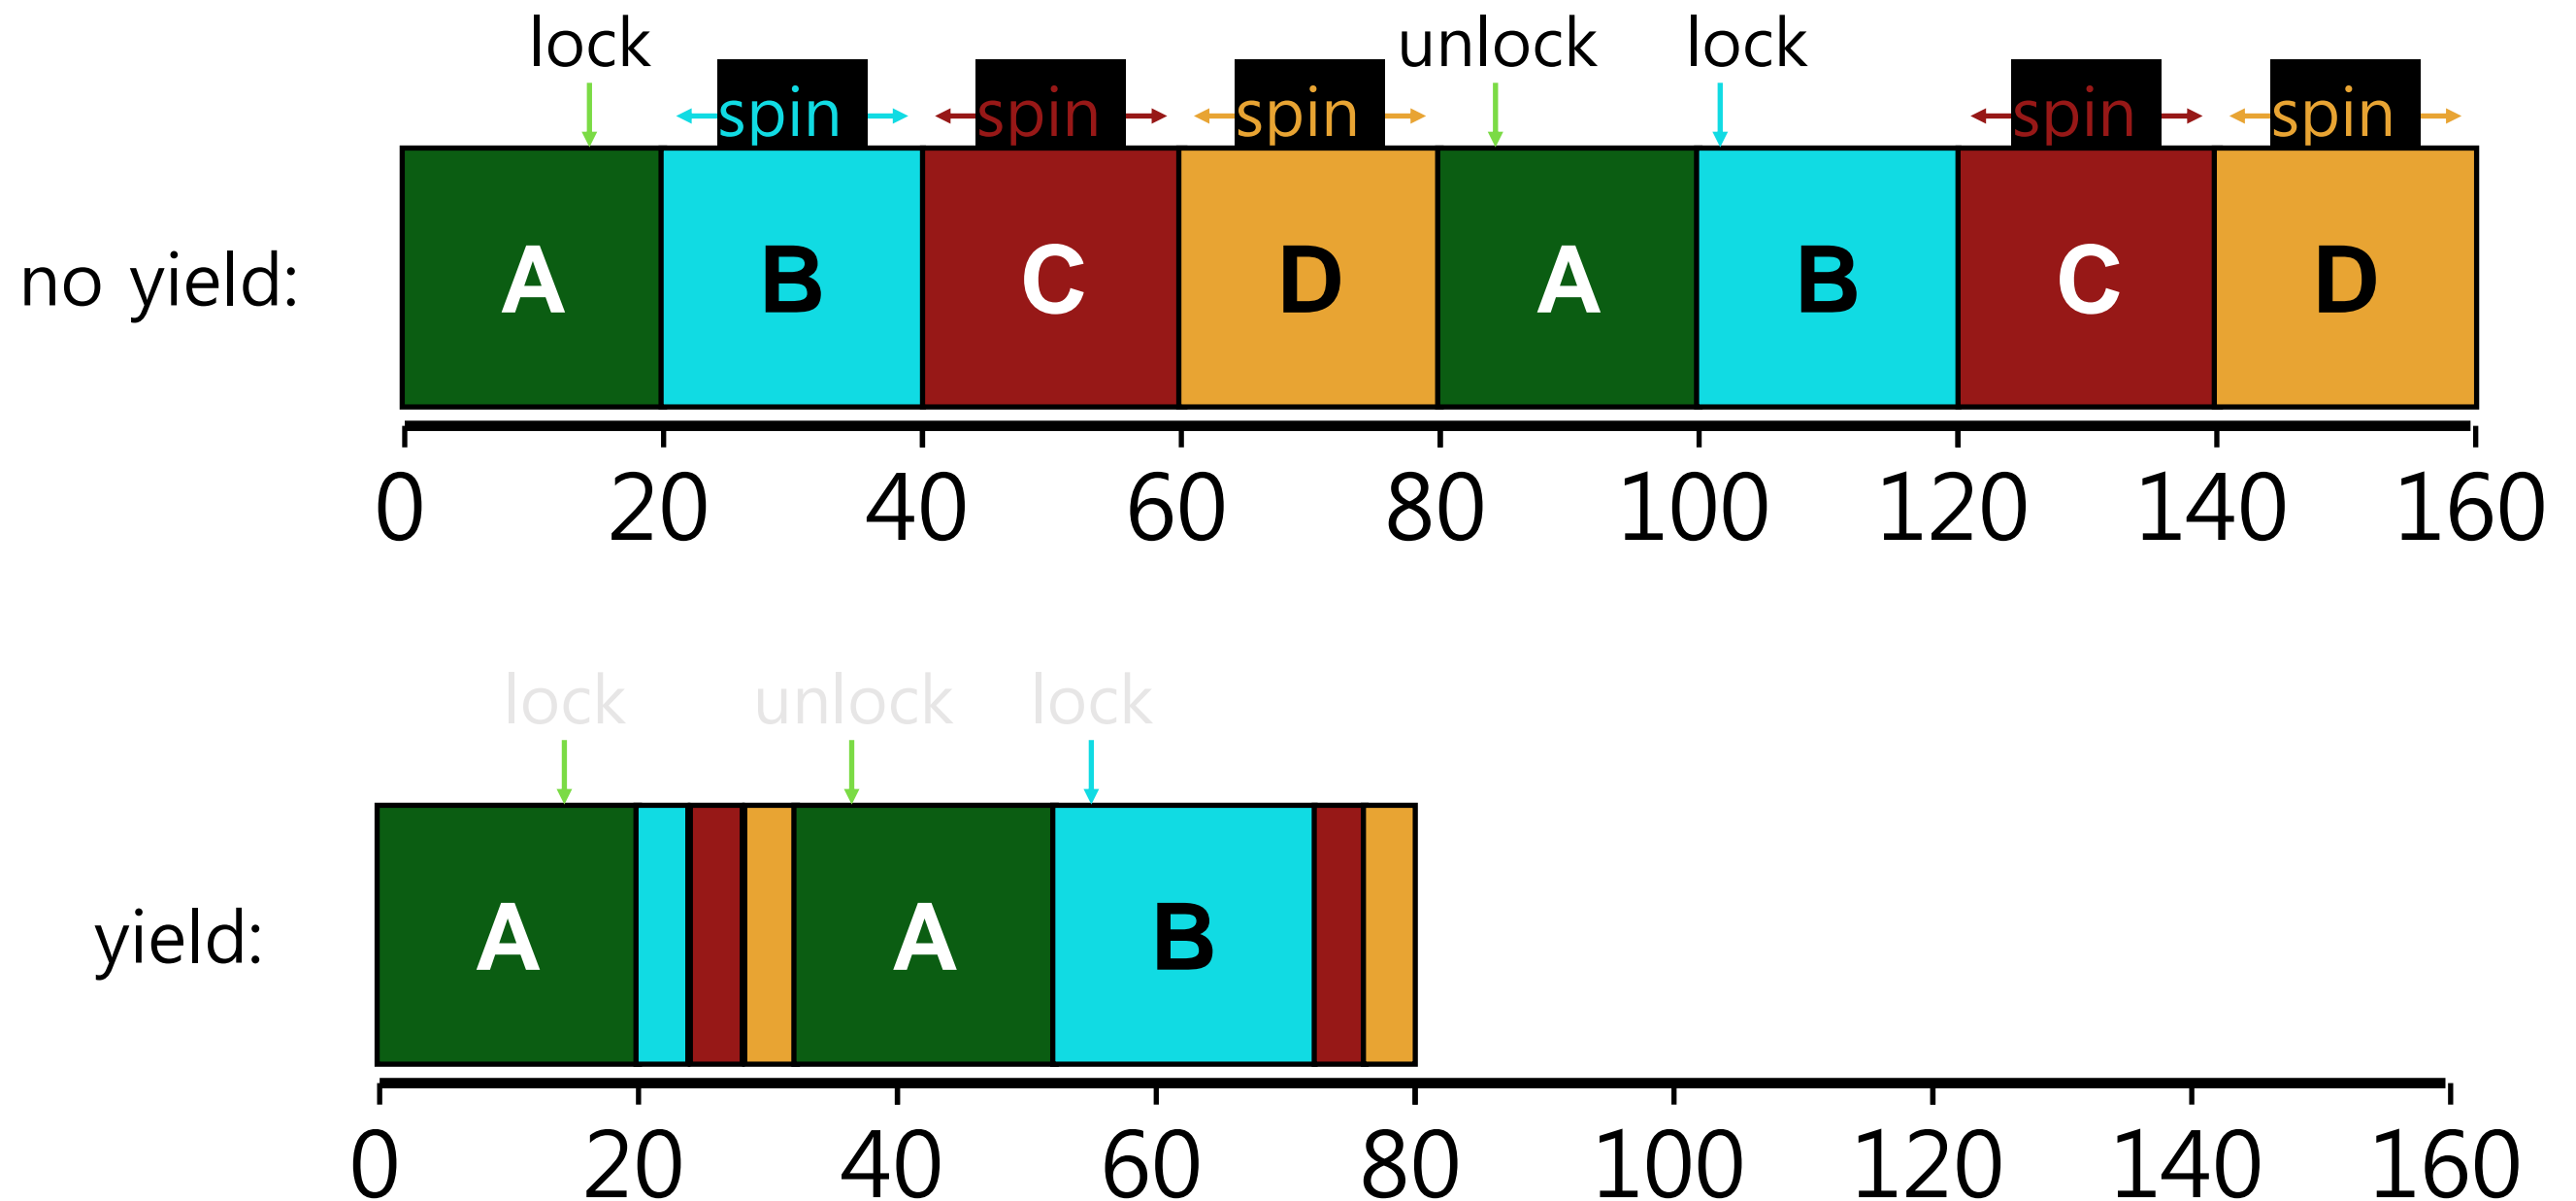
```
void lock_init(lock_t *lock) {  
    lock->ticket = 0;  
    lock->turn = 0;  
}
```

```
void acquire(lock_t *lock) {  
    int myturn = FAA(&lock->ticket);  
    while(lock->turn != myturn)
```

**yield();**

```
}  
  
void release (lock_t *lock) {  
    FAA(&lock->turn);  
}
```

# Yield Instead of Spin



# Spinlock Performance

Waste...

Without yield:  $O(\text{threads} * \text{time\_slice})$

With yield:  $O(\text{threads} * \text{context\_switch})$

So even with yield, spinning is slow with high thread contention

Next improvement: Block and put thread on waiting queue instead of spinning