

OSTEP

Concurrency

Semaphores

Questions answered in this lecture:

Review: How to implement join with condition variables?

Review: How to implement producer/consumer with condition variables?

What is the difference between **semaphores** and condition variables?

How to implement a **lock** with semaphores?

How to implement semaphores with locks and condition variables?

How to implement **join** and producer/consumer with semaphores?

How to implement **reader/writer locks** with semaphores?

Concurrency Objectives

Mutual exclusion (e.g., A and B don't run at same time)

- solved with *locks*

Ordering (e.g., B runs after A does something)

- solved with *condition variables* and *semaphores*

Condition Variables

wait(cond_t *cv, mutex_t *lock)

- assumes the lock is held when wait() is called
- puts caller to sleep + releases the lock (atomically)
- when awoken, reacquires lock before returning

signal(cond_t *cv)

- wake a single waiting thread (if ≥ 1 thread is waiting)
- if there is no waiting thread, just return, doing nothing

broadcast(cond_t *cv)

- wake all waiting threads (if ≥ 1 thread is waiting)
- if there are no waiting thread, just return, doing nothing

Join Implementation: Correct

Parent:

```
void thread_join() {
    Mutex_lock(&m);           // w
    if (done == 0)           // x
        Cond_wait(&c, &m); // y
    Mutex_unlock(&m);         // z
}
```

Child:

```
void thread_exit() {
    Mutex_lock(&m);           // a
    done = 1;                 // b
    Cond_signal(&c);           // c
    Mutex_unlock(&m);          // d
}
```

Parent: w x y z

Child:

	a	b	c
--	---	---	---

Use mutex to ensure no race between interacting with state and wait/signal

Producer/Consumer Problem

Producers generate data (like pipe writers)

Consumers grab data and process it (like pipe readers)

Producer/consumer problems are frequent in systems.

Use condition variables to:


make producers wait when buffers are full

make consumers wait when there is nothing to consume

Broken Implementation of Producer Consumer

```
void *producer(void *arg) {
    for (int i=0; i<loops; i++) {
        Mutex_lock(&m); // p1
        while(numfull == max) //p2
            Cond_wait(&cond, &m); //p3
        do_fill(i); // p4
        Cond_signal(&cond); //p5
        Mutex_unlock(&m); //p6
    }
}
```

```
void *consumer(void *arg) {
    while(1) {
        Mutex_lock(&m); // c1
        while(numfull == 0) // c2
            Cond_wait(&cond, &m); // c3
        int tmp = do_get(); // c4
        Cond_signal(&cond); // c5
        Mutex_unlock(&m); // c6
        printf("%d\n", tmp); // c7
    }
}
```



does last signal wake **producer** or **consumer2**?

Producer/Consumer: Two CVs

```

void *producer(void *arg) {
    for (int i = 0; i < loops; i++) {
        Mutex_lock(&m); // p1
        if (numfull == max) // p2
            Cond_wait(&empty, &m); // p3
        do_fill(i); // p4
        Cond_signal(&fill); // p5
        Mutex_unlock(&m); //p6
    }
}

void *consumer(void *arg) {
    while (1) {
        Mutex_lock(&m); // c1
        if (numfull == 0) // c2
            Cond_wait(&fill, &m); // c3
        int tmp = do_get(); // c4
        Cond_signal(&empty); // c5
        Mutex_unlock(&m); // c6
    }
}

```

Is this correct? Can you find a bad schedule?

1. consumer1 waits because numfull == 0
2. producer increments numfull, wakes consumer1
3. before consumer1 runs, consumer2 runs, grabs entry, sets numfull=0.
4. consumer2 then reads bad data.

Producer: p1 p2 p4 p5 p6

Consumer1: c1 c2 c3

Consumer2: c1 c2 c4 c5 c6

c4! ERROR

CV Rule of Thumb 3

Whenever a lock is acquired, recheck assumptions about state!
Use "while" instead of "if"

Possible for another thread to grab lock between signal and wakeup from wait

- Difference between Mesa (practical implementation) and Hoare (theoretical) semantics
- Signal() simply makes a thread runnable, does not guarantee that thread runs next

Note that some libraries also have "spurious wakeups"

- May wake multiple waiting threads at signal or at any time

Producer/Consumer: Two CVs and WHILE

```
void *producer(void *arg) {
    for (int i = 0; i < loops; i++) {
        Mutex_lock(&m); // p1
        while (numfull == max) // p2
            Cond_wait(&empty, &m); // p3
        do_fill(i); // p4
        Cond_signal(&fill); // p5
        Mutex_unlock(&m); //p6
    }
}

void *consumer(void *arg) {
    while (1) {
        Mutex_lock(&m);
        while (numfull == 0)
            Cond_wait(&fill, &m);
        int tmp = do_get();
        Cond_signal(&empty);
        Mutex_unlock(&m);
    }
}
```

Is this correct? Can you find a bad schedule?

Correct!

- no concurrent access to shared state
- every time lock is acquired, assumptions are reevaluated
- a consumer will get to run after every do_fill()
- a producer will get to run after every do_get()

Summary: rules of thumb for CVs

Keep **state** in addition to CV's

Always do wait/signal with **lock held**

Whenever thread wakes from waiting, **recheck state**

Semaphore: A definition

An object **with an integer value**

- We can manipulate with two routines; `sem_wait()` and `sem_post()`.
- Initialization

```
1  #include <semaphore.h>
2  sem_t s;
3  sem_init(&s, 0, 1); // initialize s to the value 1
```

- Declare a semaphore `s` and initialize it to the value 1
- The second argument, 0, indicates that the semaphore is shared between *threads in the same process*.

Semaphore: Interact with semaphore

- `sem_wait()`

```
1  int sem_wait(sem_t *s) {  
2      decrement the value of semaphore s by one  
3      wait if value of semaphore s is negative  
4  }
```

- If the value of the semaphore was *one* or *higher* when called `sem_wait()`, **return right away**.
- It will cause the caller to suspend execution waiting for a subsequent post.
- When negative, the value of the semaphore is equal to the number of waiting threads.

Semaphore: Interact with semaphore (Cont.)

- `sem_post()`

```
1  int sem_post(sem_t *s) {  
2      increment the value of semaphore s by one  
3      if there are one or more threads waiting, wake one  
4  }
```

- Simply **increments** the value of the semaphore.
- If there is a thread waiting to be woken, **wakes** one of them up.

Binary Semaphores (Locks)

- What should **x** be?
 - The initial value should be **1**.

```
1  sem_t m;  
2  sem_init(&m, 0, X); // initialize semaphore to X; what should X be?  
3  
4  sem_wait(&m);  
5  //critical section here  
6  sem_post(&m);
```

Thread Trace: Single Thread Using A Semaphore

Value of Semaphore	Thread 0	Thread 1
1		
1	call sema_wait()	
0	sem_wait() returns	
0	(crit sect)	
0	call sem_post()	
1	sem_post() returns	

Thread Trace: Two Threads Using A Semaphore

Value	Thread 0	State	Thread 1	State
1		Running		Ready
1	call sem_wait()	Running		Ready
0	sem_wait() retruns	Running		Ready
0	(crit set: begin)	Running		Ready
0	<i>Interrupt; Switch → T1</i>	Ready		Running
0		Ready	call sem_wait()	Running
-1		Ready	decrement sem	Running
-1		Ready	(sem < 0) → sleep	sleeping
-1		Running	<i>Switch → T0</i>	sleeping
-1	(crit sect: end)	Running		sleeping
-1	call sem_post()	Running		sleeping
0	increment sem	Running		sleeping
0	wake(T1)	Running		Ready
0	sem_post() returns	Running		Ready
0	<i>Interrupt; Switch → T1</i>	Ready		Running
0		Ready	sem_wait() retruns	Running
0		Ready	(crit sect)	Running
0		Ready	call sem_post()	Running
1		Ready	sem_post() returns	Running

Semaphores As Condition Variables

```
1  sem_t s;
2
3  void *
4  child(void *arg) {
5      printf("child\n");
6      sem_post(&s); // signal here: child is done
7      return NULL;
8  }
9
10 int
11 main(int argc, char *argv[]) {
12     sem_init(&s, 0, X); // what should X be?
13     printf("parent: begin\n");
14     pthread_t c;
15     pthread_create(c, NULL, child, NULL);
16     sem_wait(&s); // wait here for child
17     printf("parent: end\n");
18     return 0;
19 }
```

A Parent Waiting For Its Child

```
parent: begin
child
parent: end
```

The execution result

- What should **x** be?
 - The value of semaphore should be set to is **0**.

Thread Trace: Parent Waiting For Child (Case 1)

- The parent call `sem_wait()` before the child has called `sem_post()`.

Value	Parent	State	Child	State
0	Create(Child)	Running	<i>(Child exists; is runnable)</i>	Ready
0	call <code>sem_wait()</code>	Running		Ready
-1	decrement sem	Running		Ready
-1	$(sem < 0) \rightarrow \text{sleep}$	sleeping		Ready
-1	<i>Switch</i> →Child	sleeping	child runs	Running
-1		sleeping	call <code>sem_post()</code>	Running
0		sleeping	increment sem	Running
0		Ready	wake(Parent)	Running
0		Ready	<code>sem_post()</code> returns	Running
0		Ready	<i>Interrupt; Switch</i> →Parent	Ready
0	<code>sem_wait()</code> retruns	Running		Ready

Thread Trace: Parent Waiting For Child (Case 2)

- The child runs to completion before the parent call `sem_wait()`.

Value	Parent	State	Child	State
0	Create(Child)	Running	(Child exists; is runnable)	Ready
0	<i>Interrupt; switch→Child</i>	Ready	child runs	Running
0		Ready	call <code>sem_post()</code>	Running
1		Ready	increment sem	Running
1		Ready	wake(nobody)	Running
1		Ready	<code>sem_post()</code> returns	Running
1	parent runs	Running	<i>Interrupt; Switch→Parent</i>	Ready
1	call <code>sem_wait()</code>	Running		Ready
0	decrement sem	Running		Ready
0	(sem<0)→awake	Running		Ready
0	<code>sem_wait()</code> retruns	Running		Ready

The Producer/Consumer (Bounded-Buffer) Problem

- **Producer:** `put()` interface
 - Wait for a buffer to become *empty* in order to put data into it.
- **Consumer:** `get()` interface
 - Wait for a buffer to become *filled* before using it.

```
1  int buffer[MAX];
2  int fill = 0;
3  int use = 0;
4
5  void put(int value) {
6      buffer[fill] = value;    // line f1
7      fill = (fill + 1) % MAX; // line f2
8  }
9
10 int get() {
11     int tmp = buffer[use];    // line g1
12     use = (use + 1) % MAX;    // line g2
13     return tmp;
14 }
```

The Producer/Consumer (Bounded-Buffer) Problem

```
1  sem_t empty;
2  sem_t full;
3
4  void *producer(void *arg) {
5      int i;
6      for (i = 0; i < loops; i++) {
7          sem_wait(&empty);          // line P1
8          put(i);                    // line P2
9          sem_post(&full);           // line P3
10     }
11 }
12
13 void *consumer(void *arg) {
14     int i, tmp = 0;
15     while (tmp != -1) {
16         sem_wait(&full);            // line C1
17         tmp = get();                // line C2
18         sem_post(&empty);           // line C3
19         printf("%d\n", tmp);
20     }
21 }
22 ...
```

A Solution: Adding Mutual Exclusion

```
1  sem_t empty;
2  sem_t full;
3  sem_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          sem_wait(&mutex);           // line p0 (NEW LINE)
9          sem_wait(&empty);           // line p1
10         put(i);                     // line p2
11         sem_post(&full);             // line p3
12         sem_post(&mutex);           // line p4 (NEW LINE)
13     }
14 }
15
(Cont.)
```

Adding Mutual Exclusion (Incorrectly)

A Solution: Adding Mutual Exclusion

```
(Cont.)
16  void *consumer(void *arg) {
17      int i;
18      for (i = 0; i < loops; i++) {
19          sem_wait(&mutex);           // line c0 (NEW LINE)
20          sem_wait(&full);            // line c1
21          int tmp = get();             // line c2
22          sem_post(&empty);           // line c3
23          sem_post(&mutex);           // line c4 (NEW LINE)
24          printf("%d\n", tmp);
25      }
26 }
```

Adding Mutual Exclusion (Incorrectly)

A Solution: Adding Mutual Exclusion (Cont.)

- Imagine two thread: one producer and one consumer.
 - The consumer **acquire** the `mutex` (line c0).
 - The consumer **calls** `sem_wait()` on the full semaphore (line c1).
 - The consumer is **blocked** and **yield** the CPU.
 - The consumer still holds the mutex!
 - The producer **calls** `sem_wait()` on the binary `mutex` semaphore (line p0).
 - The producer is now **stuck** waiting too. **a classic deadlock.**

Finally, A Working Solution

```
1  sem_t empty;
2  sem_t full;
3  sem_t mutex;
4
5  void *producer(void *arg) {
6      int i;
7      for (i = 0; i < loops; i++) {
8          sem_wait(&empty);           // line p1
9          sem_wait(&mutex);           // line p1.5 (MOVED MUTEX HERE...)
10         put(i);                     // line p2
11         sem_post(&mutex);           // line p2.5 (... AND HERE)
12         sem_post(&full);            // line p3
13     }
14 }
15
(Cont.)
```

Adding Mutual Exclusion (Correctly)

Finally, A Working Solution

```
(Cont.)
16  void *consumer(void *arg) {
17      int i;
18      for (i = 0; i < loops; i++) {
19          sem_wait(&full);           // line c1
20          sem_wait(&mutex);          // line c1.5 (MOVED MUTEX HERE...)
21          int tmp = get();           // line c2
22          sem_post(&mutex);          // line c2.5 (... AND HERE)
23          sem_post(&empty);          // line c3
24          printf("%d\n", tmp);
25      }
26  }
27
28  int main(int argc, char *argv[]) {
29      // ...
30      sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with ...
31      sem_init(&full, 0, 0);    // ... and 0 are full
32      sem_init(&mutex, 0, 1);   // mutex=1 because it is a lock
33      // ...
34  }
```

Adding Mutual Exclusion (Correctly)

Reader-Writer Locks

- Imagine a number of concurrent list operations, including **inserts** and simple **lookups**.
 - **insert:**
 - Change the state of the list
 - A traditional critical section makes sense.
 - **lookup:**
 - Simply *read* the data structure.
 - As long as we can guarantee that no insert is on-going, we can allow many lookups to proceed **concurrently**.

This special type of lock is known as a **reader-write lock**.

Reader-Writer

- Only a **single writer** can acquire the lock.
- Once a reader has acquired a **read lock**,
 - **More readers** will be allowed to acquire the read lock too.
 - A writer will have to wait until all readers are finished.

```
1  typedef struct _rwlock_t {
2      sem_t lock;           // binary semaphore (basic lock)
3      sem_t writelock;      // used to allow ONE writer or MANY readers
4      int readers;          // count of readers reading in critical section
5  } rwlock_t;
6
7  void rwlock_init(rwlock_t *rw) {
8      rw->readers = 0;
9      sem_init(&rw->lock, 0, 1);
10     sem_init(&rw->writelock, 0, 1);
11 }
12
13 void rwlock_acquire_readlock(rwlock_t *rw) {
14     sem_wait(&rw->lock);
15     ...
```

A Reader-Writer Locks (Cont.)


```
15     rw->readers++;
16     if (rw->readers == 1)
17         sem_wait(&rw->writelock); // first reader acquires writelock
18     sem_post(&rw->lock);
19 }
20
21 void rwlock_release_readlock(rwlock_t *rw) {
22     sem_wait(&rw->lock);
23     rw->readers--;
24     if (rw->readers == 0)
25         sem_post(&rw->writelock); // last reader releases writelock
26     sem_post(&rw->lock);
27 }
28
29 void rwlock_acquire_writelock(rwlock_t *rw) {
30     sem_wait(&rw->writelock);
31 }
32
33 void rwlock_release_writelock(rwlock_t *rw) {
34     sem_post(&rw->writelock);
35 }
```

A Reader-Writer Locks (Cont.)

- The reader-writer locks have **fairness problem**.
 - It would be relatively easy for reader to **starve writer**.
 - How to prevent more readers from entering the lock once a writer is waiting?
- 숙제 (due 10/26 hard copy로 수업시간에 낼것, 1장으로)

The Dining Philosophers

- Assume there are five “**philosophers**” sitting around a table.
 - Between each pair of philosophers is a single fork (five total).
 - The philosophers each have times where they **think**, and don't need any forks, and times where they **eat**.
 - In order to *eat*, a philosopher needs **two forks**, both the one on their *left* and the one on their *right*.
 - **The contention for these forks.**



The Dining Philosophers (Cont.)

- Key challenge
 - There is **no deadlock**.
 - **No** philosopher **starves** and never gets to eat.
 - **Concurrency** is high.

```
while (1) {  
    think();  
    getforks();  
    eat();  
    putforks();  
}
```

Basic loop of each philosopher

```
// helper functions  
int left(int p) { return p; }  
  
int right(int p) {  
    return (p + 1) % 5;  
}
```

Helper functions (Downey's solutions)

- Philosopher p wishes to refer to the fork on their left \rightarrow call `left(p)`.
- Philosopher p wishes to refer to the fork on their right \rightarrow call `right(p)`.

The Dining Philosophers (Cont.)

- We need some **semaphore**, one for each fork: `sem_t forks[5]`.

```
1  void getforks() {
2      sem_wait(forks[left(p)]);
3      sem_wait(forks[right(p)]);
4  }
5
6  void putforks() {
7      sem_post(forks[left(p)]);
8      sem_post(forks[right(p)]);
9  }
```

The `getforks()` and `putforks()` Routines (Broken Solution)

- **Deadlock** occur!
 - If each philosopher happens to **grab the fork on their left** before any philosopher can grab the fork on their right.
 - Each will be stuck *holding one fork* and waiting for another, *forever*.

A Solution: Breaking The Dependency

- Change **how forks are acquired**.
 - Let's assume that philosopher 4 acquire the forks in a *different order*.

```
1  void getforks() {  
2      if (p == 4) {  
3          sem_wait(forks[right(p)]);  
4          sem_wait(forks[left(p)]);  
5      } else {  
6          sem_wait(forks[left(p)]);  
7          sem_wait(forks[right(p)]);  
8      }  
9  }
```

- There is no situation where each philosopher grabs one fork and is stuck waiting for another. **The cycle of waiting is broken.**

How To Implement Semaphores

- Build our own version of semaphores called **Zemaphores**

```
1  typedef struct __Zem_t {
2      int value;
3      pthread_cond_t cond;
4      pthread_mutex_t lock;
5  } Zem_t;
6
7  // only one thread can call this
8  void Zem_init(Zem_t *s, int value) {
9      s->value = value;
10     Cond_init(&s->cond);
11     Mutex_init(&s->lock);
12 }
13
14 void Zem_wait(Zem_t *s) {
15     Mutex_lock(&s->lock);
16     while (s->value <= 0)
17         Cond_wait(&s->cond, &s->lock);
18     s->value--;
19     Mutex_unlock(&s->lock);
20 }
21 ...
```

How To Implement Semaphores (Cont.)

```
22 void Zem_post(Zem_t *s) {  
23     Mutex_lock(&s->lock);  
24     s->value++;  
25     Cond_signal(&s->cond);  
26     Mutex_unlock(&s->lock);  
27 }
```

- Zemaphore don't maintain the invariant that *the value of* the semaphore.
 - The value never be lower than zero.
 - This behavior is **easier** to implement and **matches** the current Linux implementation.

Semaphore (LINUX actual)

- **sem_wait** (*sem*)
 - decrements (locks) the semaphore pointed to by *sem*.
 - If the semaphore's value is greater than zero, then the decrement proceeds, and the function returns, immediately.
 - If the semaphore currently has the value zero, then the call blocks until either it becomes possible to perform the decrement (i.e., the semaphore value rises above zero), or a signal handler interrupts the call
- **sem_post** (*sem*)
 - increments (unlocks) the semaphore pointed to by *sem*.
 - If the semaphore's value consequently becomes greater than zero, then another process or thread blocked in a [sem_wait\(\)](#) call will be woken up and proceed to lock the semaphore.
- **sem_init(sem_t **sem*, int *pshared*, unsigned int *value*)**
 - *pshared* indicates sharing between threads(0)/processes(nonzero)

Condition Variable

Queue:

Queue:



wait()

Queue:



Queue



wait()

Queue



Queue:



signal()

Queue:



Queue:

signal()

Queue:

Queue:

signal()

Queue:
nothing to do!

signal()

Queue:



Queue:



wait()

Queue:



If we weren't careful, C may sleep forever.

Thread Queue: Signal Queue:

Thread Queue:

Signal Queue:



wait()

Thread Queue:



Signal Queue:

Thread Queue: Signal Queue:

signal()

Thread Queue: Signal Queue:

Thread Queue:

Signal Queue:



signal

signal()

Thread Queue:

Signal Queue:



signal

Thread Queue:



Signal Queue:



wait()

Thread Queue: Signal Queue:

wait()

Thread Queue: Signal Queue:

signal was not lost do to some race condition!

wait()

Thread Queue: Signal Queue:

Actual Implementation

Use **counter** instead of Signal Queue
- all signals are the same

If the **counter** is positive, don't bother to queue a thread upon wait().

CV's don't keep **extra state**, so CV users must.
Semaphores keep **extra state**, so users sometimes don't.

Actual Definition (see handout)

```
sem_init(sem_t *s, int initval) {  
    s->value = initval  
}
```

```
sem_wait(sem_t *s) {  
    s->value -= 1  
    wait if s->value < 0  
}
```

wait and post are atomic

```
sem_post(sem_t *s) {  
    s->value += 1  
    wake one waiting thread (if there  
    are any)  
}
```

Actual Definition (see handout)

```
sem_init(sem_t *s, int initval) {  
    s->value = initval  
}
```

```
sem_wait(sem_t *s) {  
    s->value -= 1  
    wait if s->value < 0  
}
```

value = 4:	4 waiting signals
value = -3:	3 waiting threads

```
sem_post(sem_t *s) {  
    s->value += 1  
    wake one waiting thread (if there  
    are any)  
}
```

Join example

Join is simpler with semaphores than CV's.

Join w/ CV

```
int done = 0;
mutex_t m = MUTEX_INIT;
cond_t c = COND_INIT;
```

```
    Mutex_lock(&m);
    done = 1;  cond_s
    ignal(&c);  Mutex
    _unlock(&m);
```

```
pthread_create(c, NULL, child, NULL);
Mutex_lock(&m); while(
done == 0)
    Cond_wait(&c, &m);
Mutex_unlock(&m);
```

Join w/ Semaphore

```
sem_t s;  
  
sem_post(&s);  
  
sem_init(&s, ?);  
  
Pthread_create(c, NULL, child, NULL);  
sem_wait(&s);
```

Join w/ Semaphore

```
sem_t s;
```

```
sem_post(&s);
```

```
sem_init(&s, ?, );
```

```
Pthread_create(c, NULL, child, NULL);  
sem_wait(&s);
```

Join w/ Semaphore

```
sem_t s;
```

```
sem_post(&s);
```

```
sem_init(&s,  ?); What is this int?
```

```
Pthread_create(c, NULL, child, NULL);  
sem_wait(&s);
```

Join w/ Semaphore

```
sem_t s;
```

```
sem_post(&s);
```

```
sem_init(&s, ?, );
```

```
Pthread_create(c, NULL, child, NULL);  
sem_wait(&s);
```


Join w/ Semaphore

```
sem_t s;
```

```
sem_post(&s);
```

Run it!
(sem-join.c)

```
sem_init(&s, 0);
```

```
pthread_create(c, NULL, child, NULL);  
sem_wait(&s);
```

Equivalence Claim

Semaphores are **equally powerful** to Locks+CVs.

- what does this mean?

Either may be more convenient, but that's not relevant.

Equivalence means we can **build each over the other**.

Condition Variables vs Semaphores

Condition variables have no state (other than waiting queue)

- Programmer must track additional state

Semaphores have state: track integer value

- State cannot be directly accessed by user program, but state determines behavior of semaphore operations

Semaphore Operations

Allocate and Initialize

```
sem_t sem;  
sem_init(sem_t *s, int initval) {  
    s->value = initval;  
}
```

User cannot read or write value directly after initialization

Wait or Test (sometime P() for Dutch word)

Waits until value of sem is > 0 , then decrements sem value

Signal or Increment or Post (sometime V() for Dutch)

Increment sem value, then wake a single waiter

Join with CV vs Semaphores

CVs:

```
void thread_join() {  
    Mutex_lock(&m);           // w  
    if (done == 0)           // x  
        Cond_wait(&c, &m); // y  
    Mutex_unlock(&m);         // z  
}
```

```
void thread_exit() {  
    Mutex_lock(&m);           // a  
    done = 1;                // b  
    Cond_signal(&c);          // c  
    Mutex_unlock(&m);         // d  
}
```

Semaphores:

Sem_wait(): Waits until value > 0, then decrement
Sem_post(): Increment value, then wake a single waiter

`sem_t s;`

`sem_init(&s, ???);`

Initialize to 0 (so sem_wait() must wait...)

```
void thread_join() {  
    sem_wait(&s);  
}
```

```
void thread_exit() {  
    sem_post(&s)  
}
```

Equivalence Claim

Semaphores are equally powerful to **Locks+CVs**




- what does this mean?

One might be more convenient, but that's not relevant

Equivalence means each can be built from the other

Proof Steps

Want to show we can do these three things:



Build Lock from Semaphore

```
typedef struct __lock_t {  
    // whatever data structs you need go here  
} lock_t;
```

```
void init(lock_t *lock) {  
}
```

```
void acquire(lock_t *lock) {  
}
```

```
void release(lock_t *lock) {  
}
```

Sem_wait(): Waits until value > 0, then decrement

Sem_post(): Increment value, then wake a single waiter

Locks

Semaphores

Build Lock from Semaphore

```
typedef struct __lock_t {  
    sem_t sem;  
} lock_t;  
  
void init(lock_t *lock) {  
    sem_init(&lock->sem, ??);  
}  
void acquire(lock_t *lock) {  
    sem_wait(&lock->sem);  
}  
void release(lock_t *lock) {  
    sem_post(&lock->sem);  
}
```

1 → 1 thread can grab lock

Sem_wait(): Waits until value > 0, then decrement


Sem_post(): Increment value, then wake a single waiter

Locks

Semaphores

Building CV's over Semaphores

Possible, but really hard to do right



Read about Microsoft Research's attempts:

<http://research.microsoft.com/pubs/64242/ImplementingCVs.pdf>

Build Semaphore from Lock and CV

```
Typedef struct {  
    // what goes here?
```

```
} sem_t;
```

```
Void sem_init(sem_t *s, int value) {  
    // what goes here?
```

```
}
```

Sem_wait(): Waits until value > 0, then decrement
Sem_post(): Increment value, then wake a single waiter

Semaphores

Locks

CV's

Build Semaphore from Lock and CV

```
Typedef struct {  
    int value;  
    cond_t cond;  
    lock_t lock;  
} sem_t;
```

```
Void sem_init(sem_t *s, int value) {  
    s->value = value;  
    cond_init(&s->cond);  
    lock_init(&s->lock);  
}
```

Sem_wait(): Waits until value > 0, then decrement
Sem_post(): Increment value, then wake a single waiter

Semaphores

Locks

CV's

Build Semaphore from Lock and CV

```
Sem_wait(sem_t *s) {  
    // what goes here?
```

```
Sem_post(sem_t *s) {  
    // what goes here?
```

```
}
```

```
}
```

Semaphores

Locks

CV's

Sem_wait(): Waits until value > 0 , then decrement

Sem_post(): Increment value, then wake a single waiter

Build Semaphore from Lock and CV

```
Sem_wait(sem_t *s) {  
    lock_acquire(&s->lock);  
    // this stuff is atomic  
  
    lock_release(&s->lock);  
}
```

```
Sem_post(sem_t *s) {  
    lock_acquire(&s->lock);  
    // this stuff is atomic  
  
    lock_release(&s->lock);  
}
```

Sem_wait(): Waits until value > 0, then decrement
Sem_post(): Increment value, then wake a single waiter

Semaphores

Locks

CV's

Build Semaphore from Lock and CV

```
Sem_wait(sem_t *s) {  
    lock_acquire(&s->lock);  
    while (s->value <= 0)  
        cond_wait(&s->cond);  
    s->value--;  
    lock_release(&s->lock);  
}
```

```
Sem_post(sem_t *s) {  
    lock_acquire(&s->lock);  
    // this stuff is atomic  
    lock_release(&s->lock);  
}
```

Sem_wait(): Waits until value > 0, then decrement
Sem_post(): Increment value, then wake a single waiter

Semaphores

Locks

CV's

Build Semaphore from Lock and CV

```
Sem_wait(sem_t *s) {  
    lock_acquire(&s->lock);  
    while (s->value <= 0)  
        cond_wait(&s->cond);  
    s->value--;  
    lock_release(&s->lock);  
}
```

```
Sem_post(sem_t *s) {  
    lock_acquire(&s->lock);  
    s->value++;  
    cond_signal(&s->cond);  
    lock_release(&s->lock);  
}
```

Sem_wait(): Waits until value > 0, then decrement
Sem_post(): Increment value, then wake a single waiter

Semaphores

Locks

CV's

Producer/Consumer: Semaphores #1

Simplest case:

- Single producer thread, single consumer thread
- Single shared buffer between producer and consumer

Requirements

- Consumer must wait for producer to fill buffer
- Producer must wait for consumer to empty buffer (if filled)

Requires 2 semaphores

- emptyBuffer: Initialize to ??? **1 → 1 empty buffer; producer can run 1 time first**
- fullBuffer: Initialize to ??? **0 → 0 full buffers; consumer can run 0 times first**

Producer

While (1) {

```
    sem_wait(&emptyBuffer);  
    Fill(&buffer);  
    sem_signal(&fullBuffer);
```

}

Consumer

While (1) {

```
    sem_wait(&fullBuffer);  
    Use(&buffer);  
    sem_signal(&emptyBuffer);
```

}

Producer/Consumer: Semaphores #2

Next case: **Circular Buffer**

- Single producer thread, single consumer thread
- Shared buffer with **N** elements between producer and consumer

Requires 2 semaphores

- emptyBuffer: Initialize to ???
 - fullBuffer: Initialize to ???
- $N \rightarrow N$ empty buffers; producer can run N times first
 $0 \rightarrow 0$ full buffers; consumer can run 0 times first

Producer

```
i = 0;
While (1) {
    sem_wait(&emptyBuffer);
    Fill(&buffer[i]);
    i = (i+1)%N;
    sem_signal(&fullBuffer);
}
```

Consumer

```
j = 0;
While (1) {
    sem_wait(&fullBuffer);
    Use(&buffer[j]);
    j = (j+1)%N;
    sem_signal(&emptyBuffer);
}
```

Producer/Consumer: Semaphore #3

Final case:

- **Multiple producer threads, multiple consumer threads**
- Shared buffer with N elements between producer and consumer

Requirements

- Each consumer must grab unique filled element
- Each producer must grab unique empty element
- **Why will previous code (shown below) not work???**

Producer

```
i = 0;
While (1) {
    sem_wait(&emptyBuffer);
    Fill(&buffer[i]);
    i = (i+1)%N;
    sem_signal(&fullBuffer);
}
```

Consumer

```
j = 0;
While (1) {
    sem_wait(&fullBuffer);
    Use(&buffer[j]);
    j = (j+1)%N;
    sem_signal(&emptyBuffer);
}
```

Are i and j private or shared? Need each producer to grab unique buffer

Producer/Consumer: Multiple Threads

Final case:

- Multiple producer threads, multiple consumer threads
- Shared buffer with N elements between producer and consumer

Requirements

- Each consumer must grab unique filled element
- Each producer must grab unique empty element

Producer

```
While (1) {  
    sem_wait(&emptyBuffer);  
    myi = findempty(&buffer);  
    Fill(&buffer[myi]);  
    sem_signal(&fullBuffer);  
}
```

Consumer

```
While (1) {  
    sem_wait(&fullBuffer);  
    myj = findfull(&buffer);  
    Use(&buffer[myj]);  
    sem_signal(&emptyBuffer);  
}
```

Are myi and myj private or shared? Where is mutual exclusion needed???

Producer/Consumer: Multiple Threads

Consider three possible locations for mutual exclusion
Which work??? Which is best???

Producer #1

```
sem_wait(&mutex);  
sem_wait(&emptyBuffer);  
myi = findempty(&buffer);  
Fill(&buffer[myi]);  
sem_signal(&fullBuffer);  
sem_signal(&mutex);
```

Consumer #1

```
sem_wait(&mutex);  
sem_wait(&fullBuffer);  
myj = findfull(&buffer);  
Use(&buffer[myj]);  
sem_signal(&emptyBuffer);  
sem_signal(&mutex);
```

Problem: Deadlock at mutex (e.g., consumer runs first; won't release mutex)

Producer/Consumer: Multiple Threads

Consider three possible locations for mutual exclusion

Which work??? Which is best???

Producer #2

```
sem_wait(&emptyBuffer);  
sem_wait(&mutex);  
myi = findempty(&buffer);  
Fill(&buffer[myi]);  
sem_signal(&mutex);  
sem_signal(&fullBuffer);
```

Consumer #2

```
sem_wait(&fullBuffer);  
sem_wait(&mutex);  
myj = findfull(&buffer);  
Use(&buffer[myj]);  
sem_signal(&mutex);  
sem_signal(&emptyBuffer);
```

Works, but limits concurrency:
Only 1 thread at a time can be using or filling different buffers

Producer/Consumer: Multiple Threads

Consider three possible locations for mutual exclusion

Which work??? Which is best???

Producer #3

```
sem_wait(&emptyBuffer);  
sem_wait(&mutex);  
myi = findempty(&buffer);  
sem_signal(&mutex);  
Fill(&buffer[myi]);  
sem_signal(&fullBuffer);
```

Consumer #3

```
sem_wait(&fullBuffer);  
sem_wait(&mutex);  
myj = findfull(&buffer);  
sem_signal(&mutex);  
Use(&buffer[myj]);  
sem_signal(&emptyBuffer);
```

Works and increases concurrency; only finding a buffer is protected by mutex;
Filling or Using different buffers can proceed concurrently

Reader/Writer Locks

Goal:

Let multiple reader threads grab lock (shared)

Only one writer thread can grab lock (exclusive)

- No reader threads
- No other writer threads

Let us see if we can understand code...

Reader/Writer Locks

```
1 typedef struct _rwlock_t {  
2     sem_t lock;  
3     sem_t writelock;  
4     int readers;  
5 } rwlock_t;  
6  
7 void rwlock_init(rwlock_t *rw) {  
8     rw->readers = 0;  
9     sem_init(&rw->lock, 1);  
10    sem_init(&rw->writelock, 1);  
11 }  
12
```

Reader/Writer Locks

```
13 void rwlock_acquire_readlock(rwlock_t *rw) {
14     sem_wait(&rw->lock);
15     rw->readers++;
16     if (rw->readers == 1)
17         sem_wait(&rw->writelock);
18     sem_post(&rw->lock);
19 }
21 void rwlock_release_readlock(rwlock_t *rw) {
22     sem_wait(&rw->lock);
23     rw->readers--;
24     if (rw->readers == 0)
25         sem_post(&rw->writelock);
26     sem_post(&rw->lock);
27 }
29 rwlock_acquire_writelock(rwlock_t *rw) { sem_wait(&rw->writelock); }
31 rwlock_release_writelock(rwlock_t *rw) { sem_post(&rw->writelock); }
```

T1: acquire_readlock()
T2: acquire_readlock()
T3: acquire_writelock()
T2: release_readlock()
T1: release_readlock()
T4: acquire_readlock()
T5: acquire_readlock() // ???
T3: release_writelock()
// what happens???

Semaphores

Semaphores are equivalent to locks + condition variables

- Can be used for both mutual exclusion and ordering

Semaphores contain **state**

- How they are initialized depends on how they will be used
- Init to 1: Mutex
- Init to 0: Join (1 thread must arrive first, then other)
- Init to N: Number of available resources

Sem_wait(): Waits until value > 0 , then decrement (atomic)

Sem_post(): Increment value, then wake a single waiter (atomic)

Can use semaphores in producer/consumer relationships and for reader/writer locks

Summary

Locks+CVs are good primitives, but not always **convenient**.

Possible to build other abstractions such as semaphores.

Advice: if you always use an abstraction the same way,
build another abstraction **over the first!**